

Flocking and inverse design

these slides:

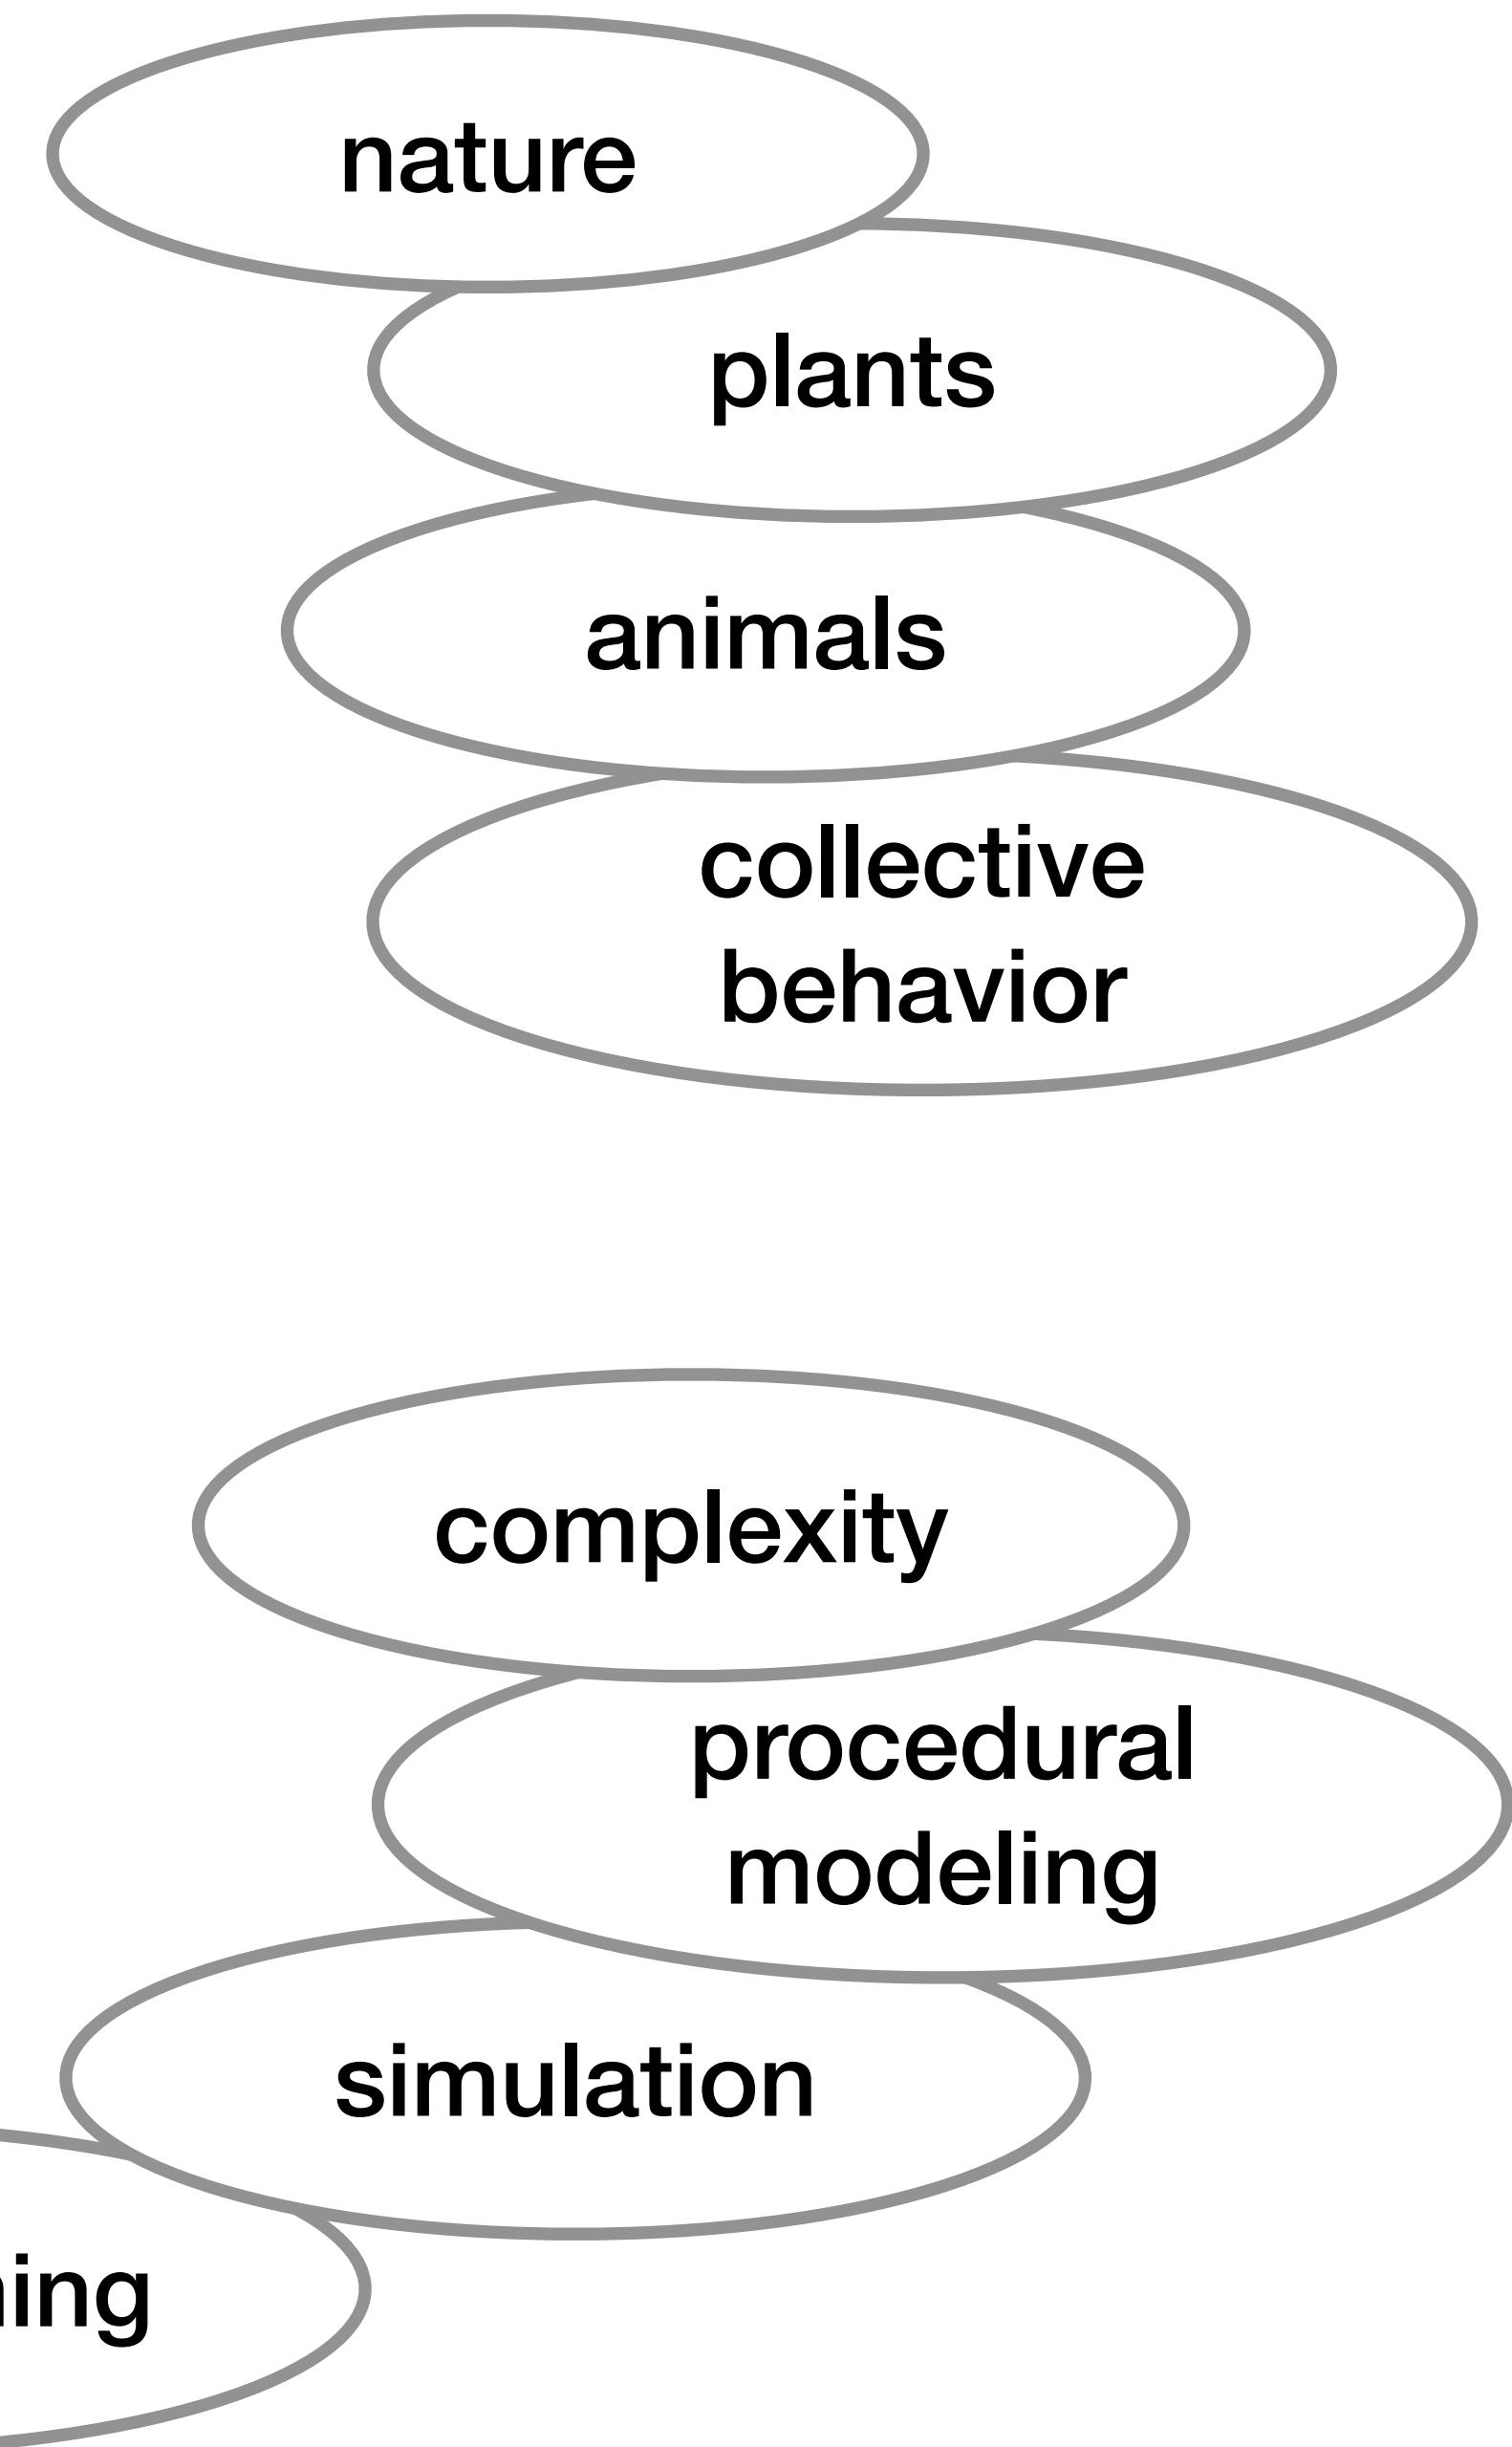
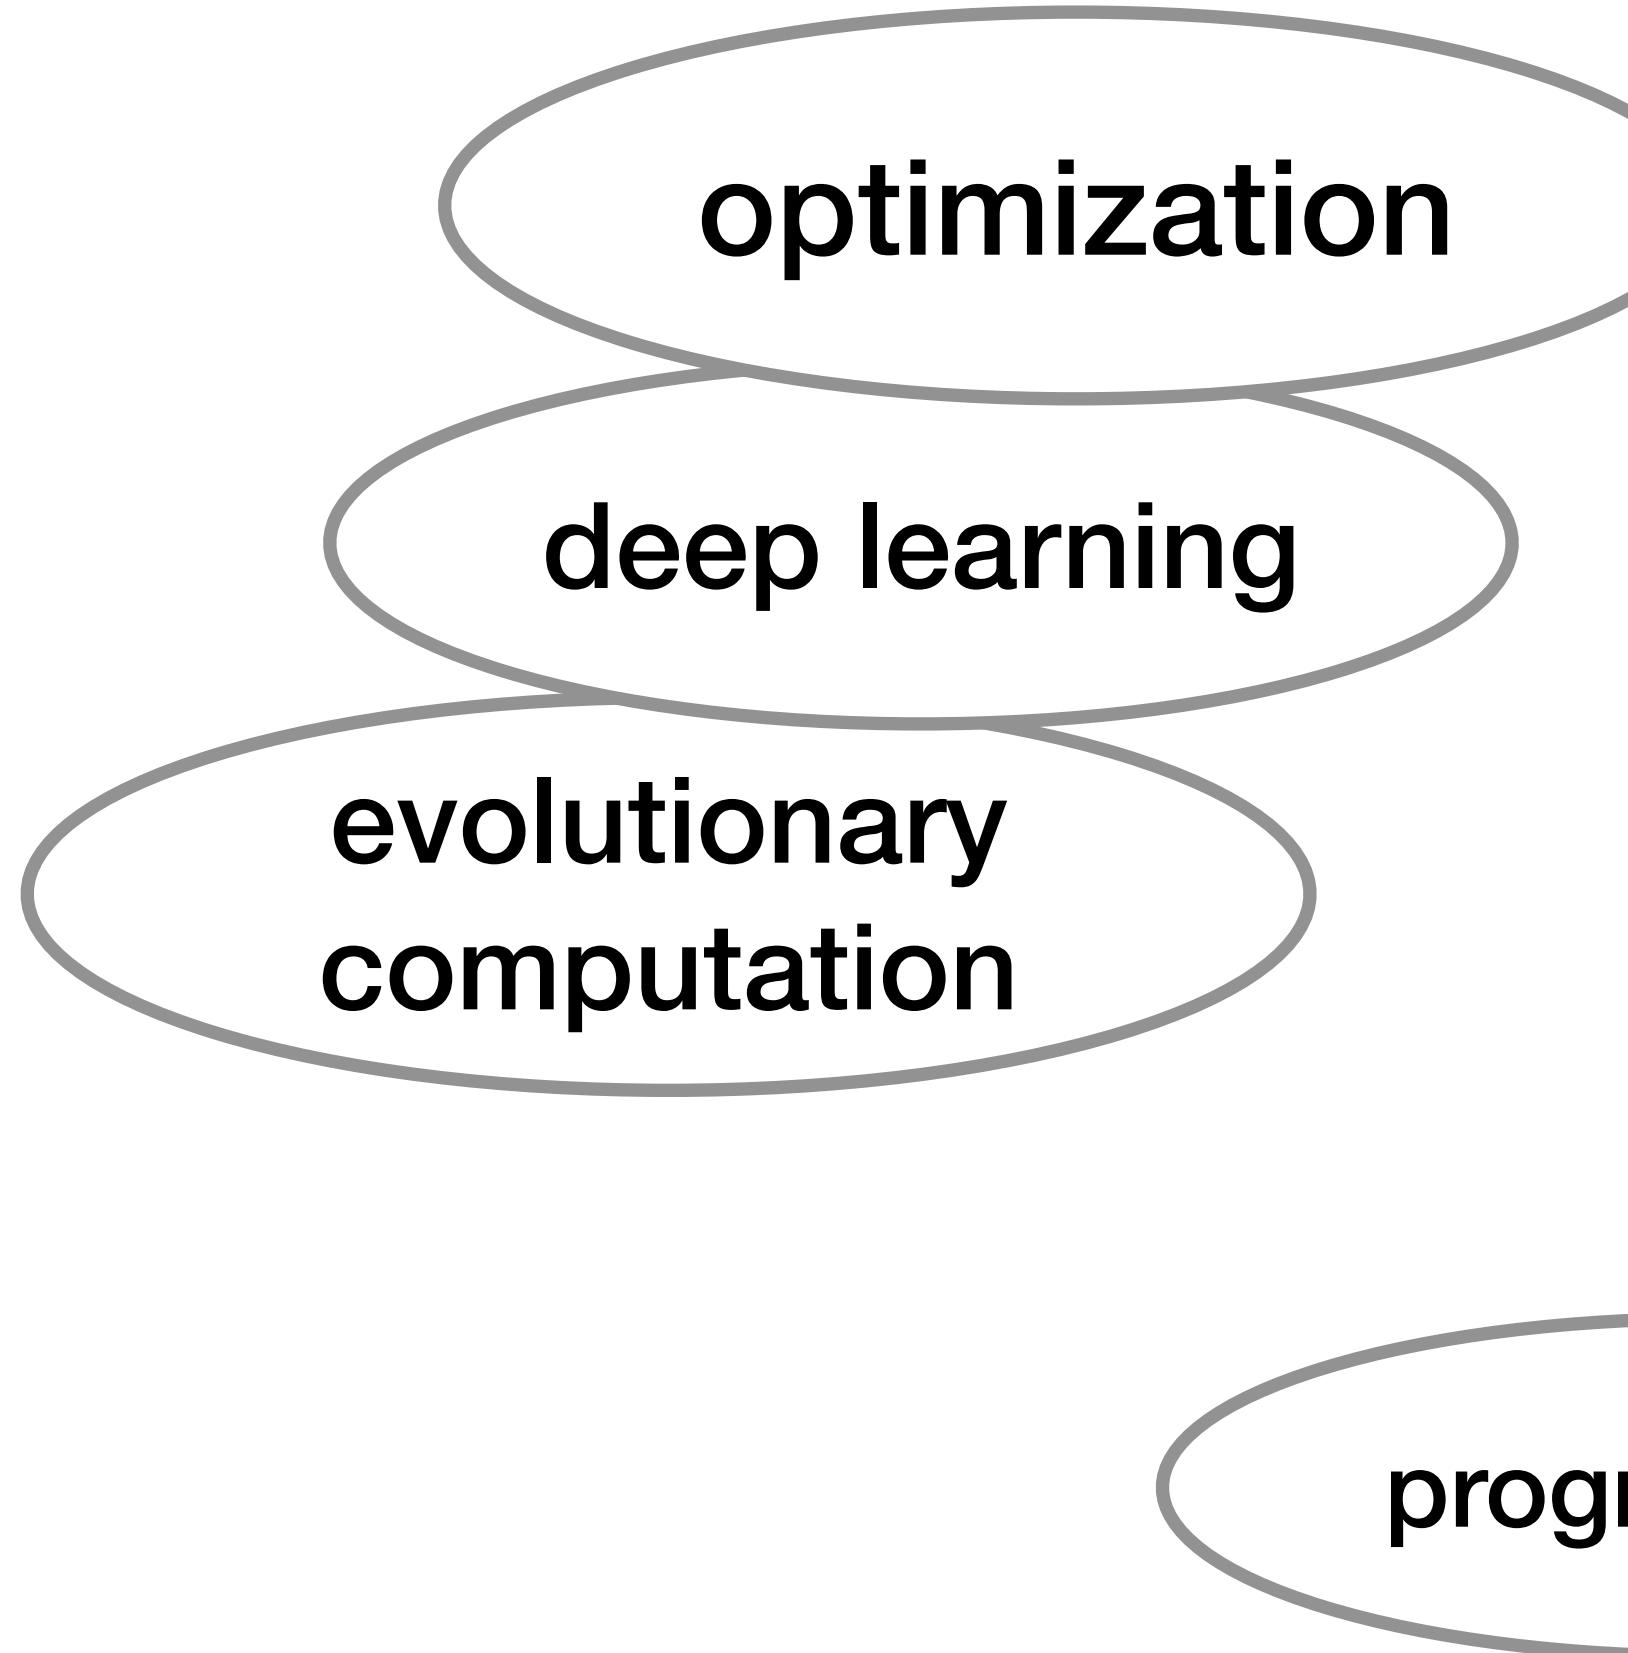
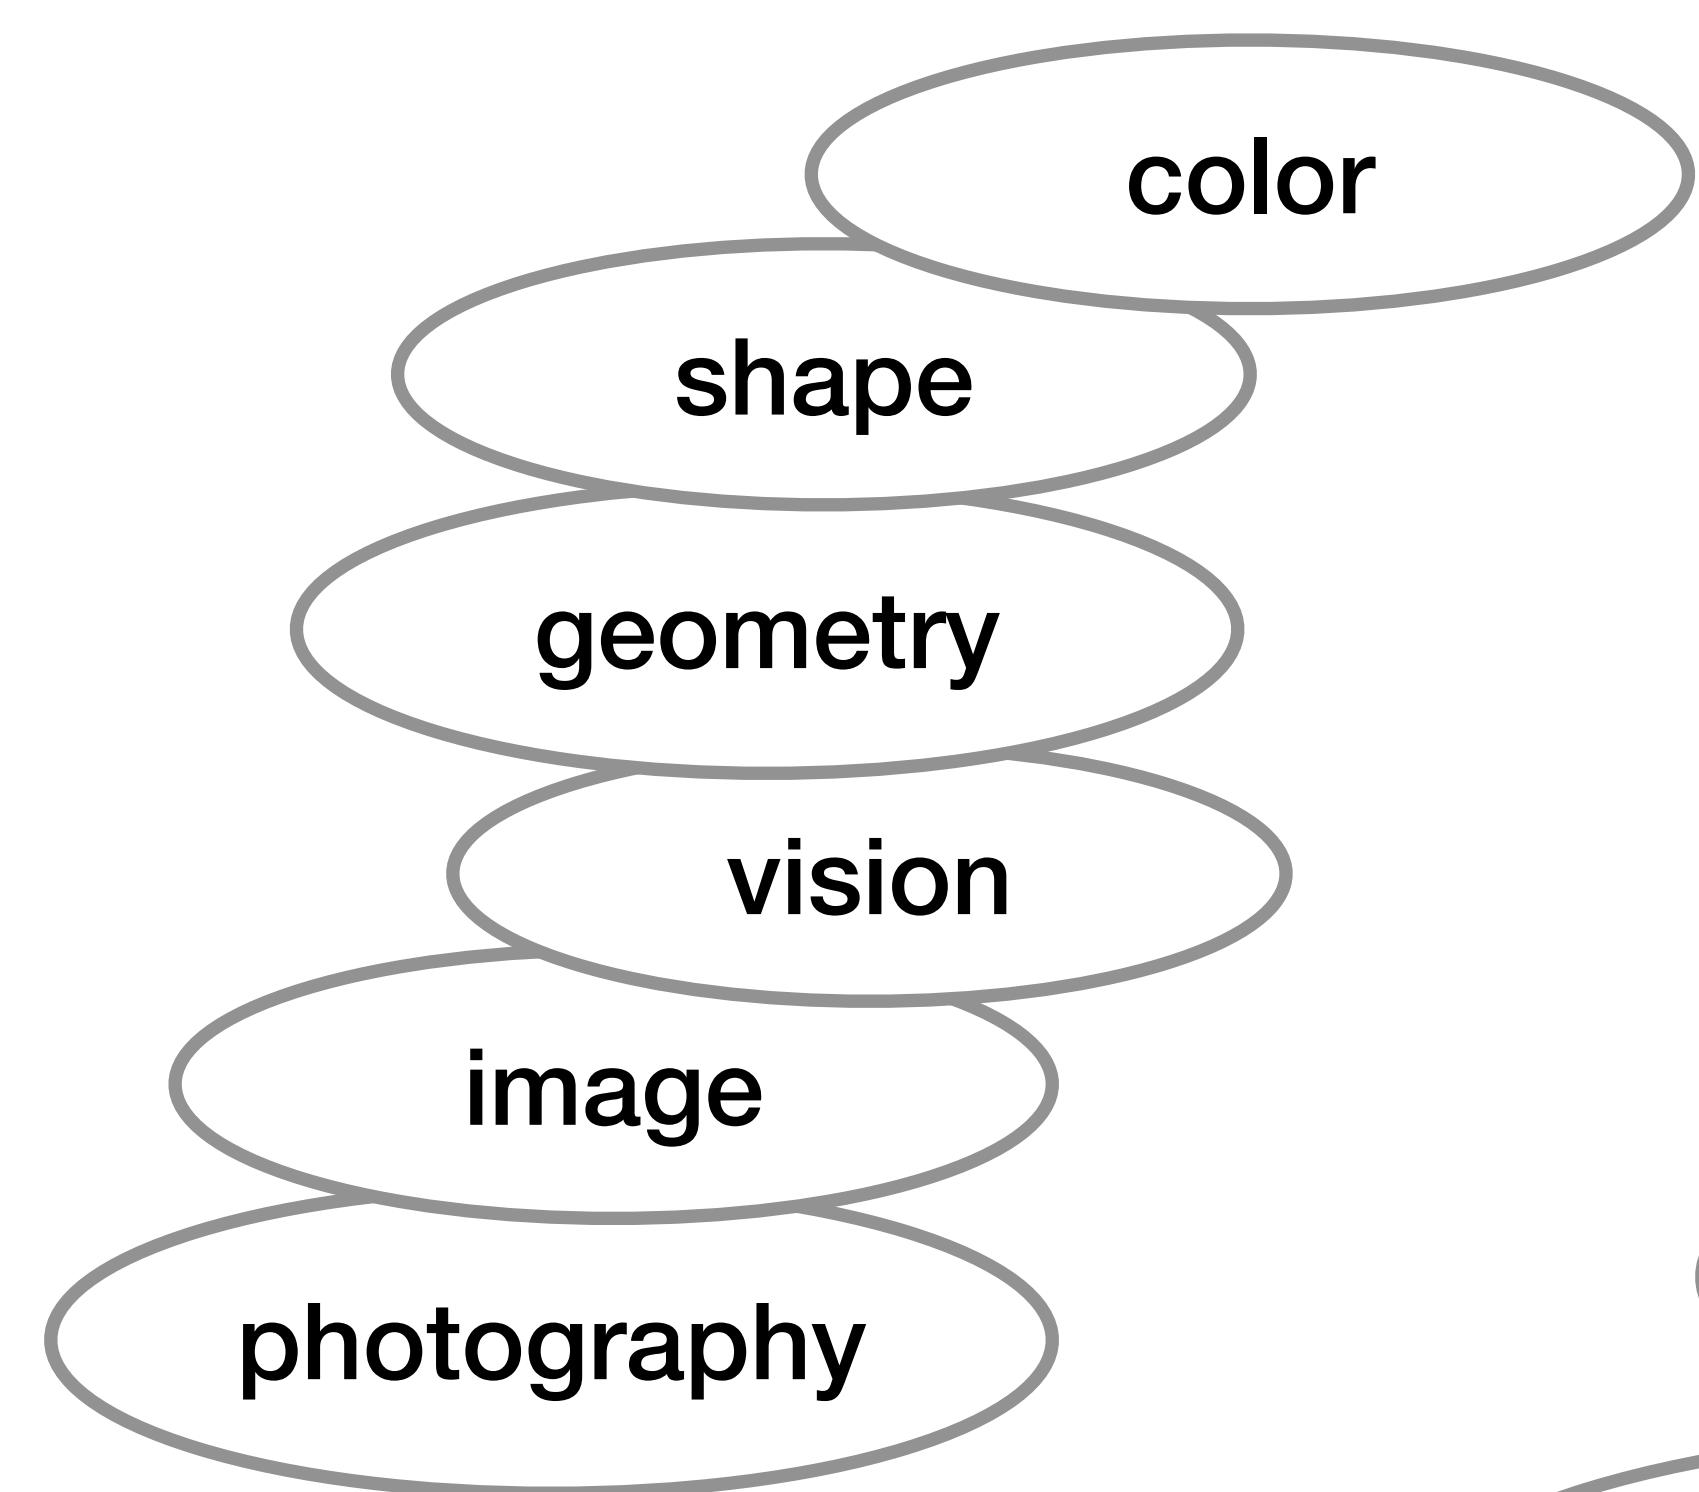


Craig Reynolds – Game AI – UCSC – August 6, 2025

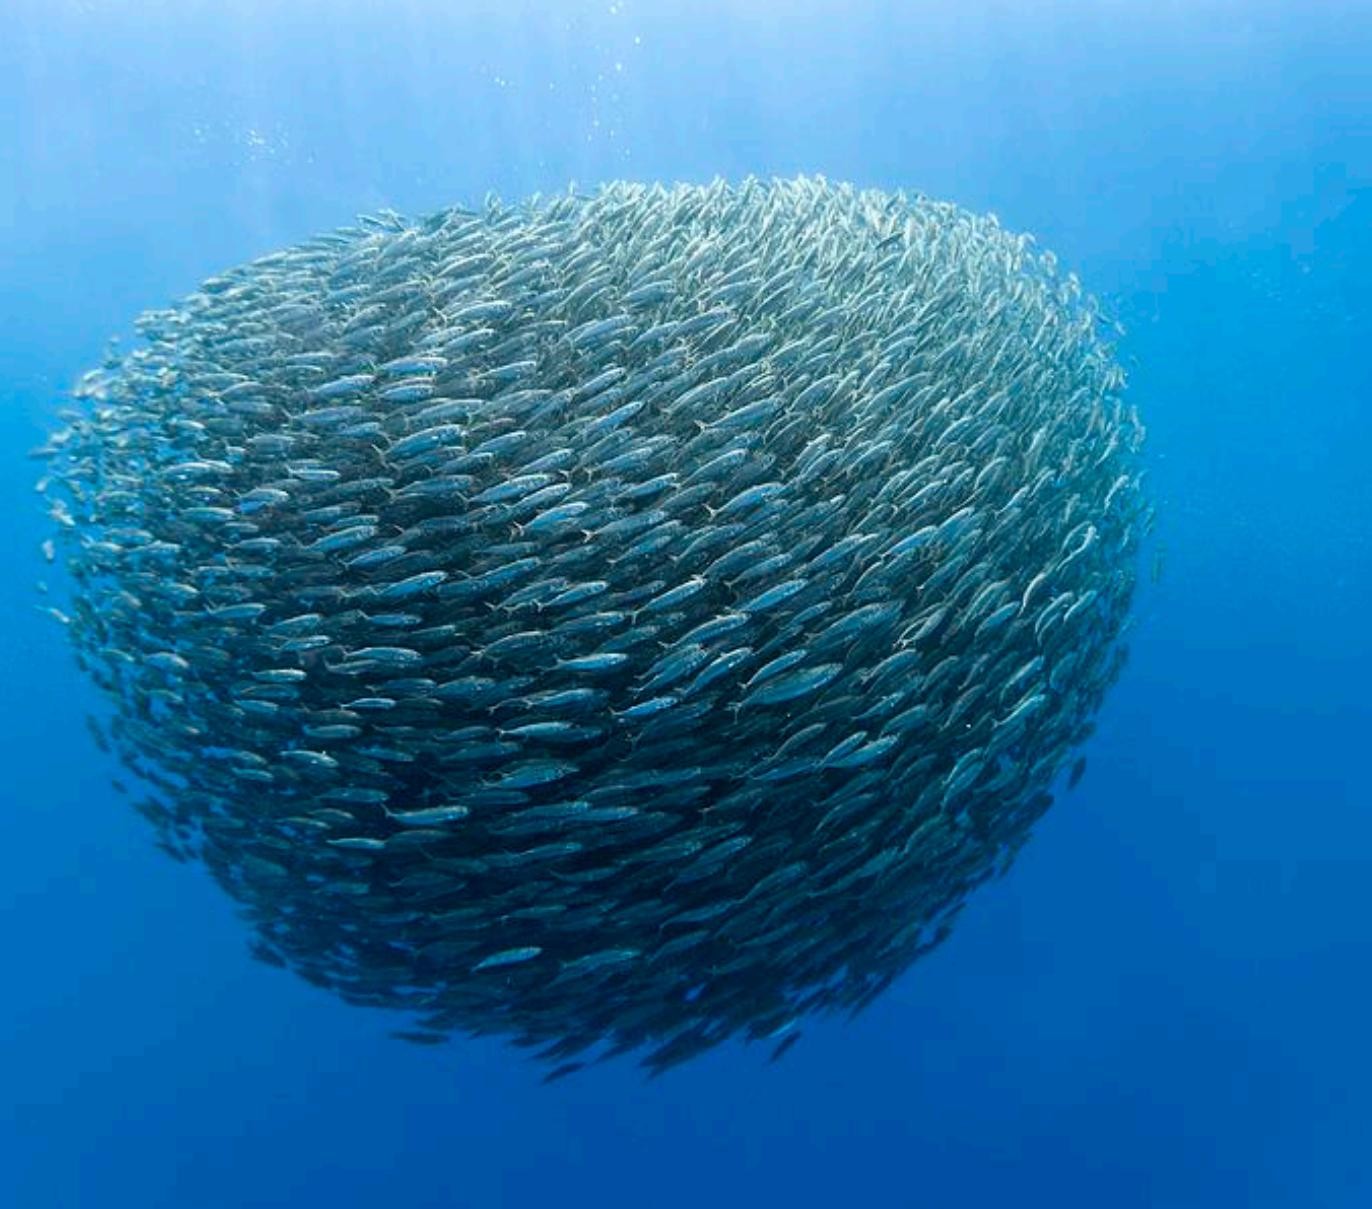
Background

I am particularly interested
in procedural models of
complex natural systems.

- **MIT**: CS, procedural animation (BS 1975, MS 1978)
- **triple-I**: *Juggler*, *Looker*, *TRON*
- **Symbolics**: digital content creation tools, boids, *Breaking the Ice*
- **Electronic Arts**: “non-player characters” for games, steering behaviors
- **SGI Silicon Studio**: DCC tools for games, steering behaviors
- **PlayStation US R&D**: PScrowd, evolution of camouflage
- **UC Santa Cruz (Playable Media)**: human behaviors for “serious games”
- **RightHook**: vehicle/pedestrian behaviors for virtual testing of self-driving
- “**unaffiliated researcher**” retired since 2020



A childhood fascination with natural complexity





bird shaped murmuration on YouTube

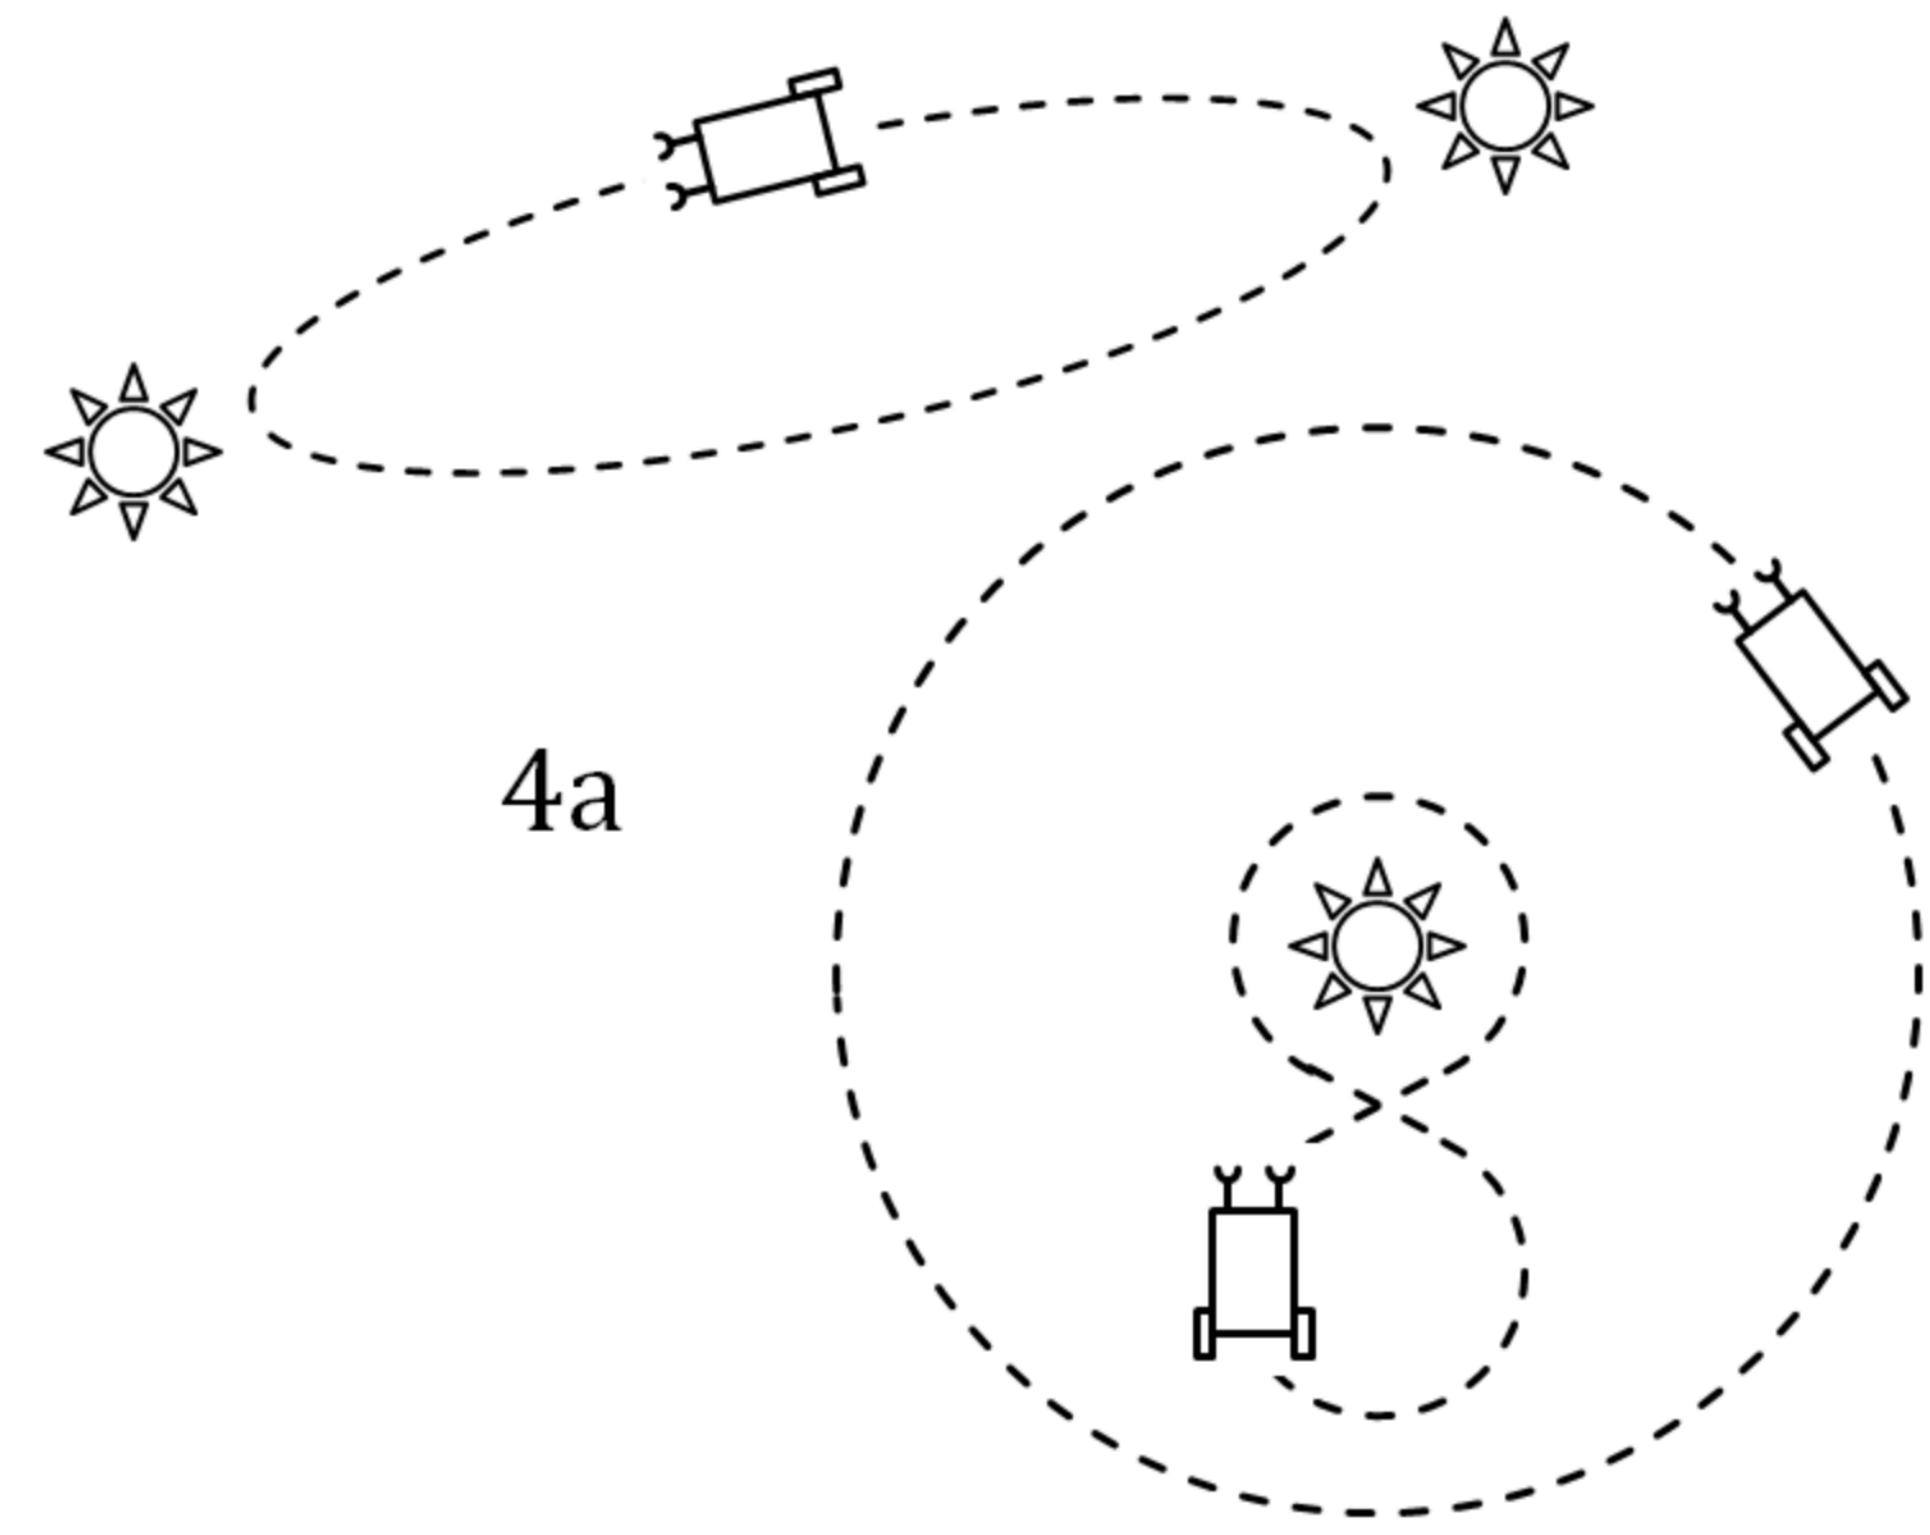
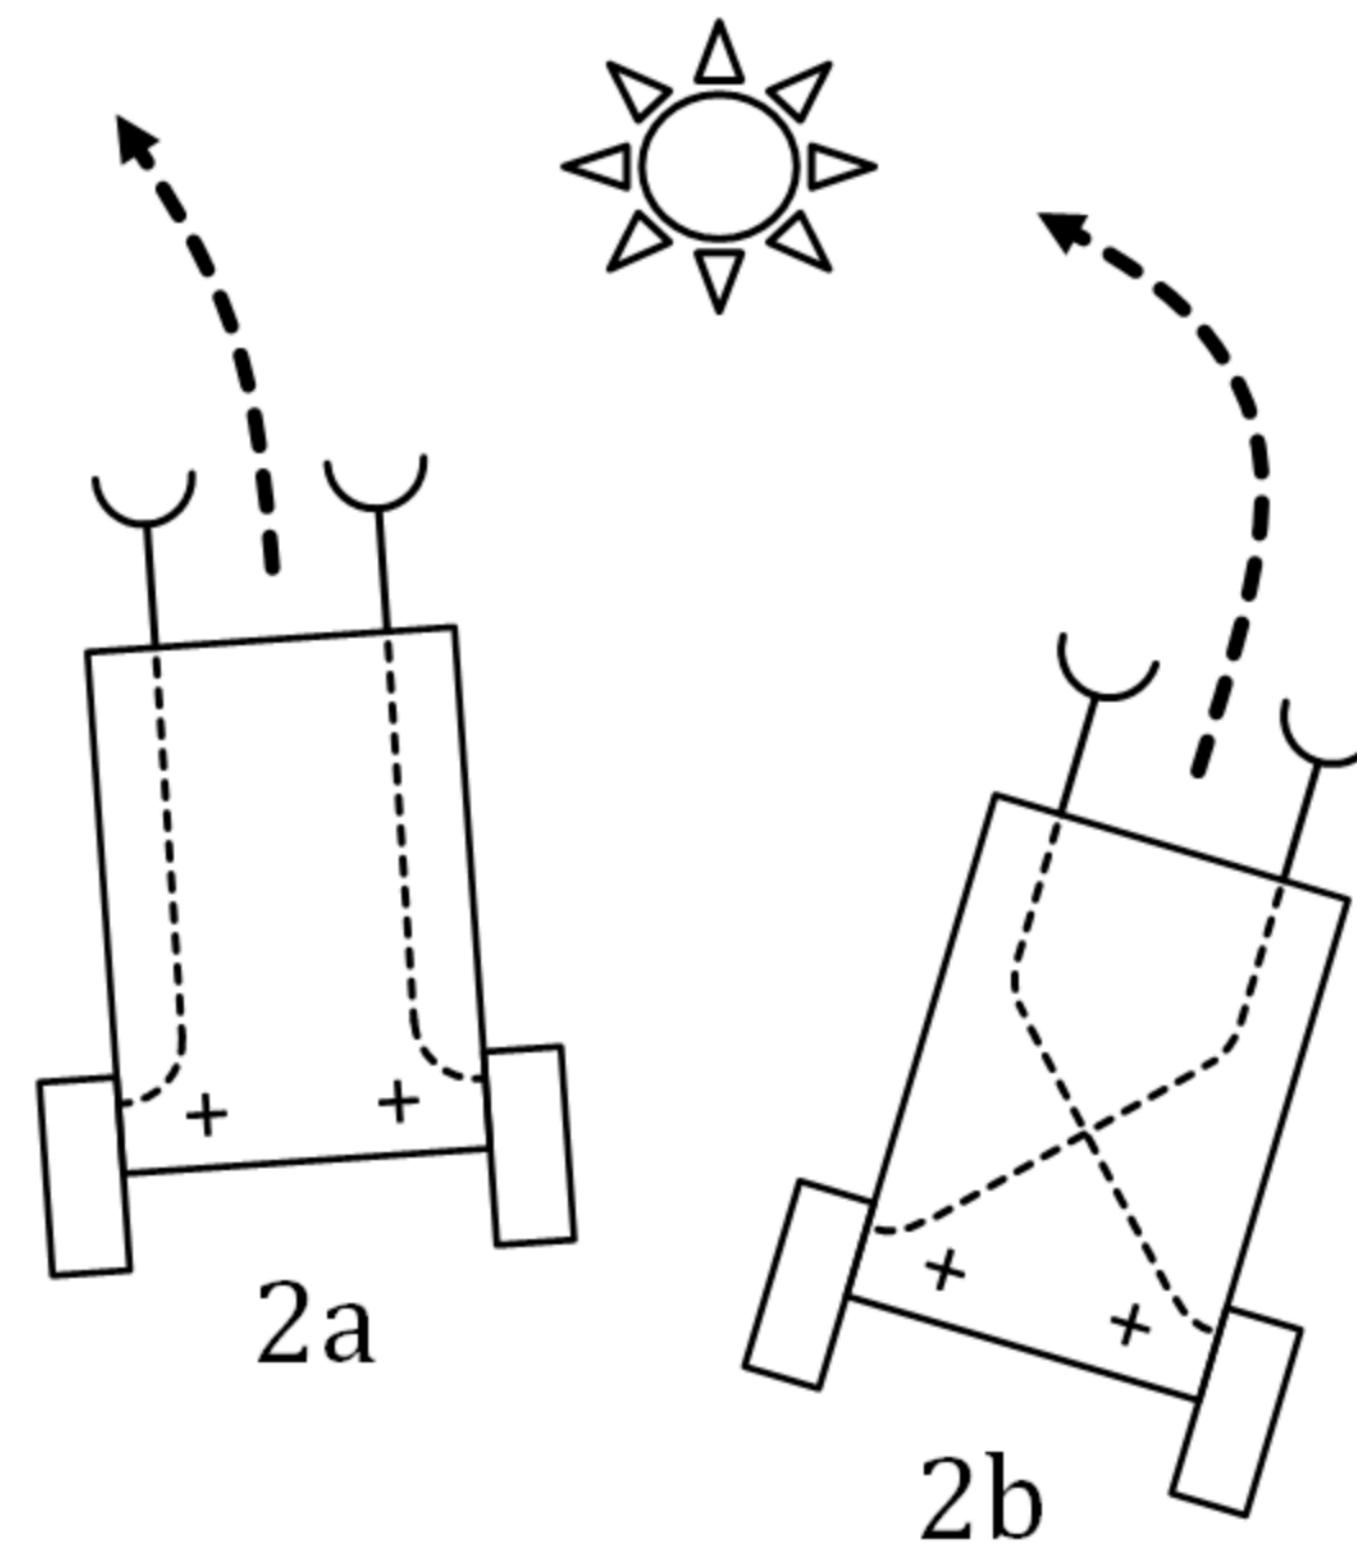
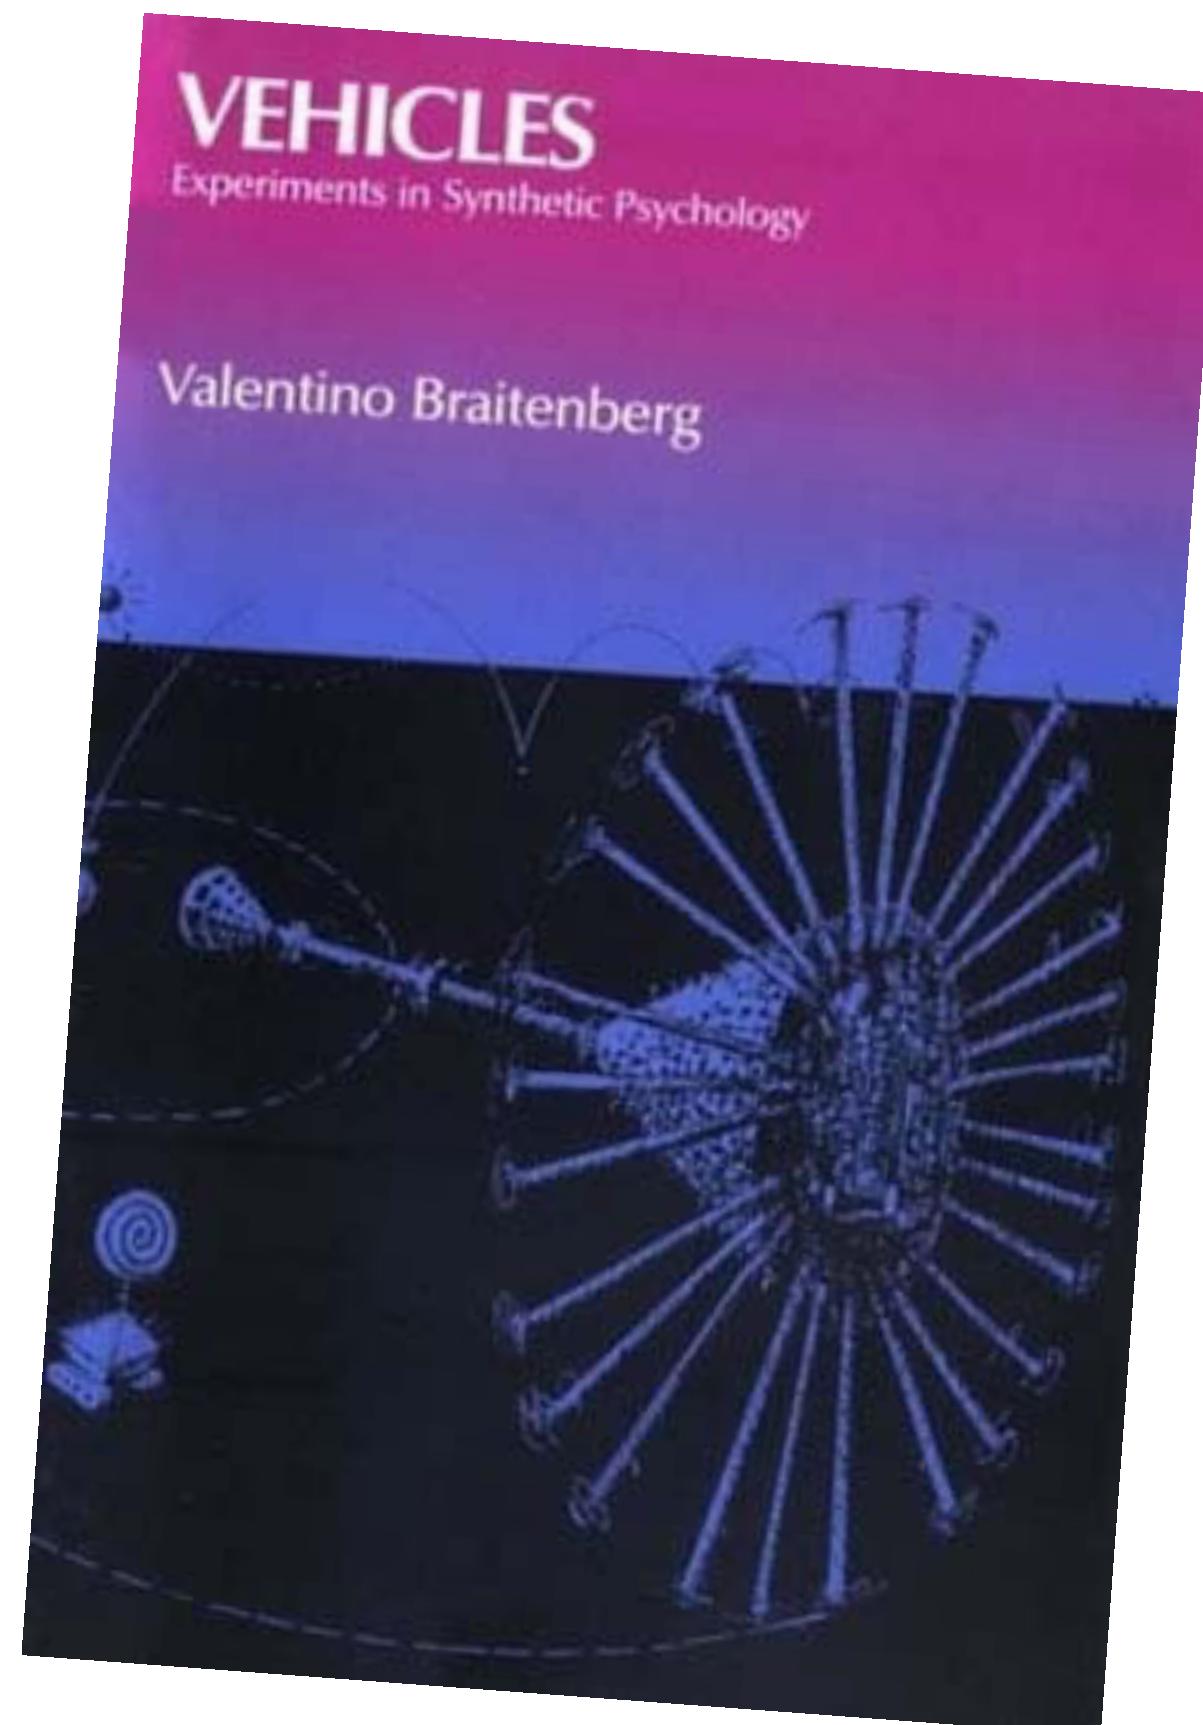
Early boids
1986-1987

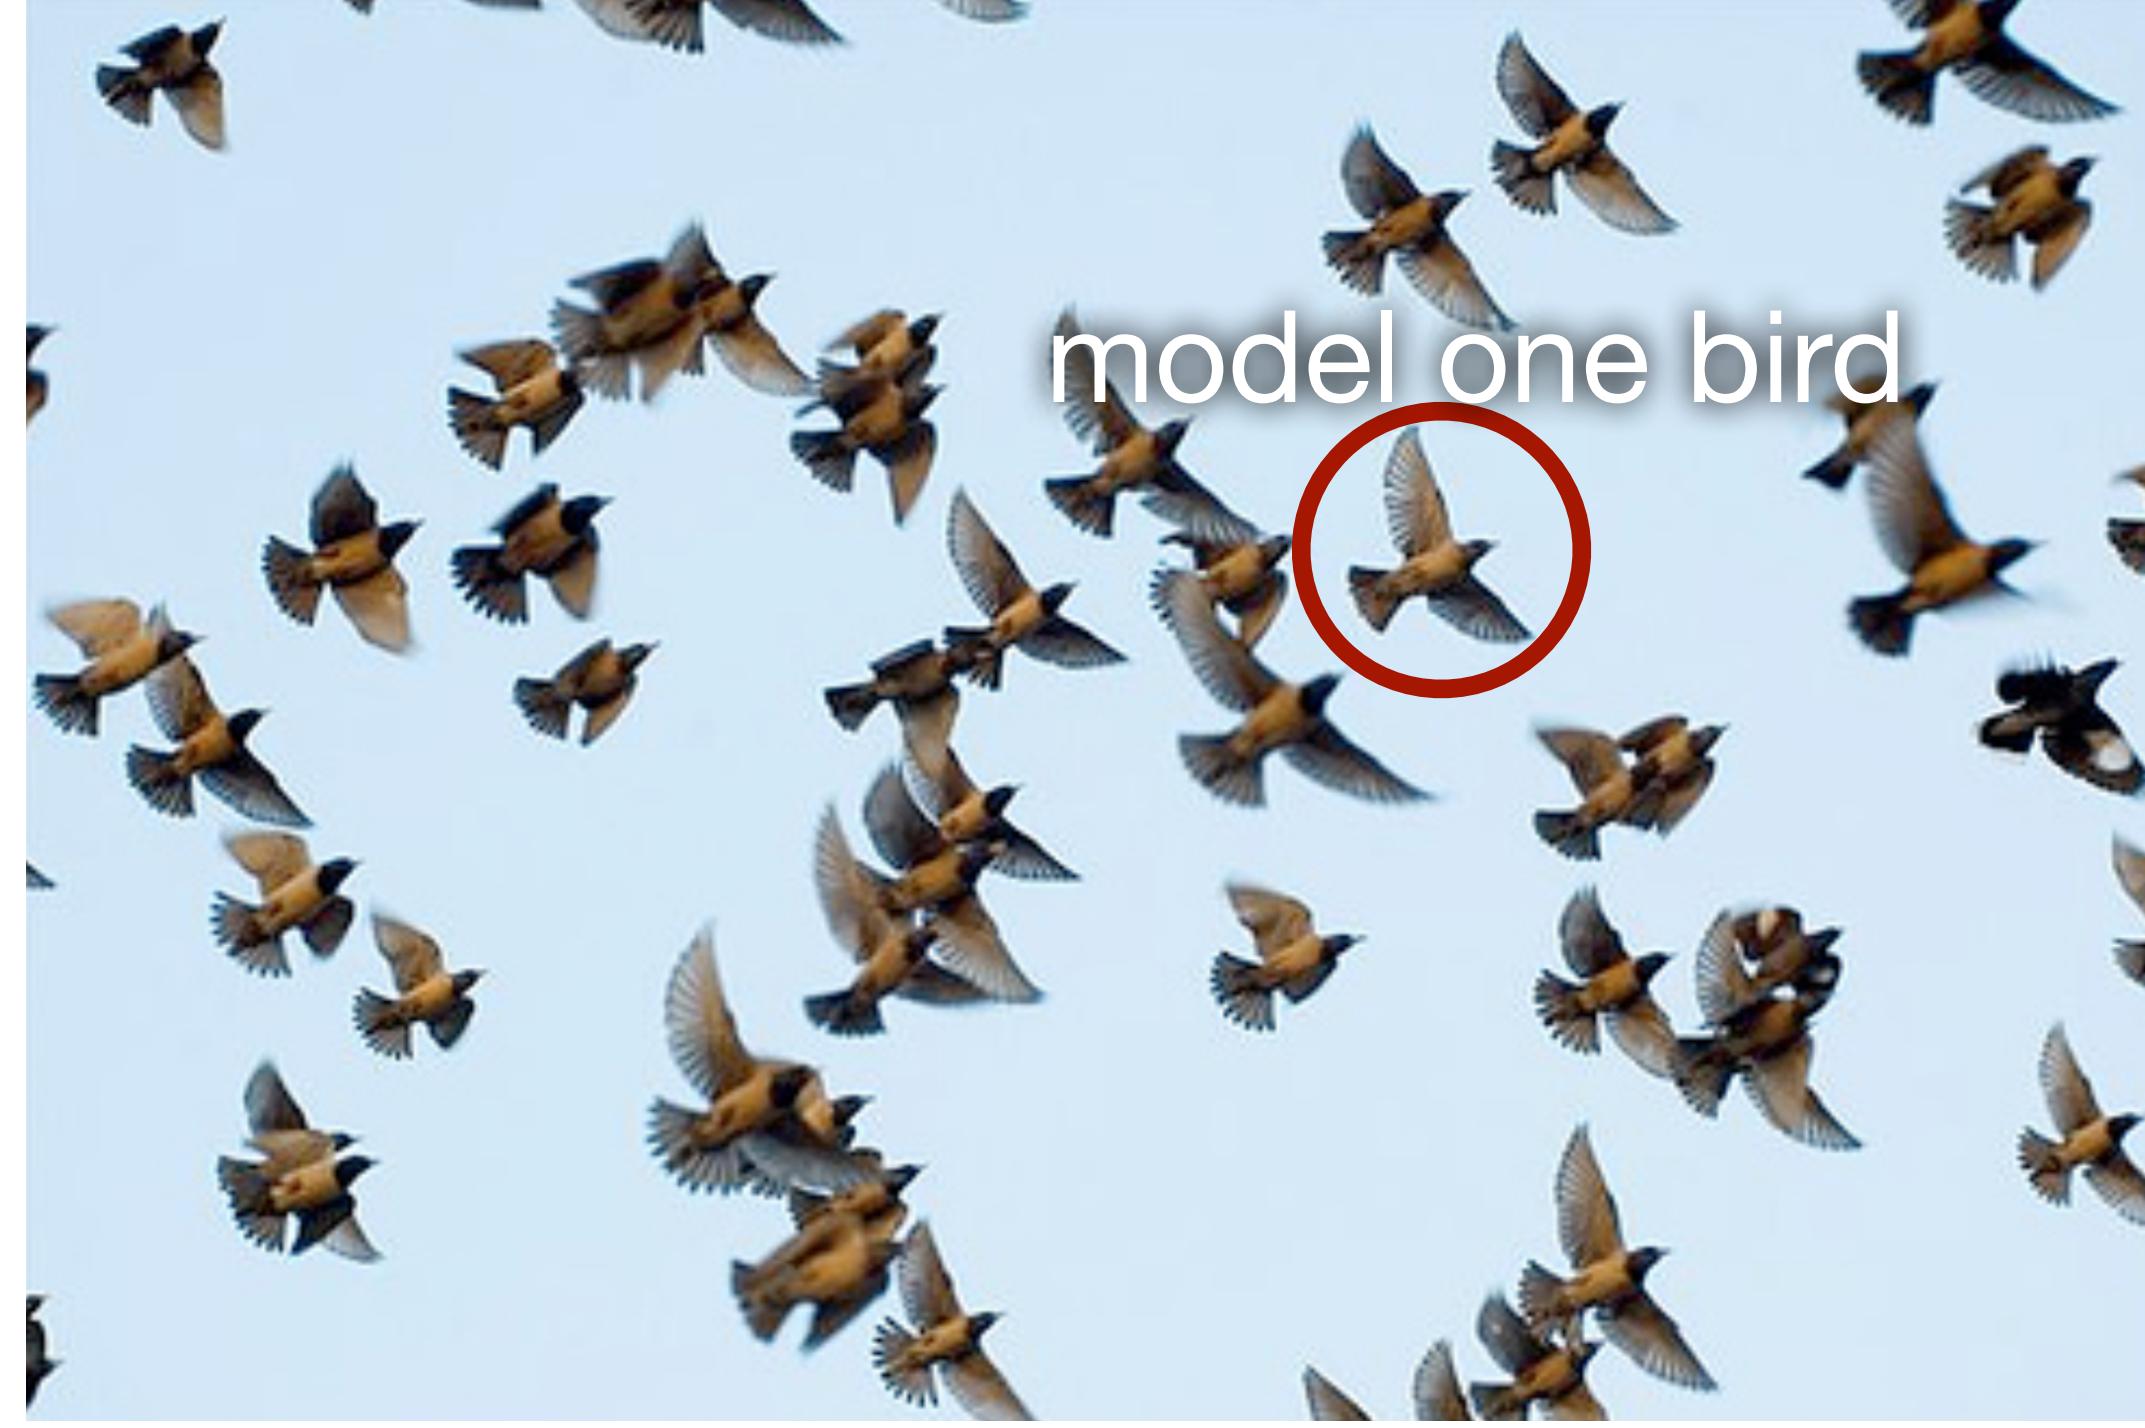
What is it like to be a bird in a flock?

Could the behavior of flocking birds
be simulated in software?



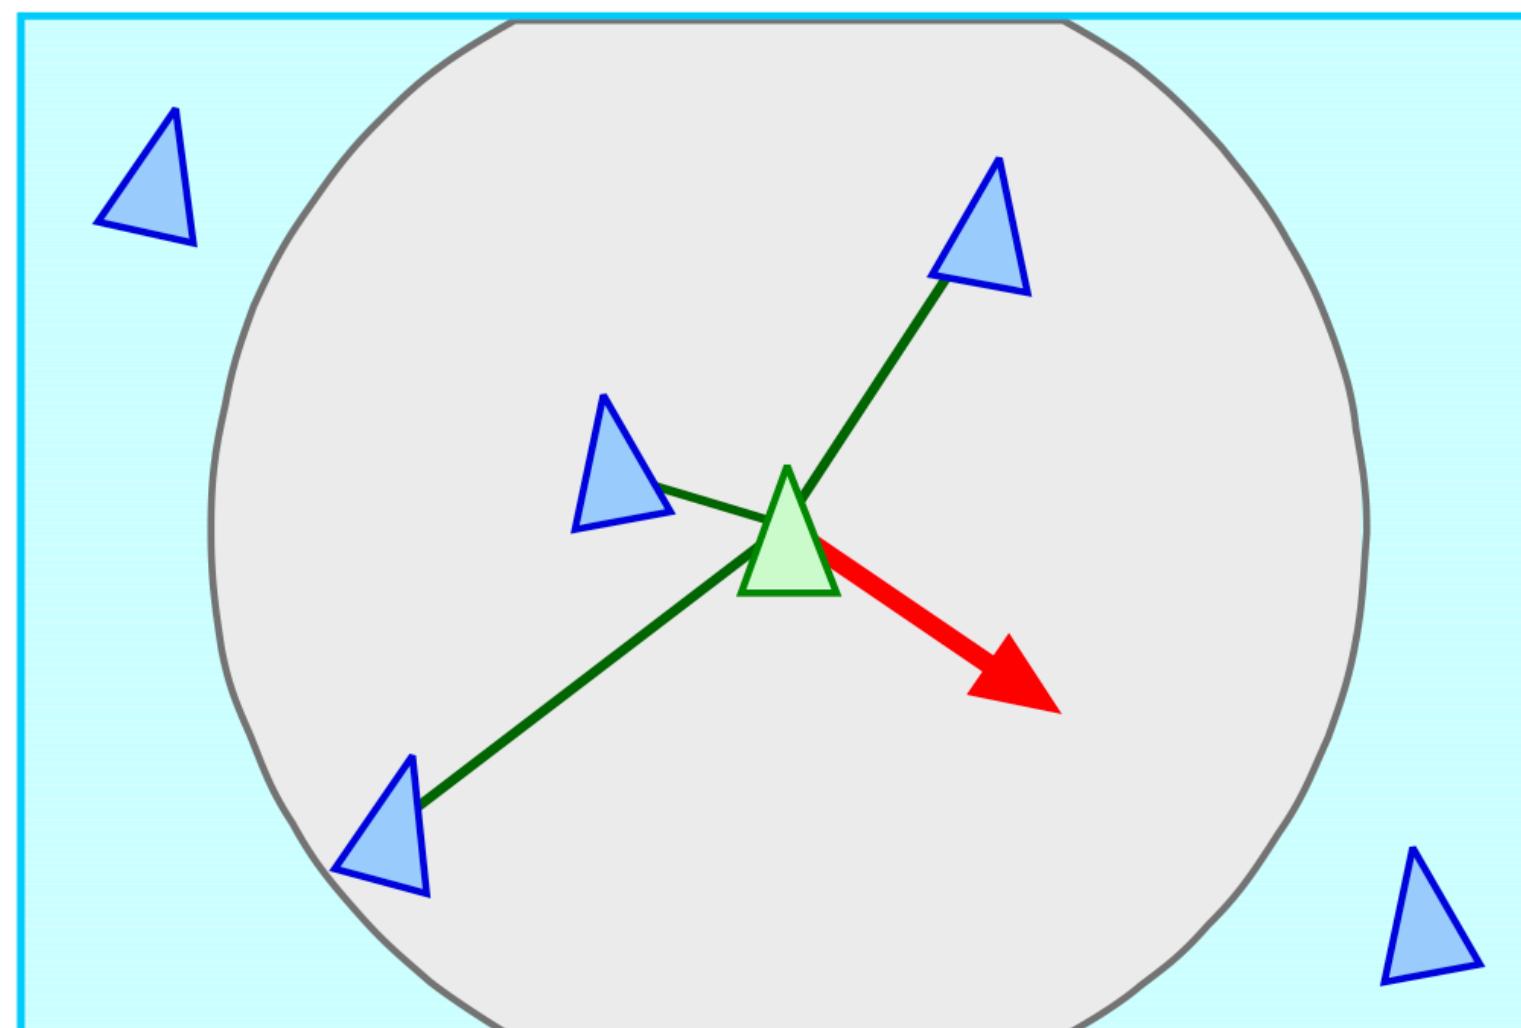
Key idea: a shift in perspective
from an external view of a complex system
to an internal local view of a few neighbors.



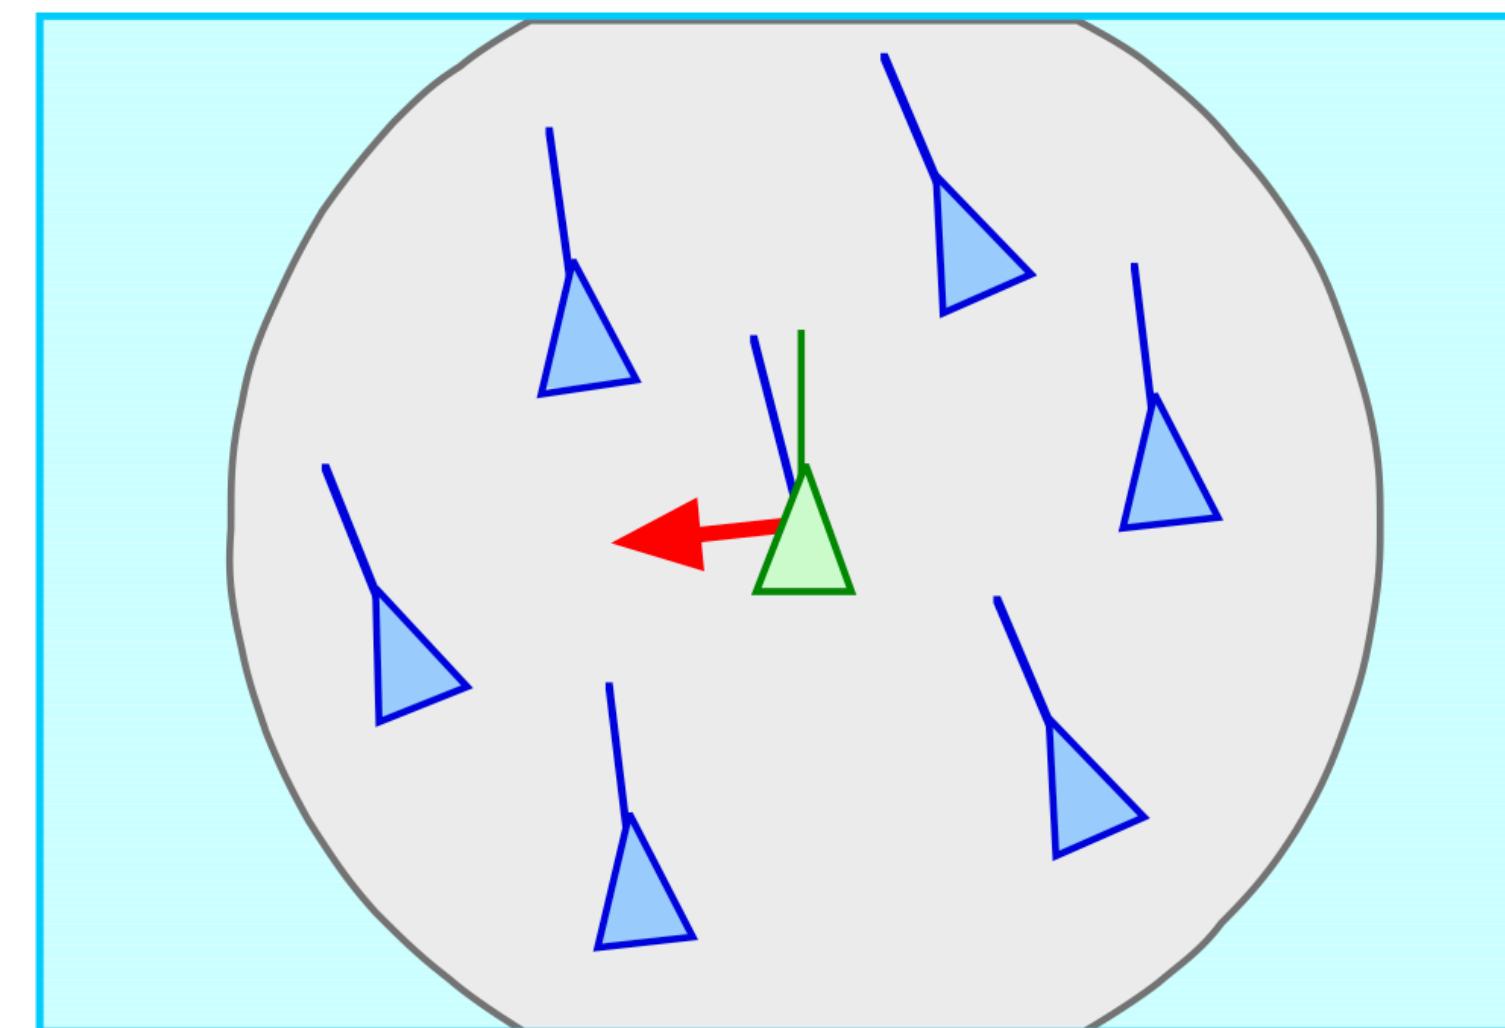


Three components of flocking

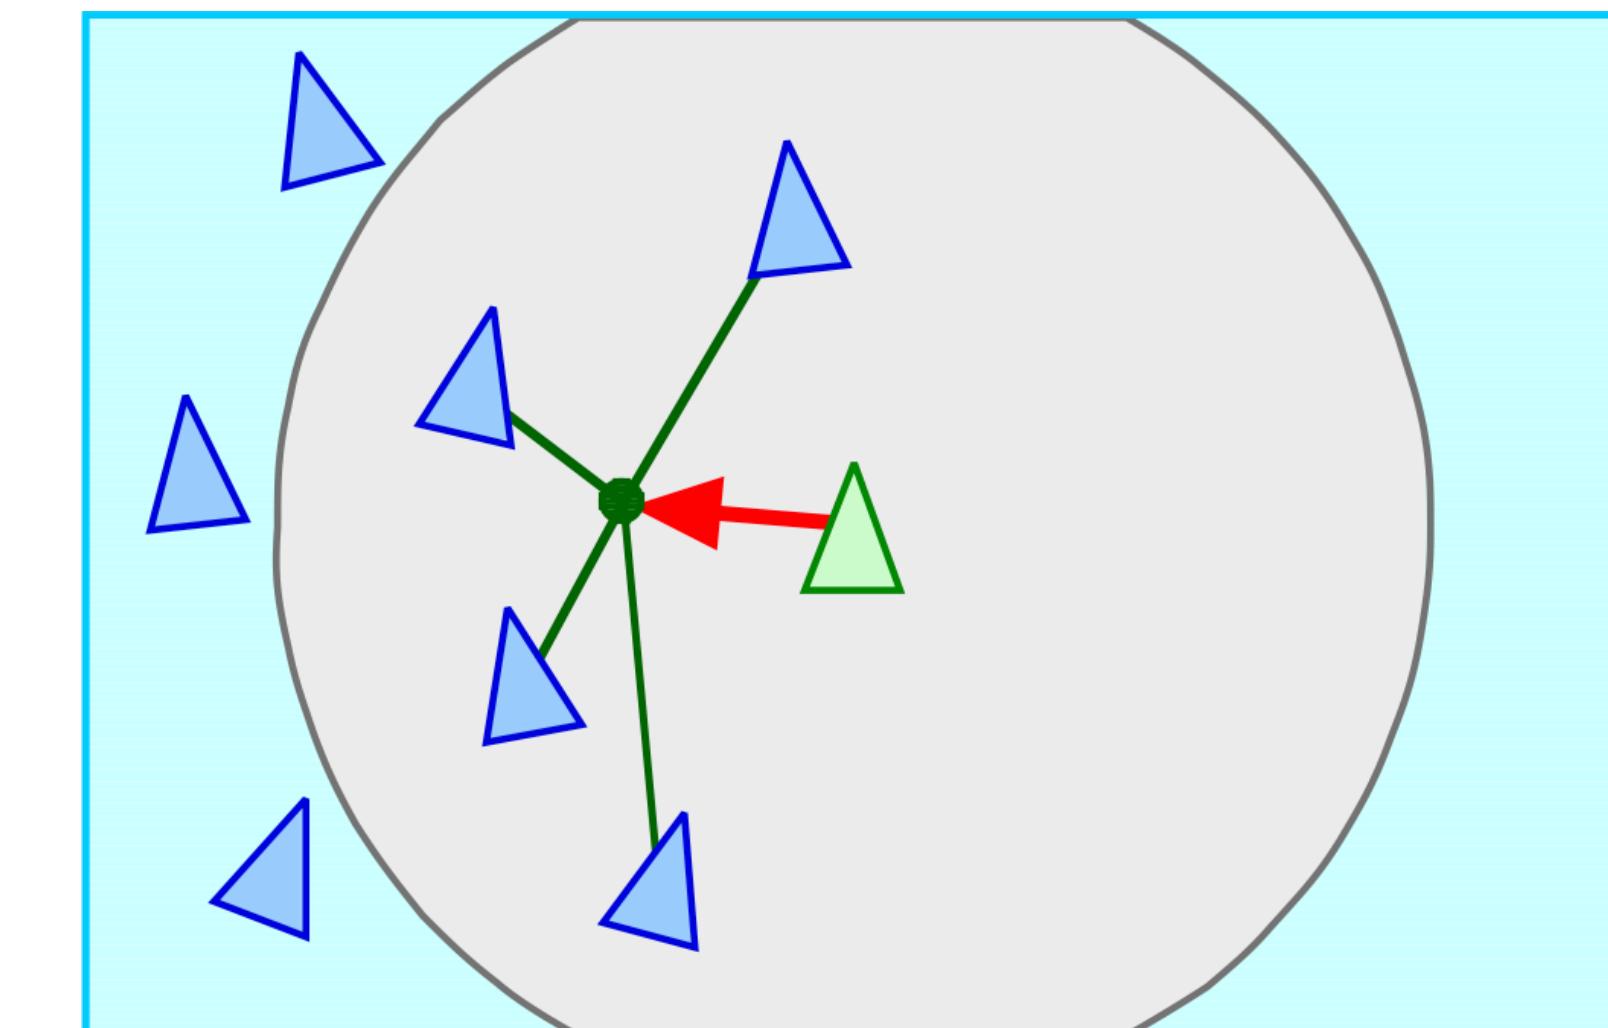
Seemed necessary, but were they sufficient?



Separation

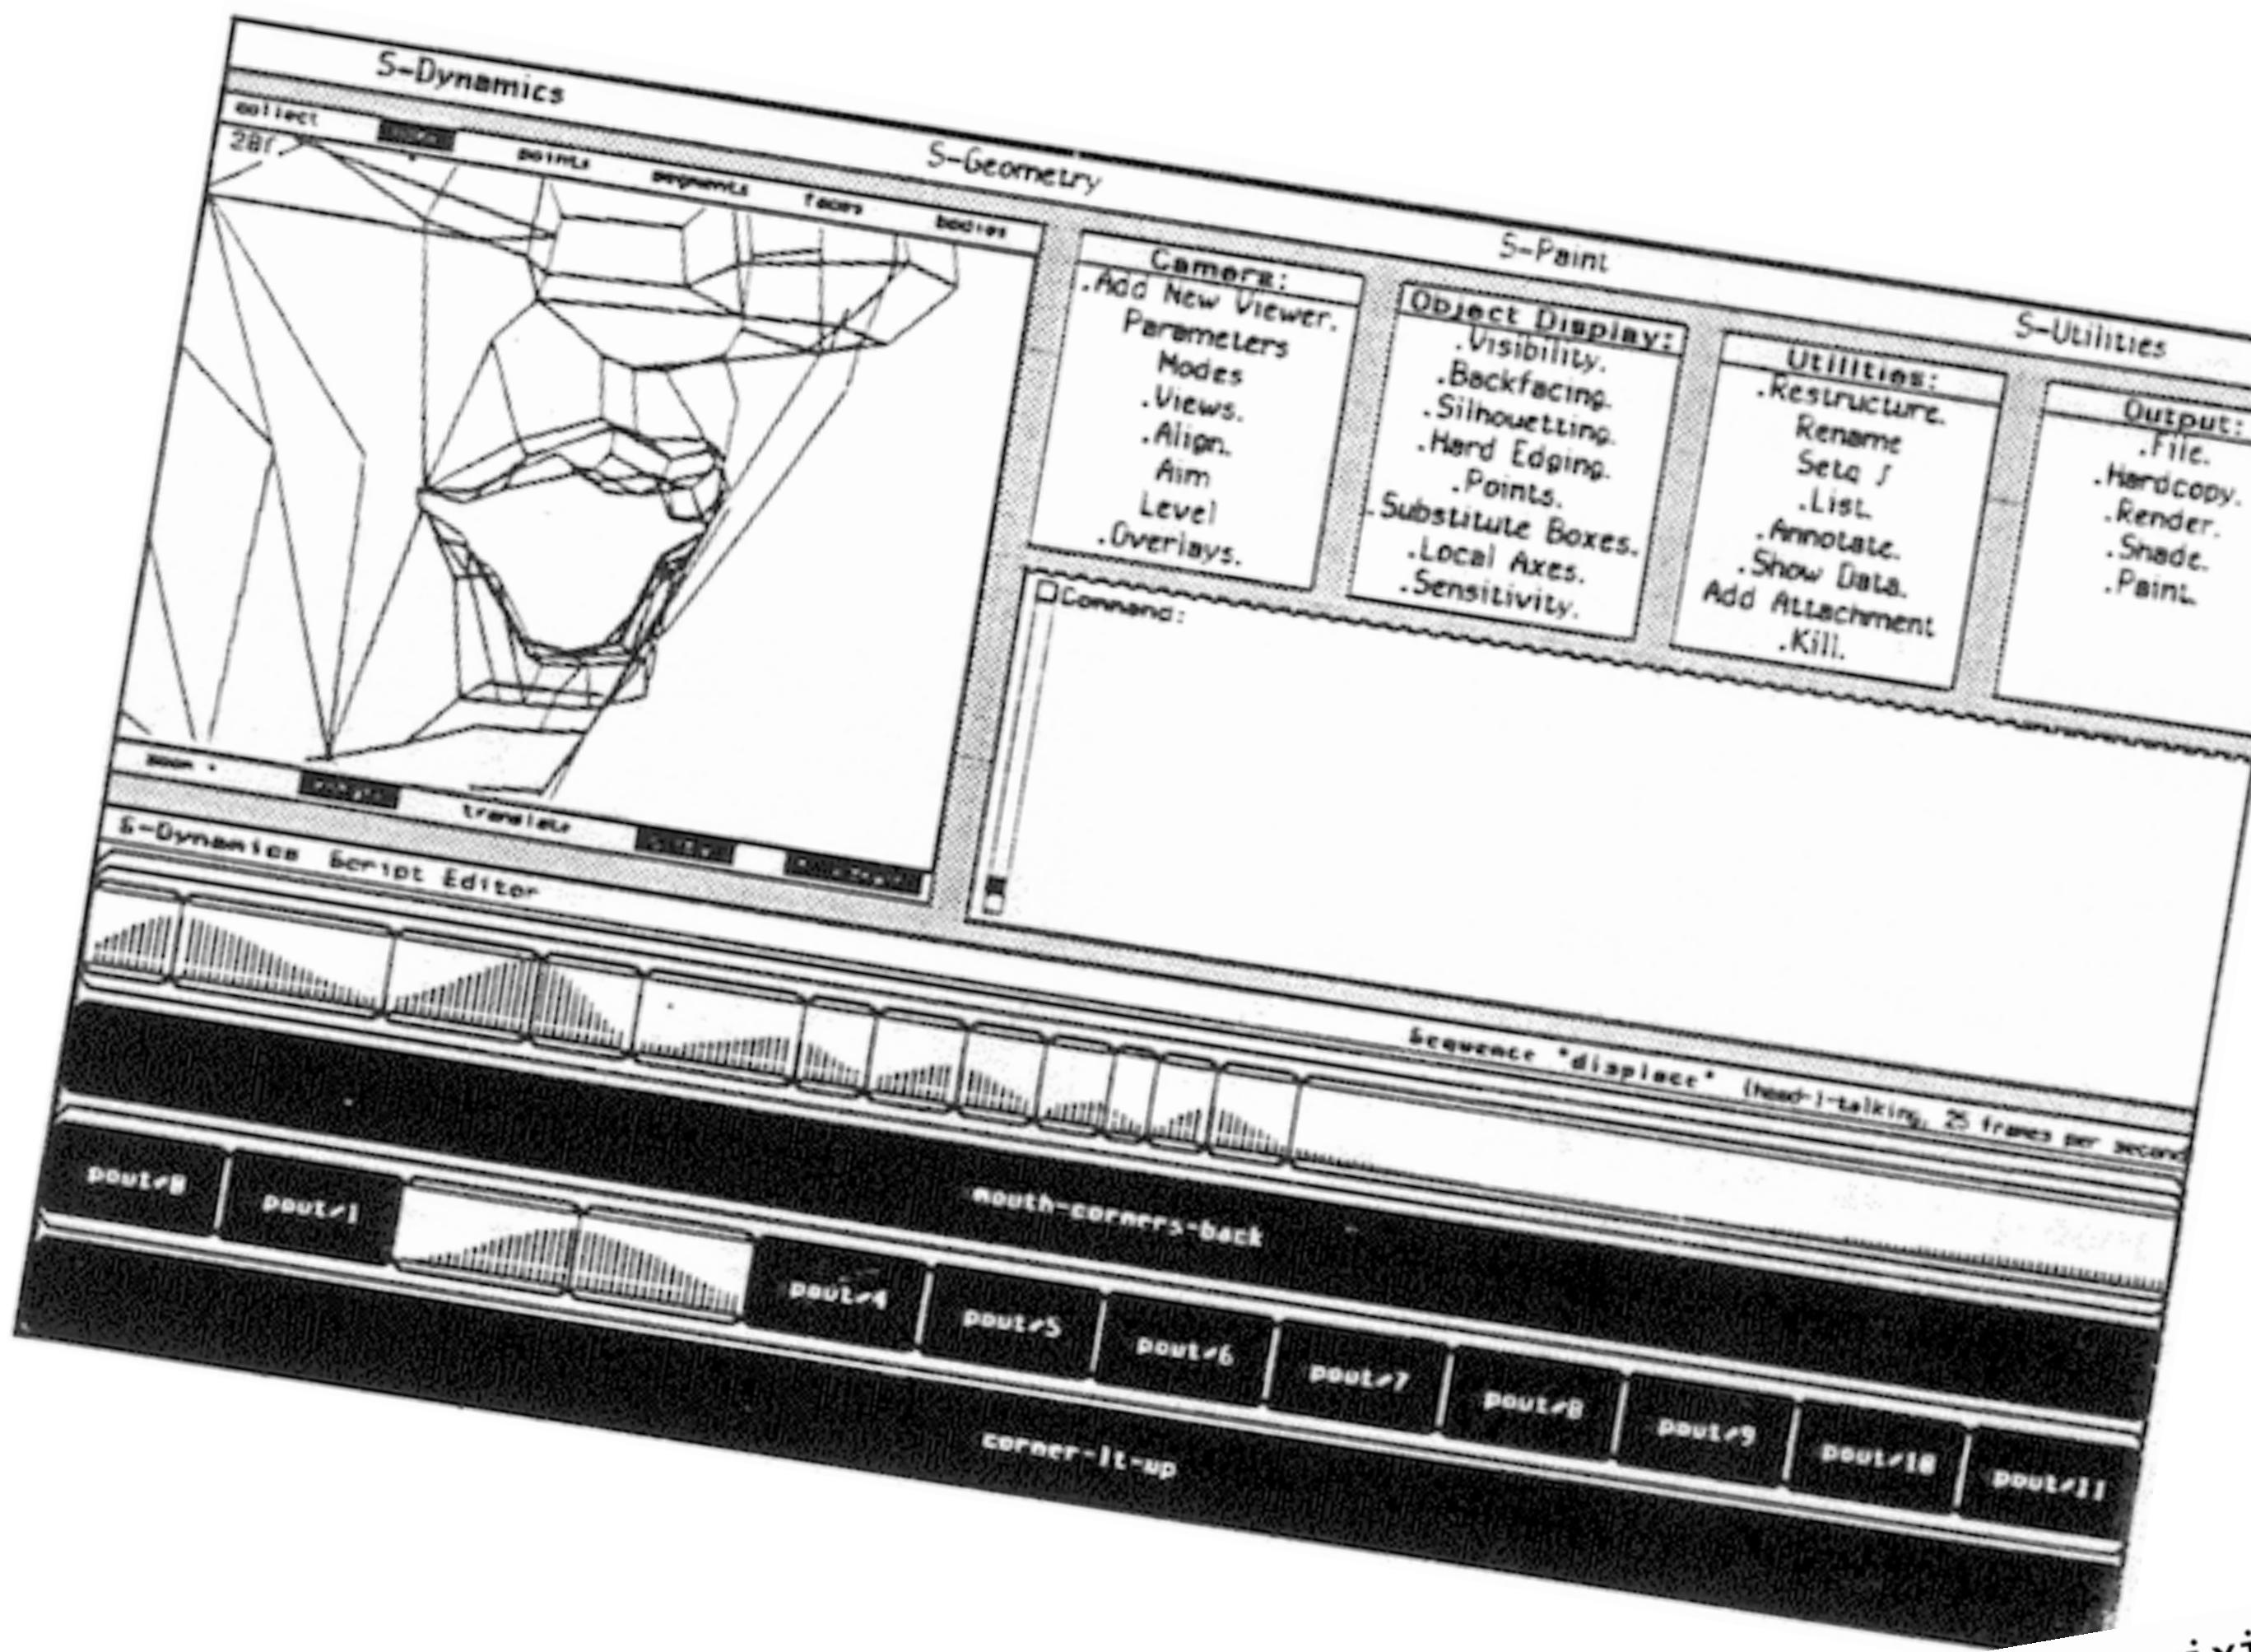


Alignment



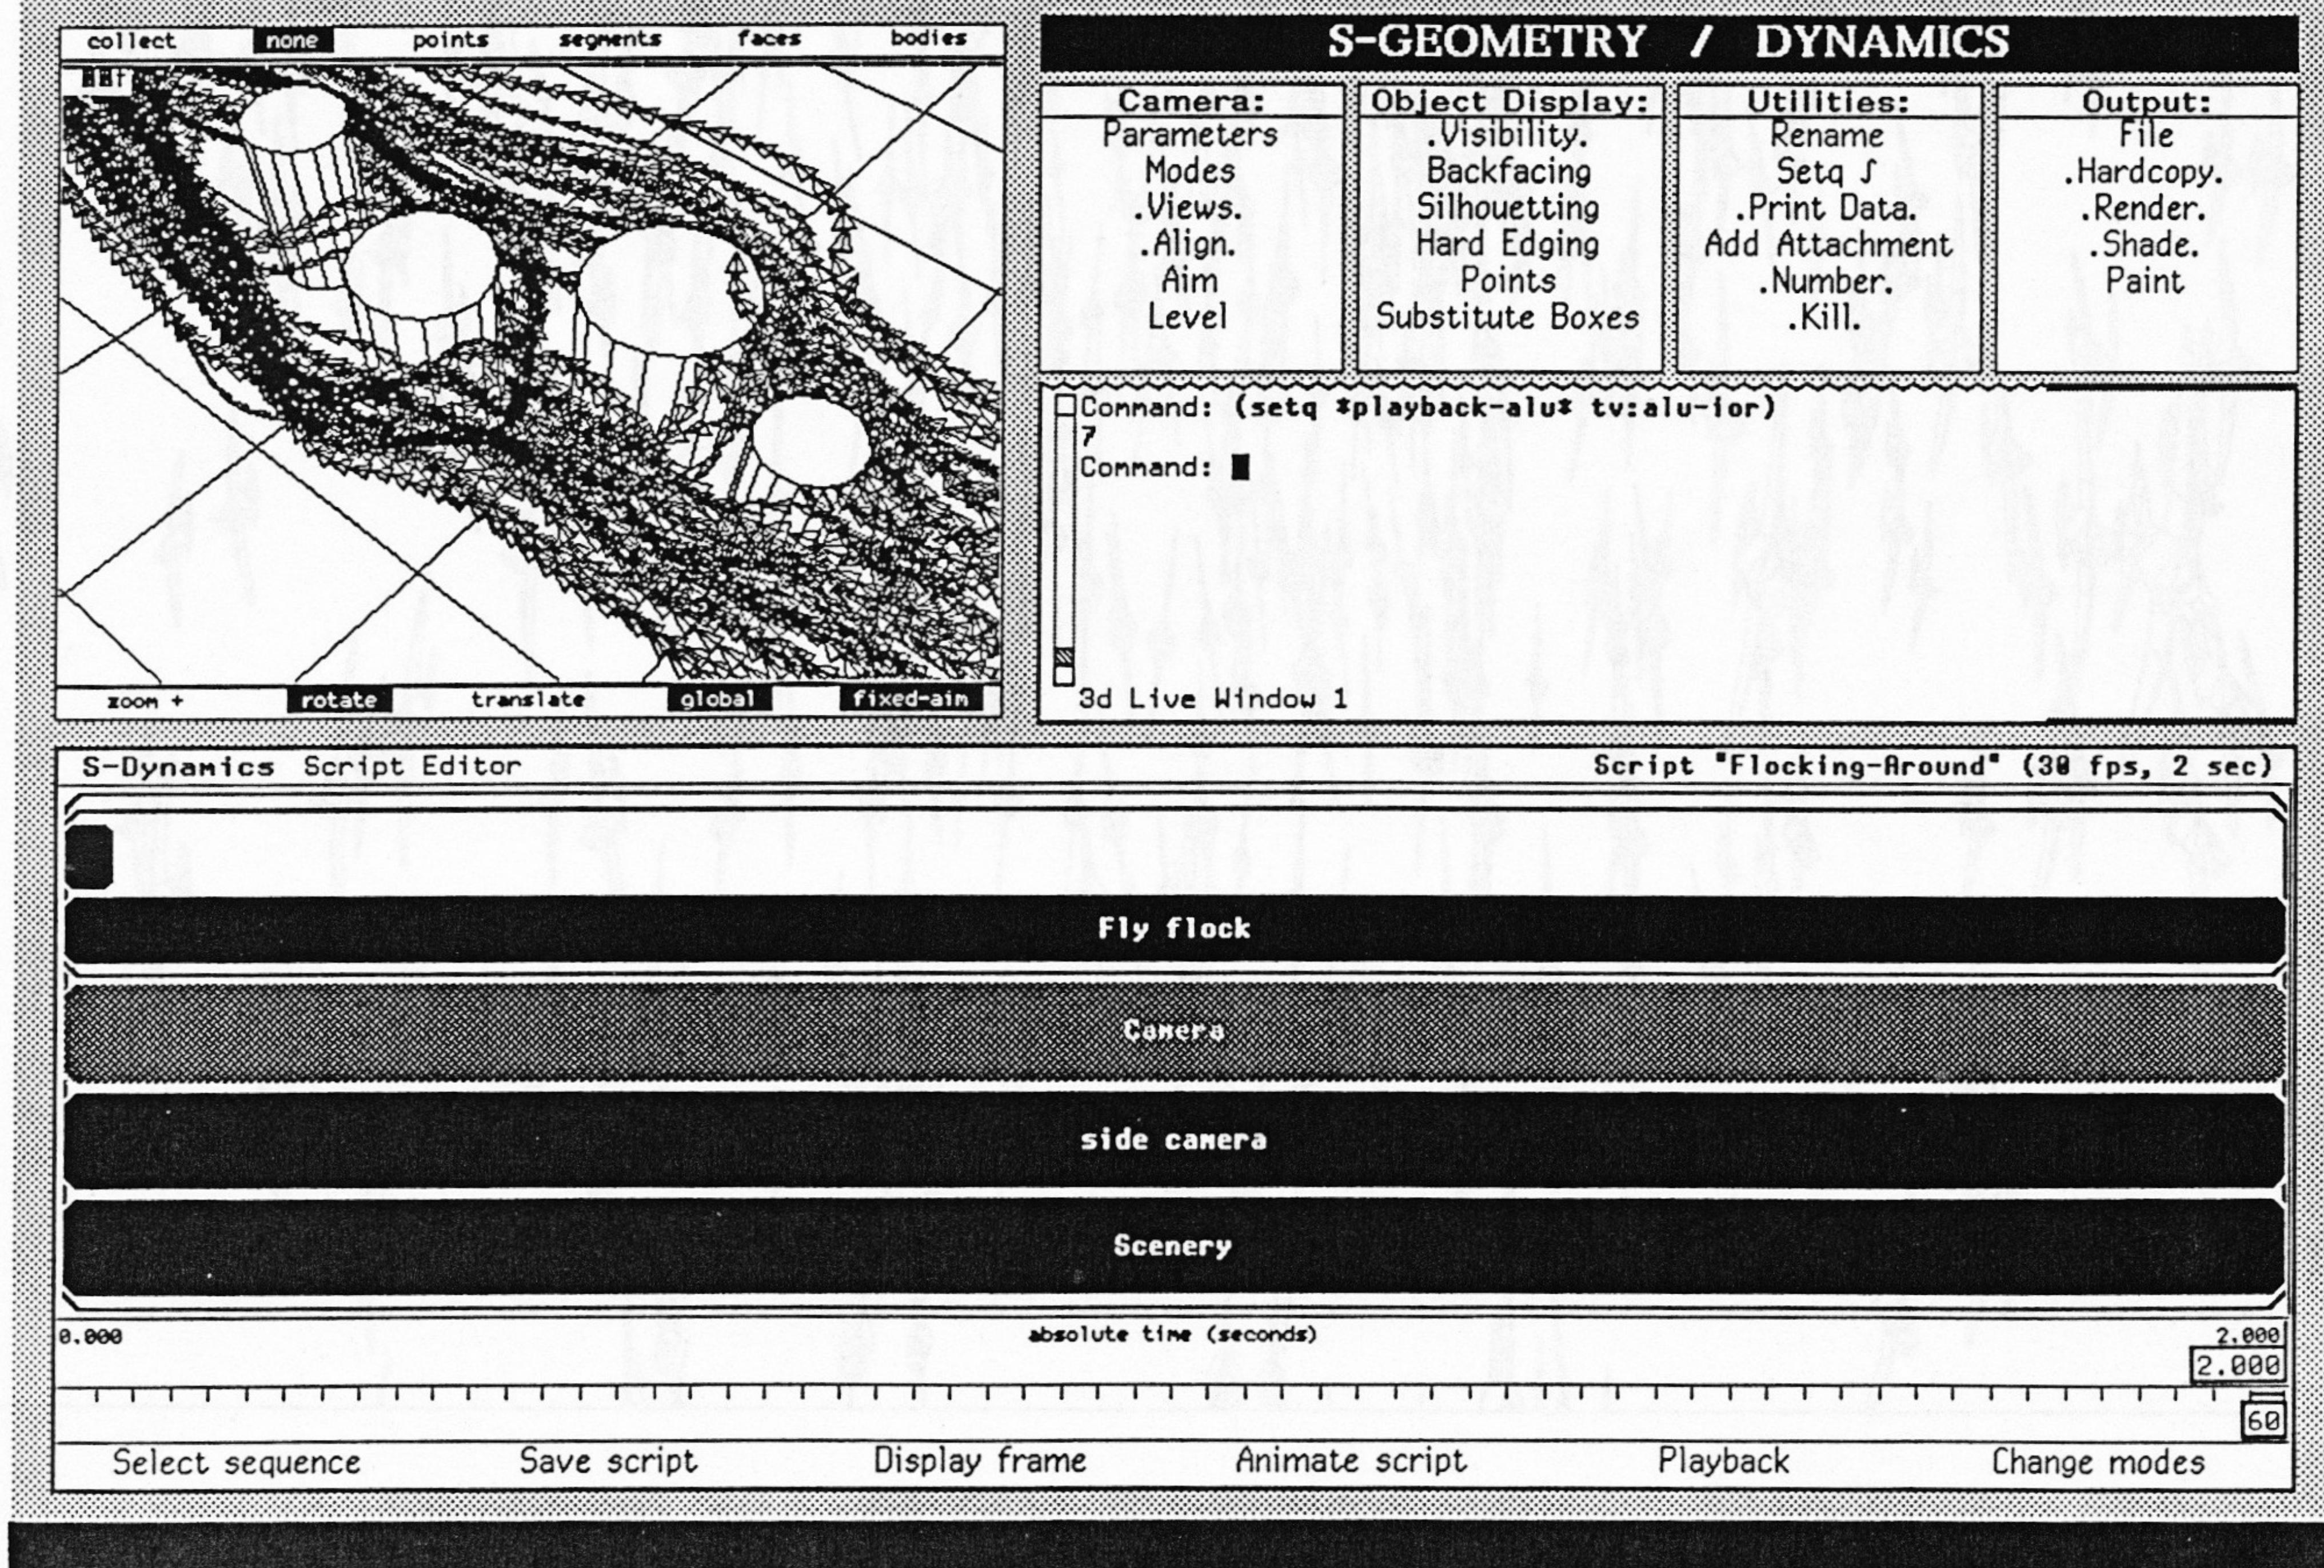
Cohesion

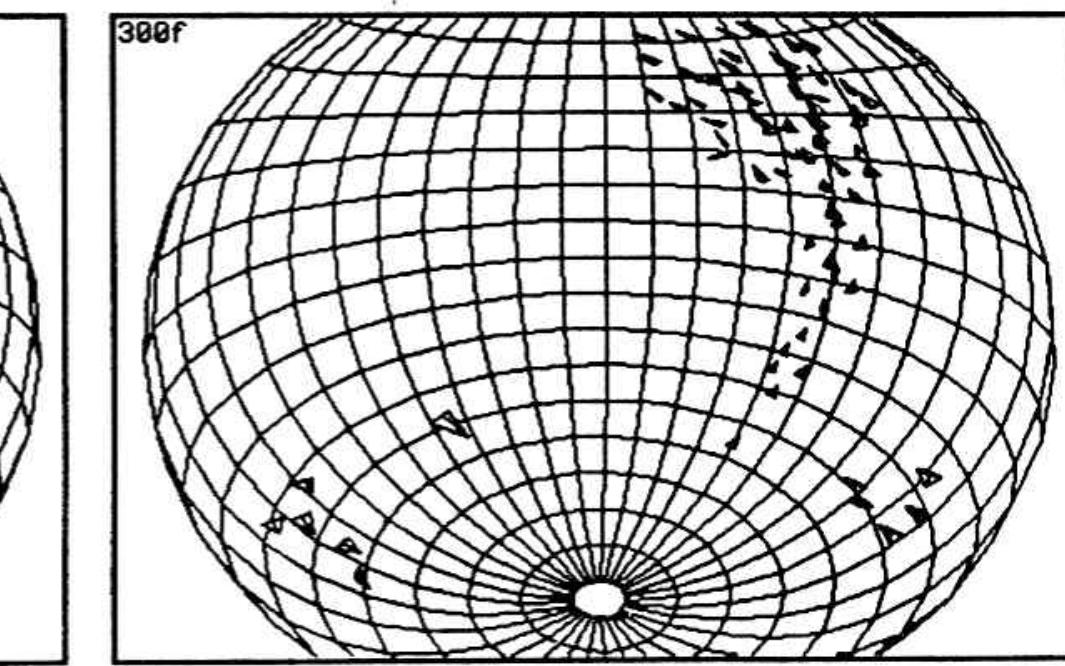
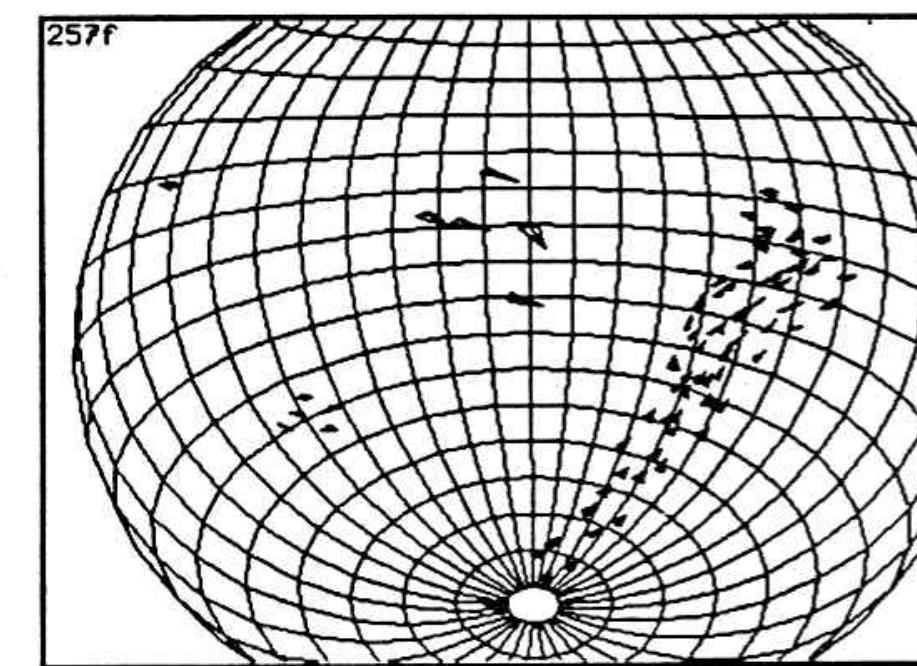
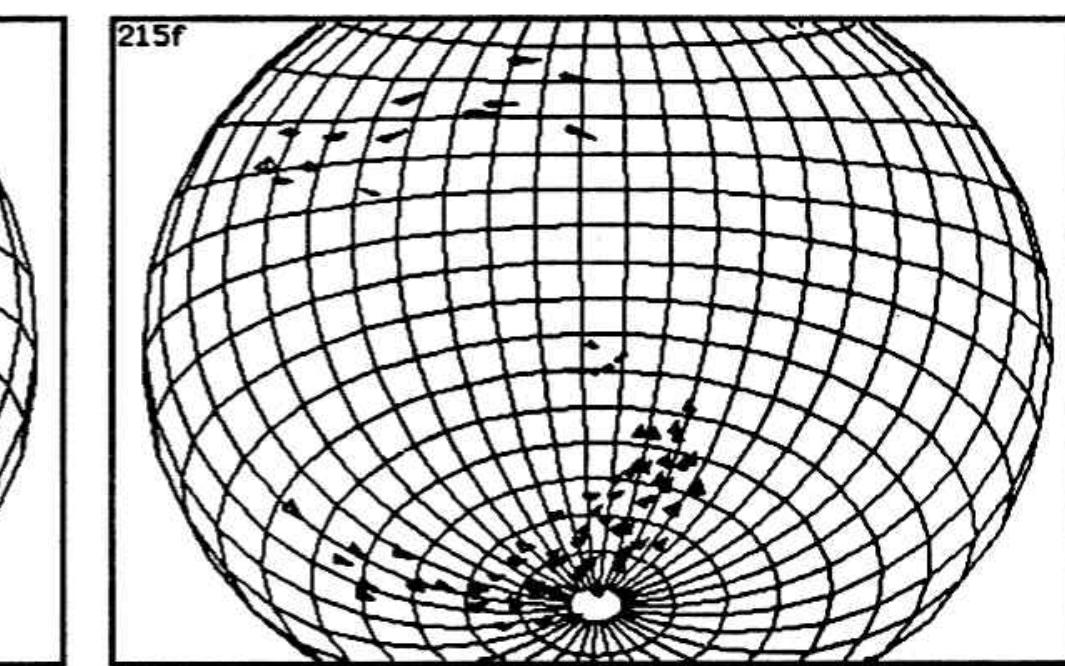
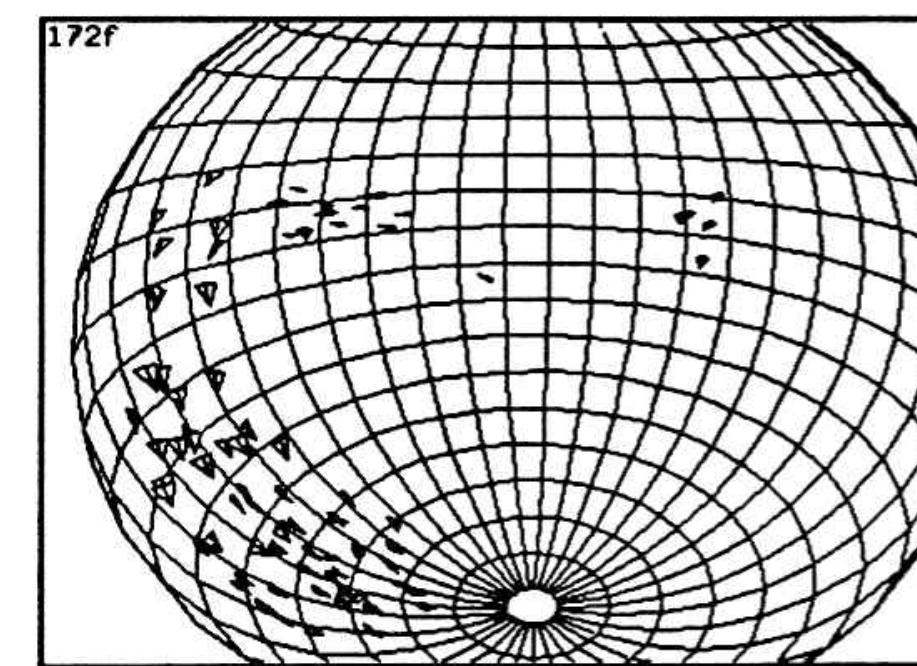
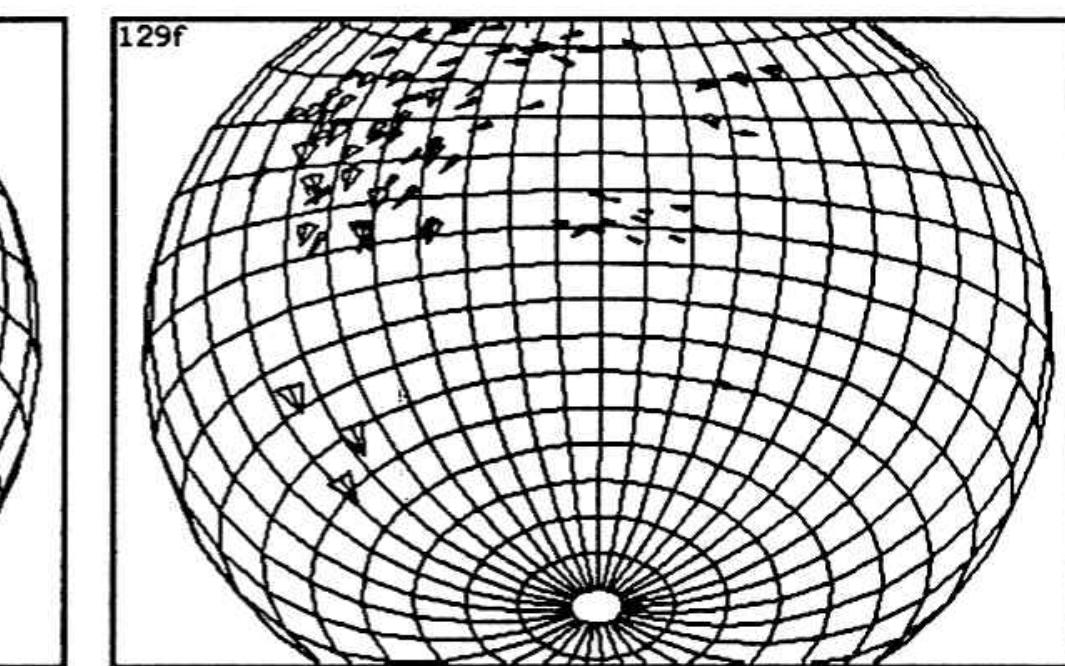
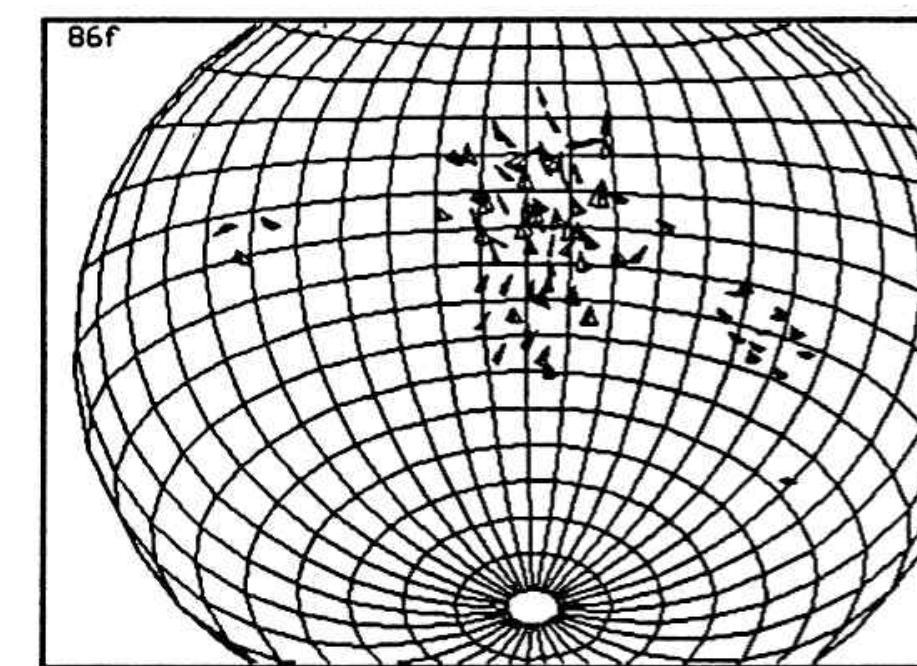
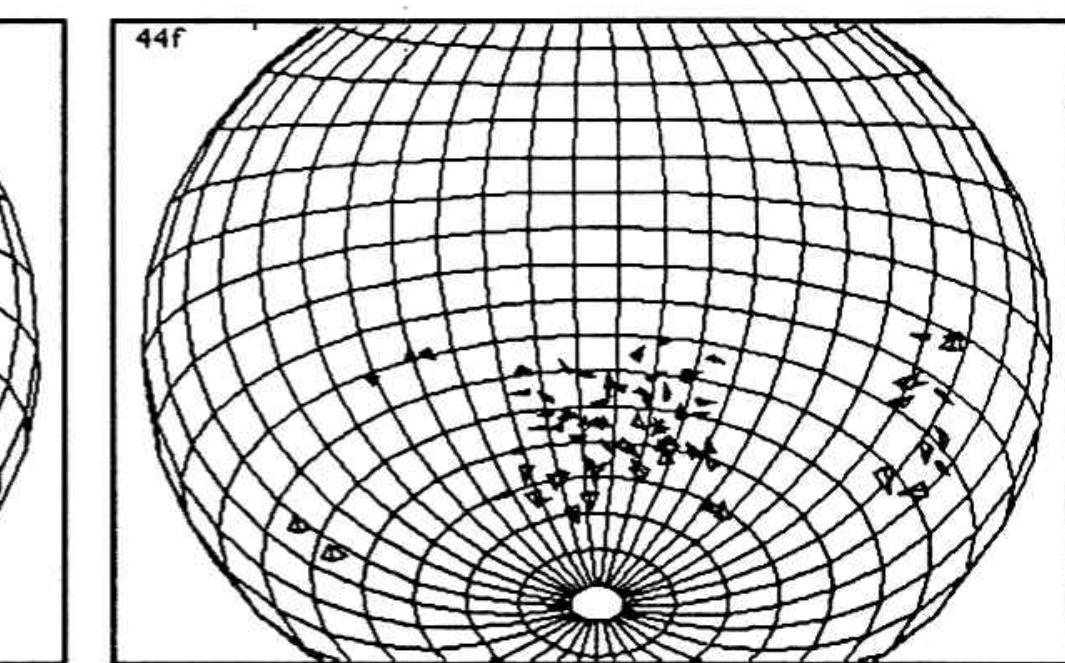
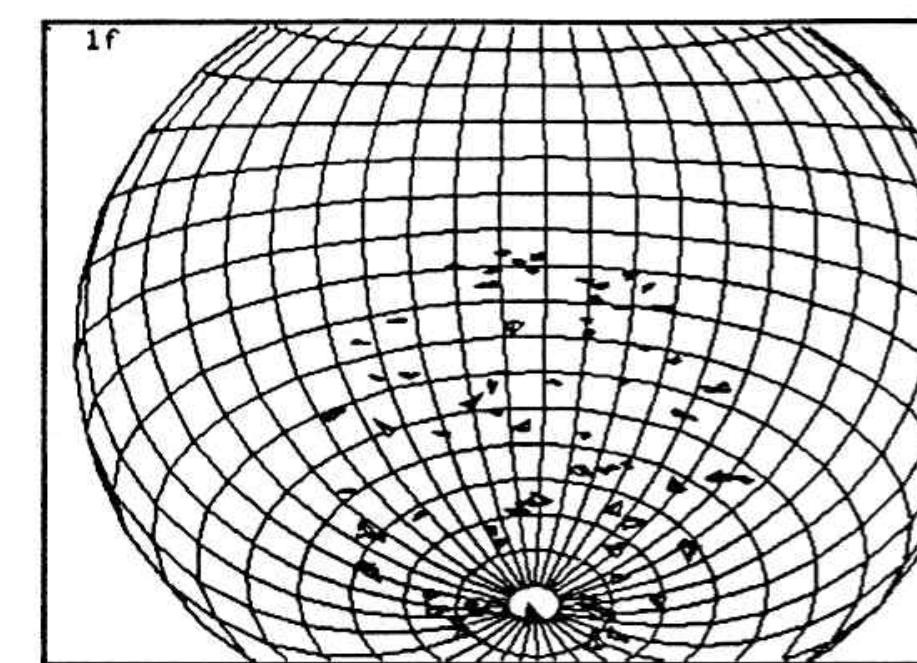
SIGGRAPH 1986:
writing tutorial notes:
“...it would be easy...”

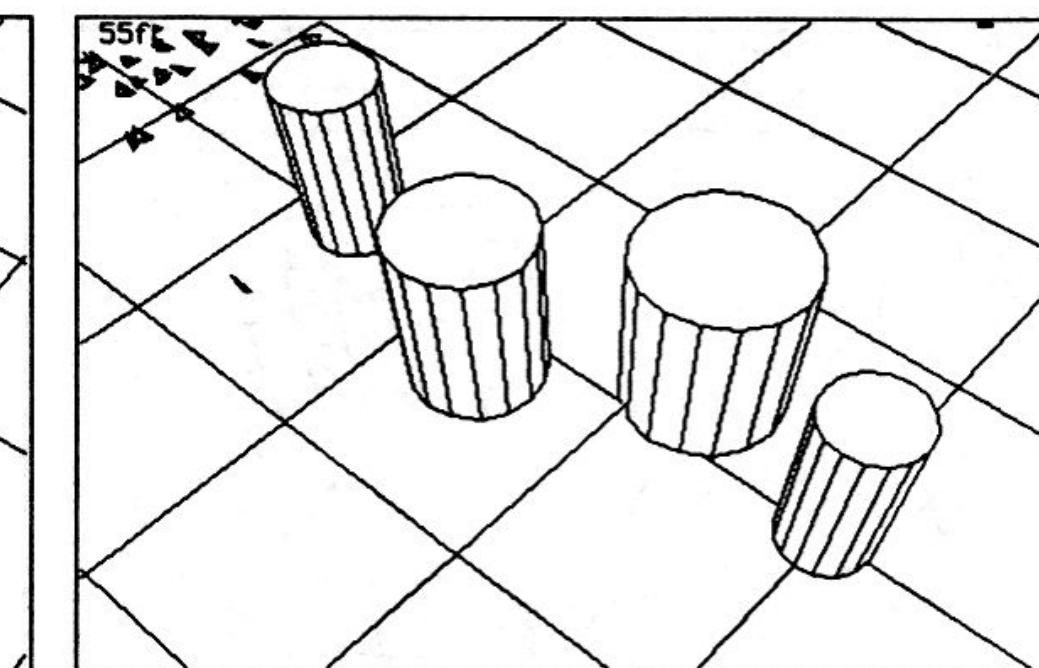
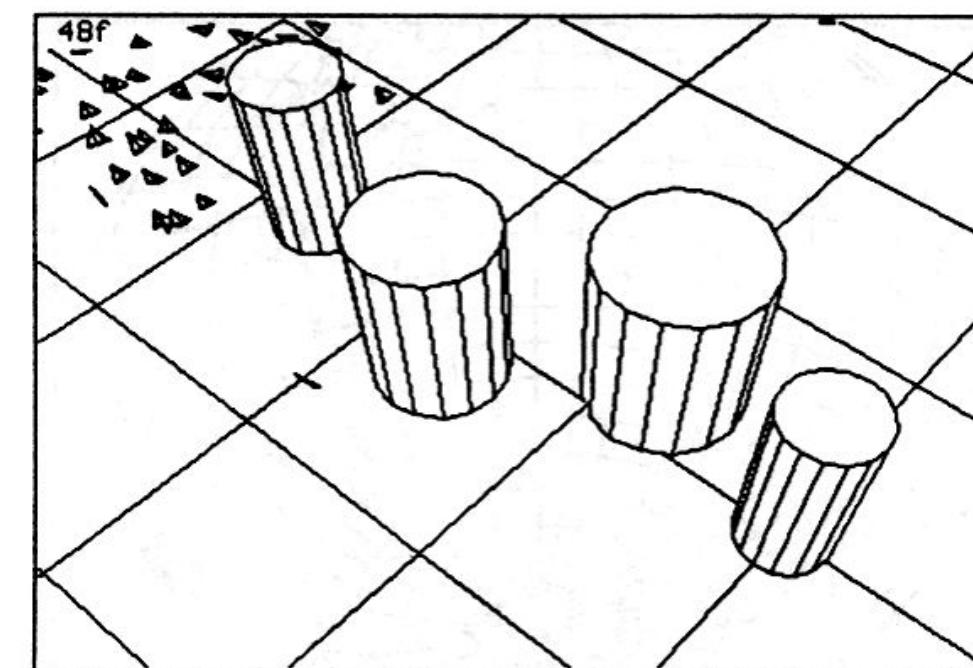
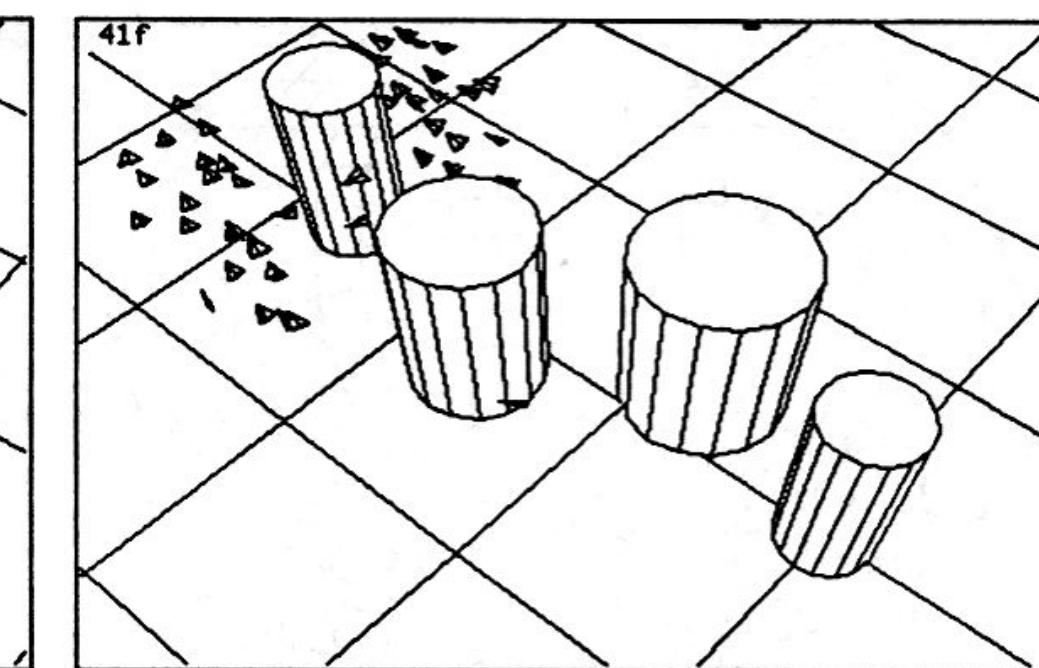
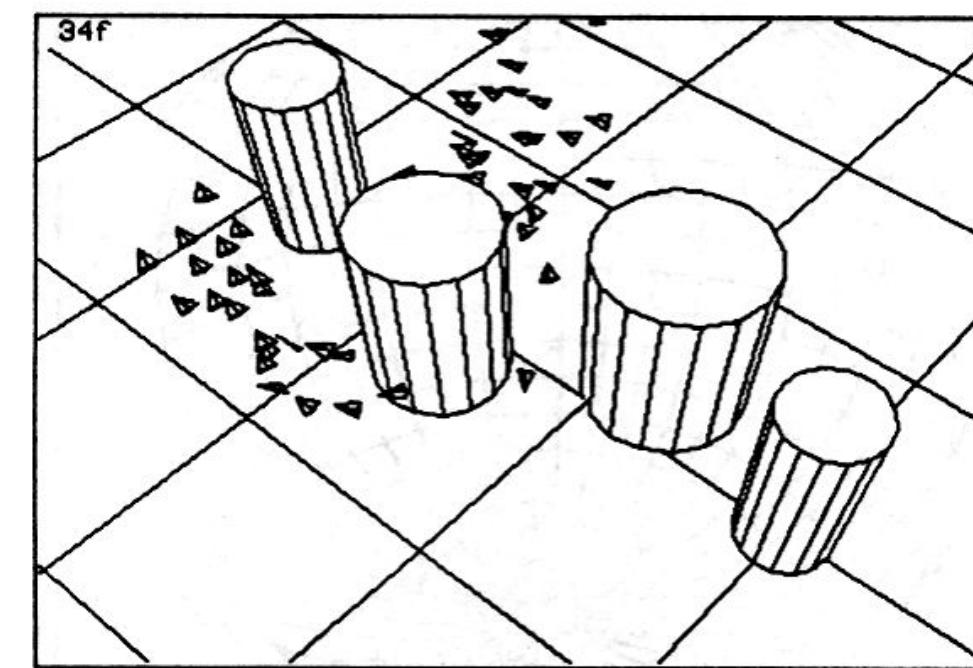
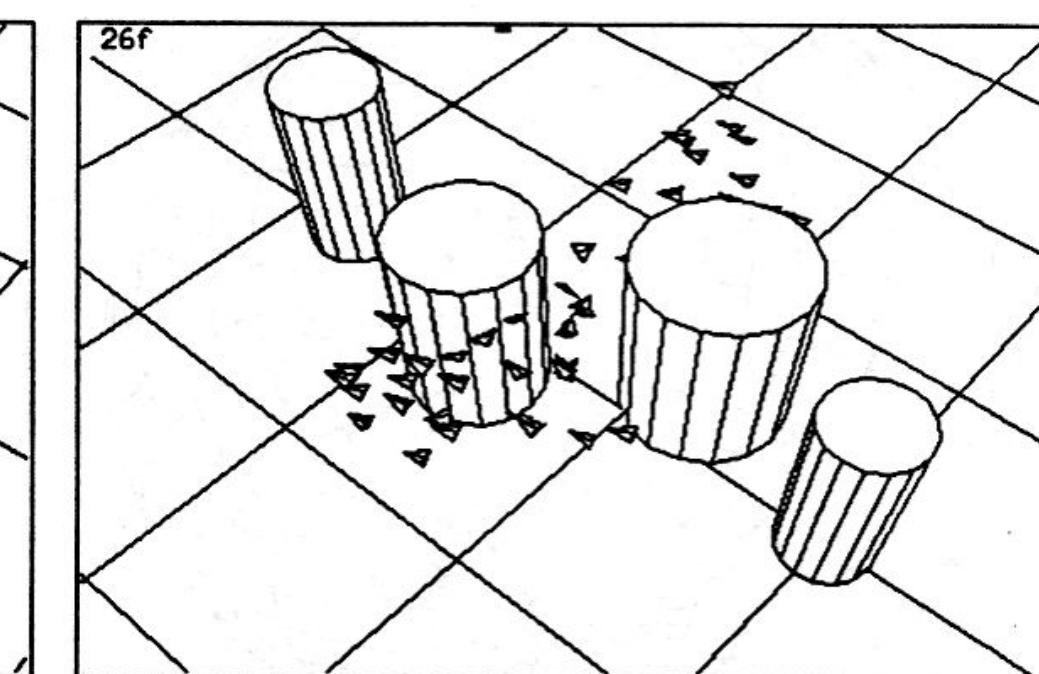
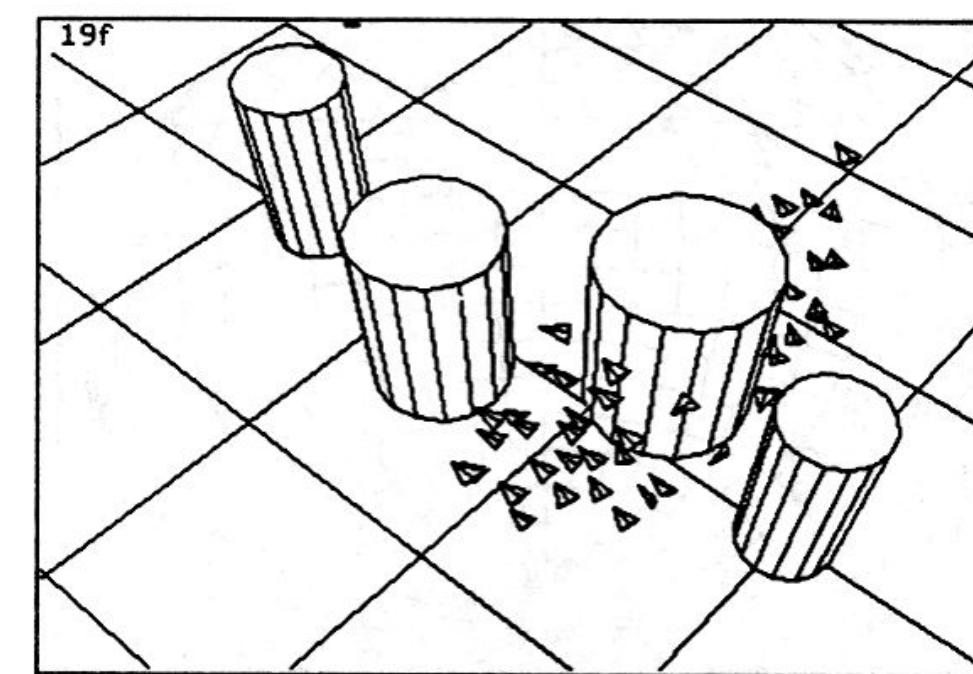
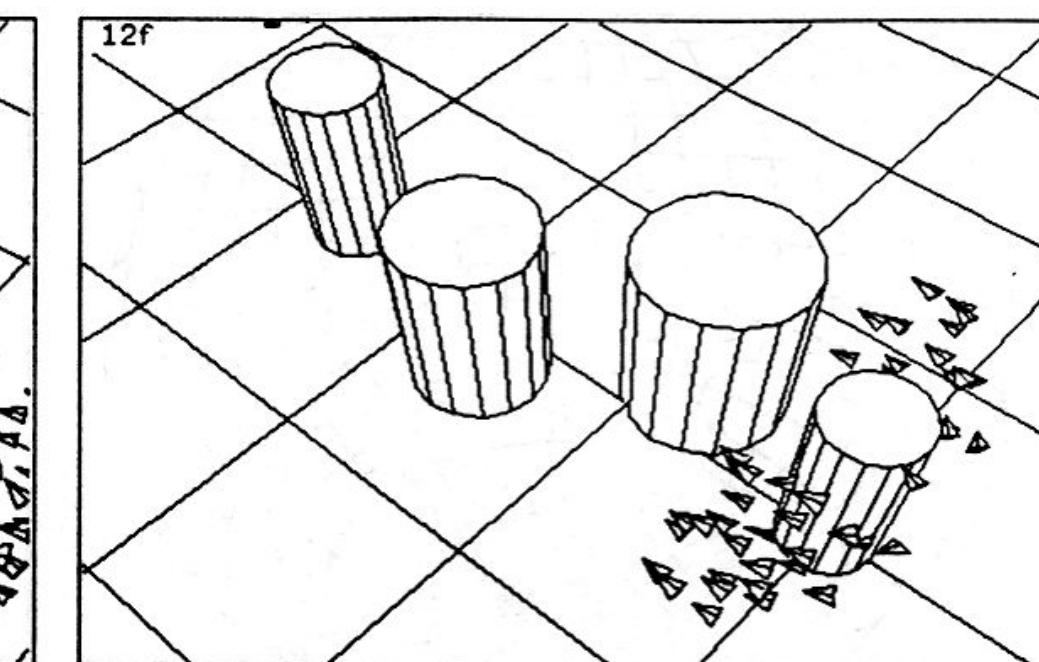
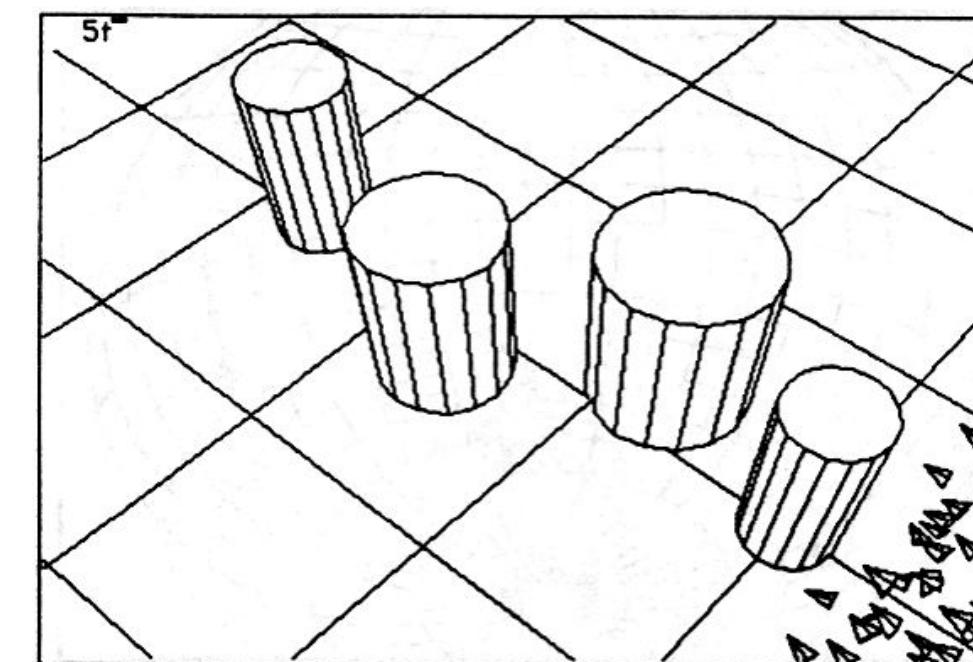


```

(defmethod (:steer flight-mixin) (old-velocity new-velocity)
  (vlet* ((new-adjusted-velocity (send self :constrain-steering new-velocity))
          (new-adj-acceleration (3d-vector-sub new-adjusted-velocity old-velocity))
          (actual-steering-acceleration (truncate-magnitude new-adj-acceleration
                                                               (send self :max-acceleration)))
          (actual-new-velocity (3d-vector-add old-velocity actual-new-velocity))
          (when (plusp (magnitude-squared actual-new-velocity))
                ;; PITCH and YAW the object (local X and Y rotations)
                (send self :align-to-local actual-new-velocity)
                ;; Adjust the roll for a "coordinated turn".
                (send self :bank actual-steering-acceleration)))
    (defmethod (:steer flight-mixin) (old-velocity new-velocity)
      (vlet* ((new-adjusted-velocity (send self :constrain-steering new-velocity))
              (new-adj-acceleration (3d-vector-sub new-adjusted-velocity old-velocity))
              (actual-steering-acceleration (truncate-magnitude new-adj-acceleration
                                                               (send self :max-acceleration)))
              (actual-new-velocity (3d-vector-add old-velocity actual-new-velocity))
              (when (plusp (magnitude-squared actual-new-velocity))
                    ;; PITCH and YAW the object (local X and Y rotations)
                    (send self :align-to-local actual-new-velocity)
                    ;; Adjust the roll for a "coordinated turn".
                    (send self :bank actual-steering-acceleration))))
```



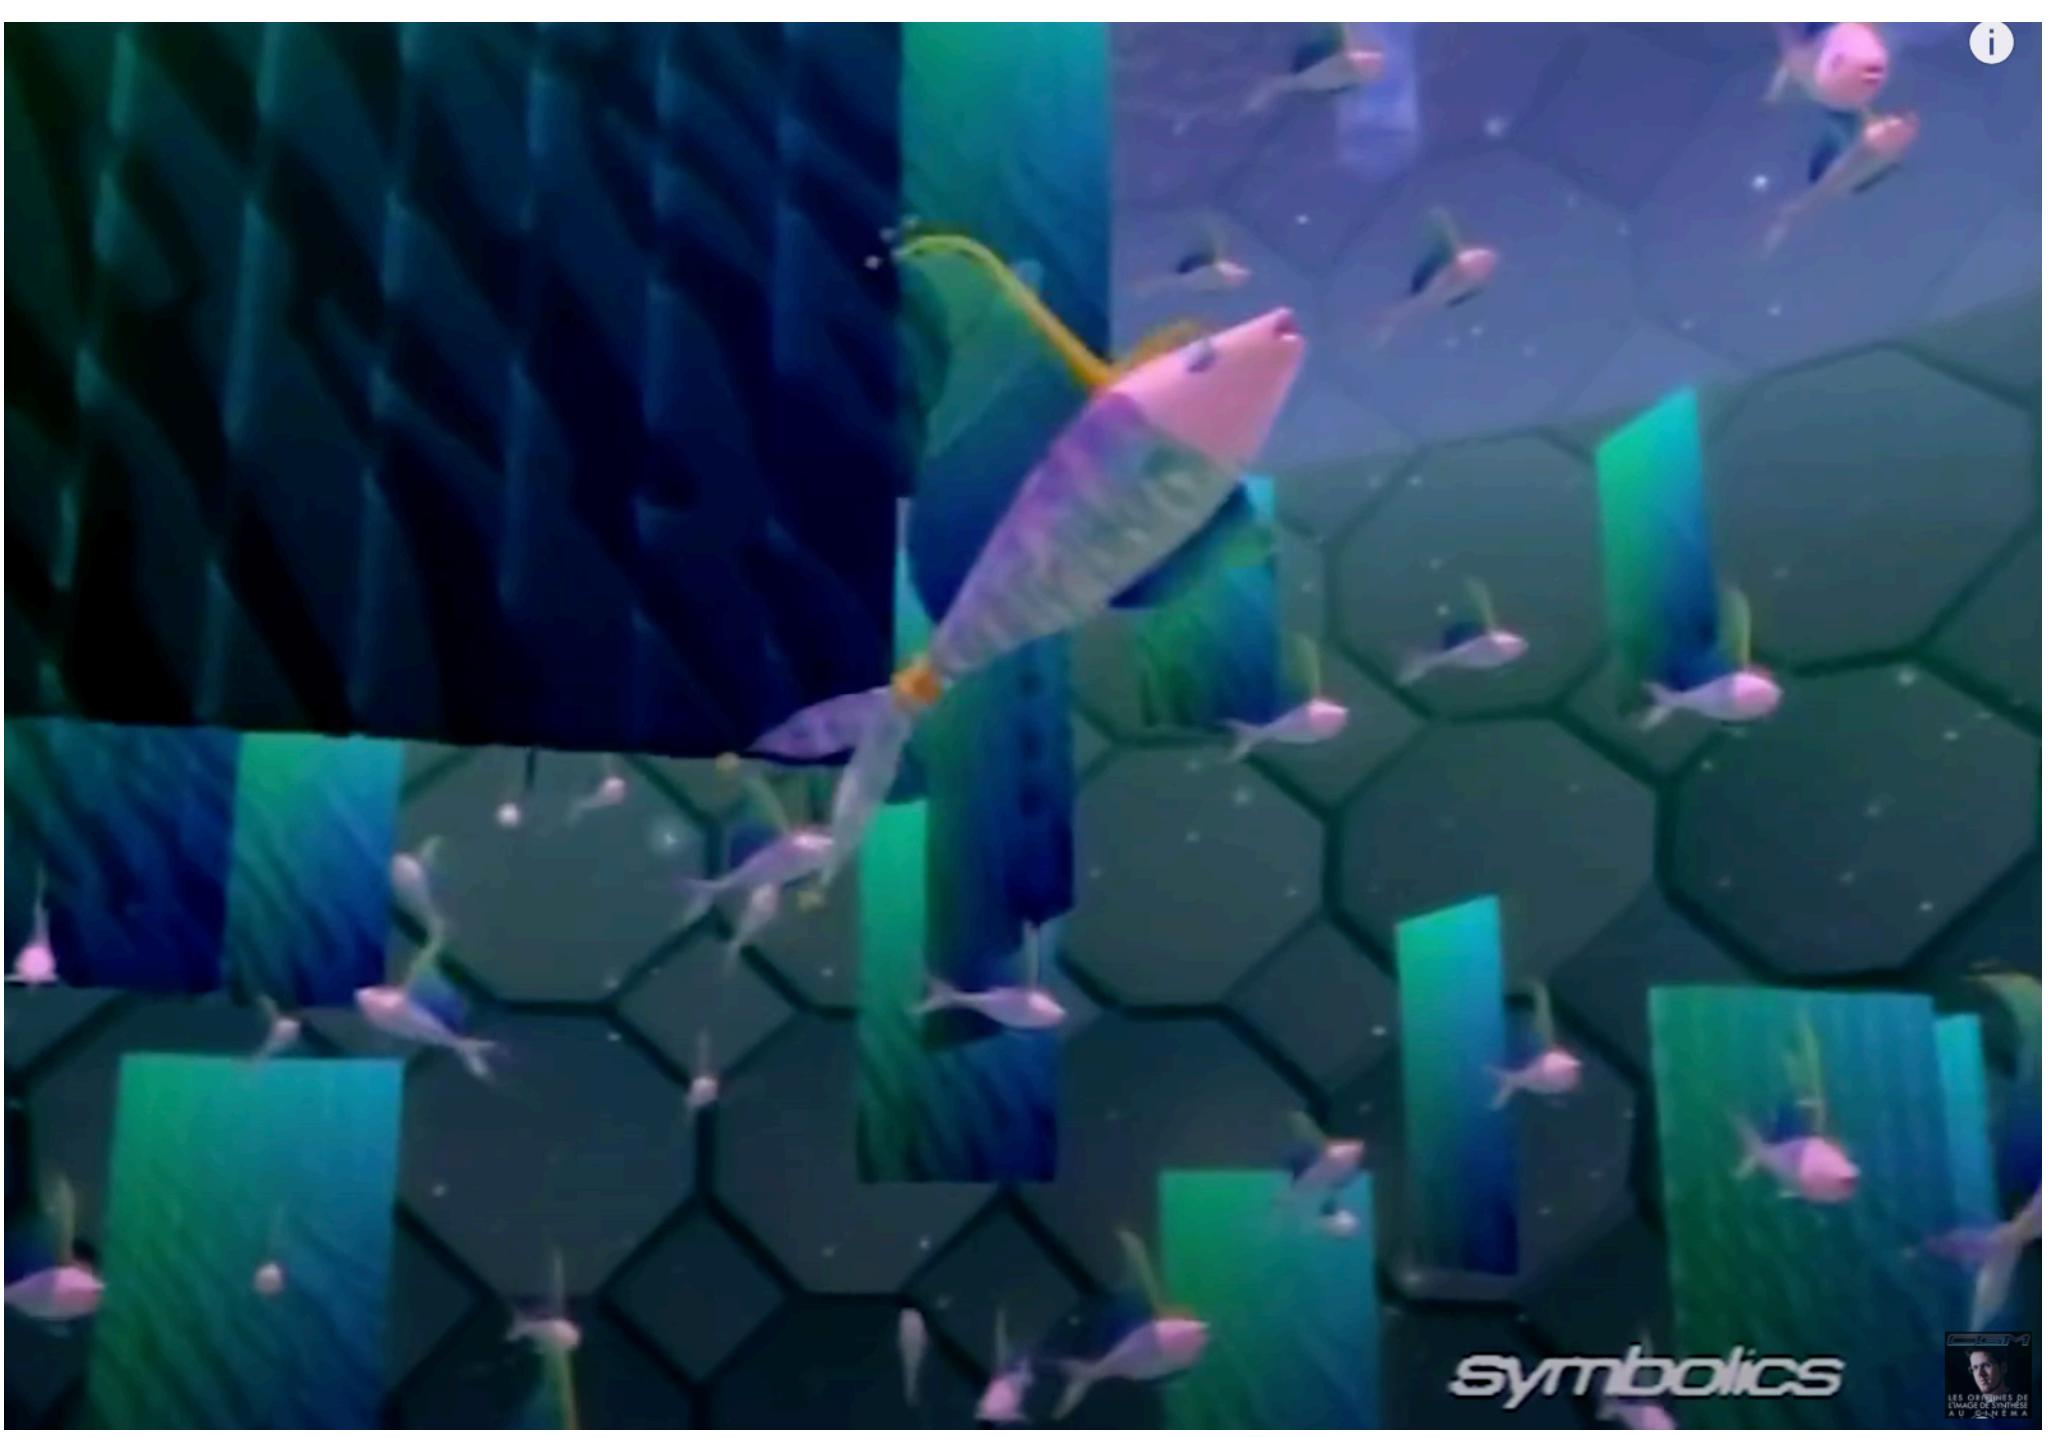
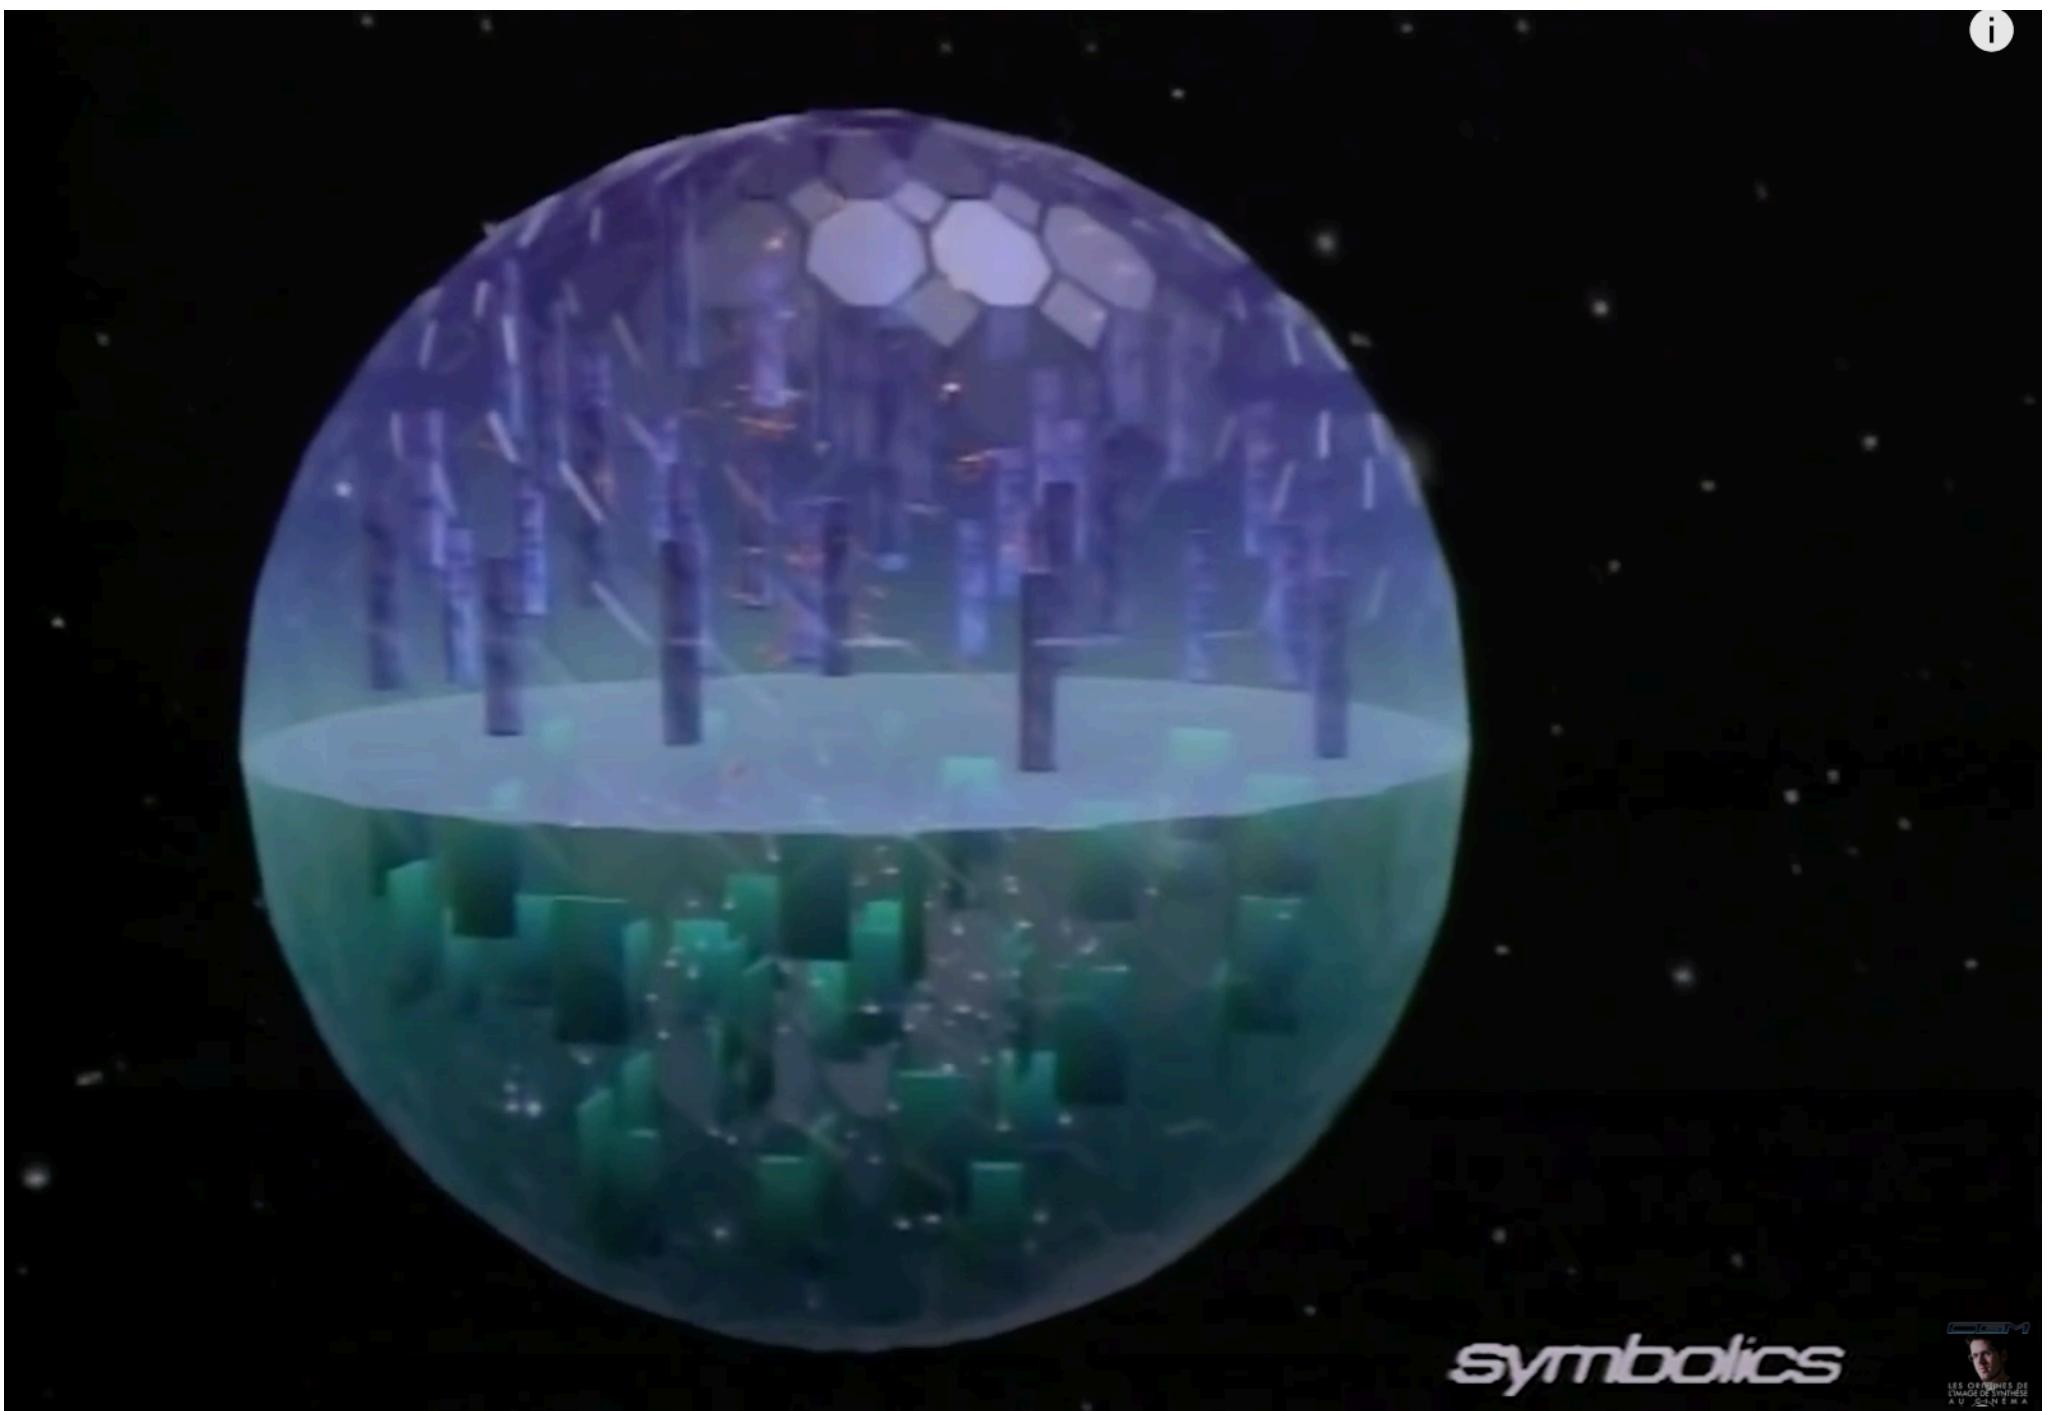


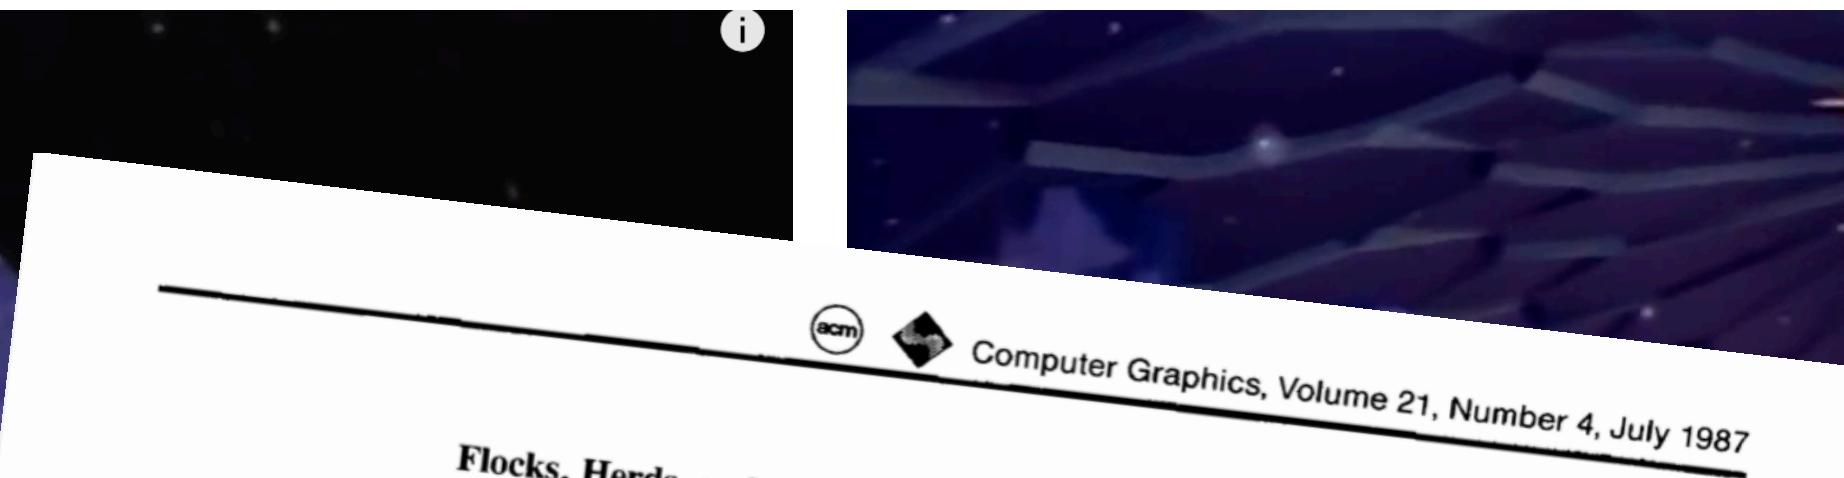
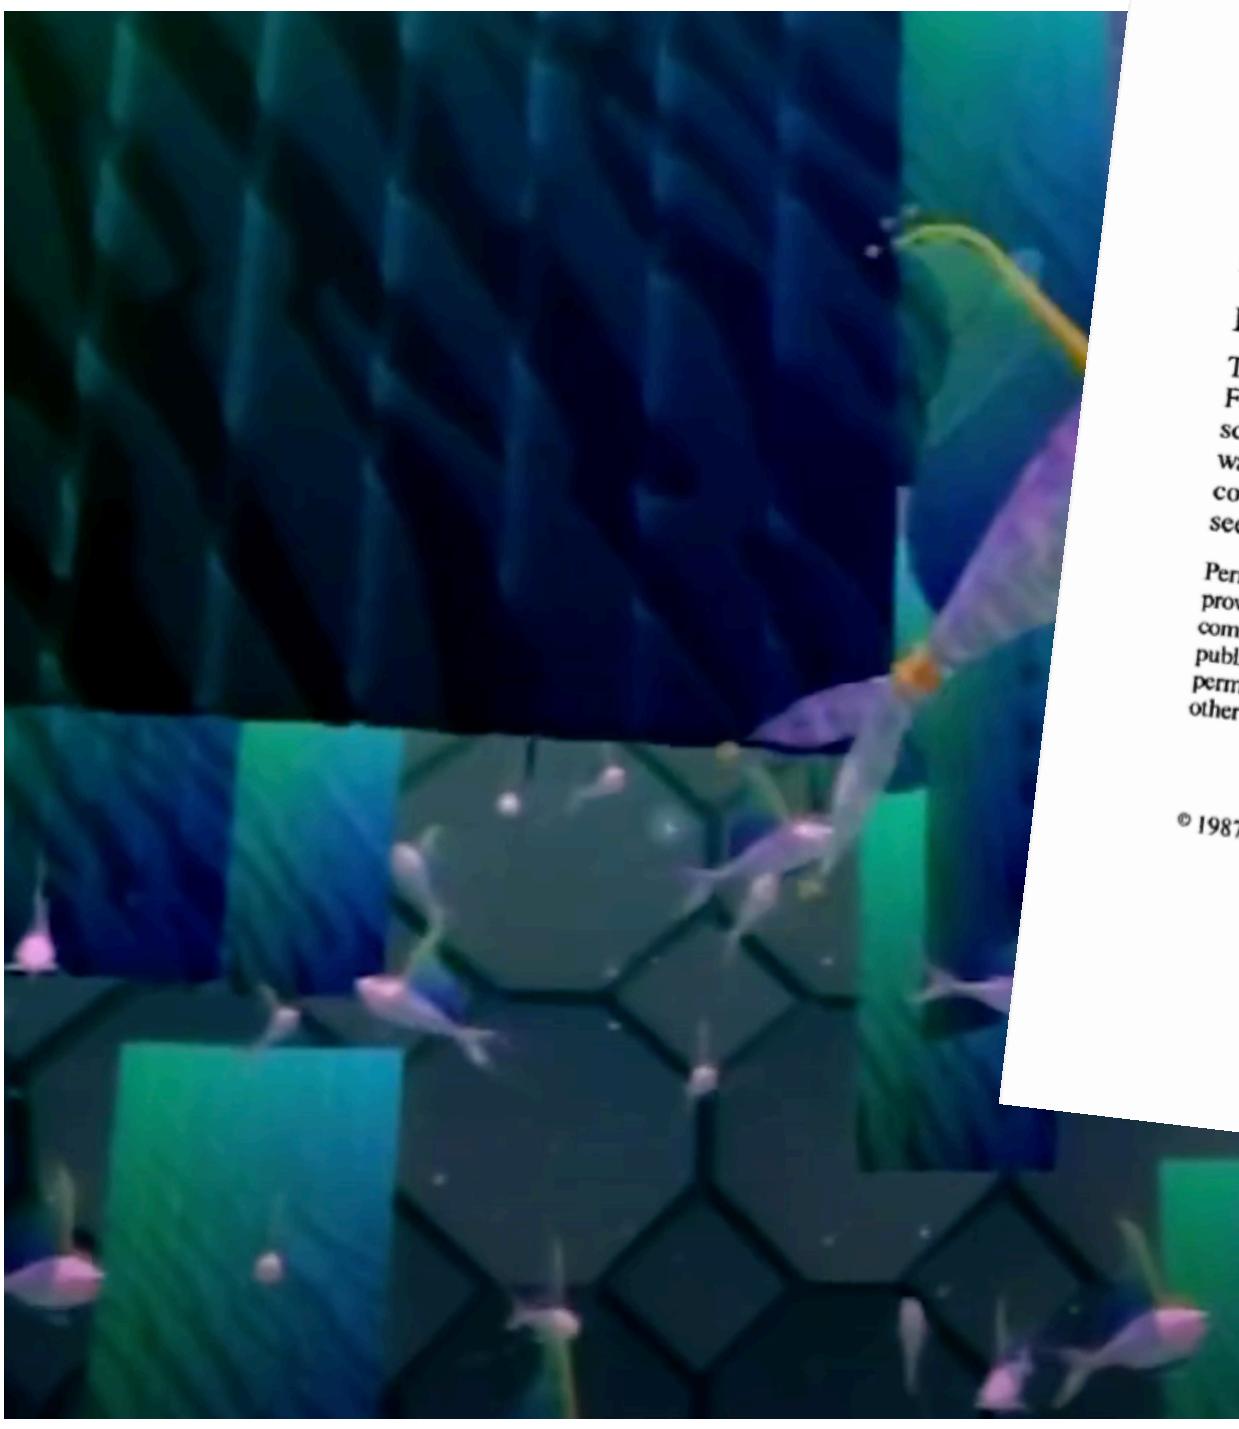
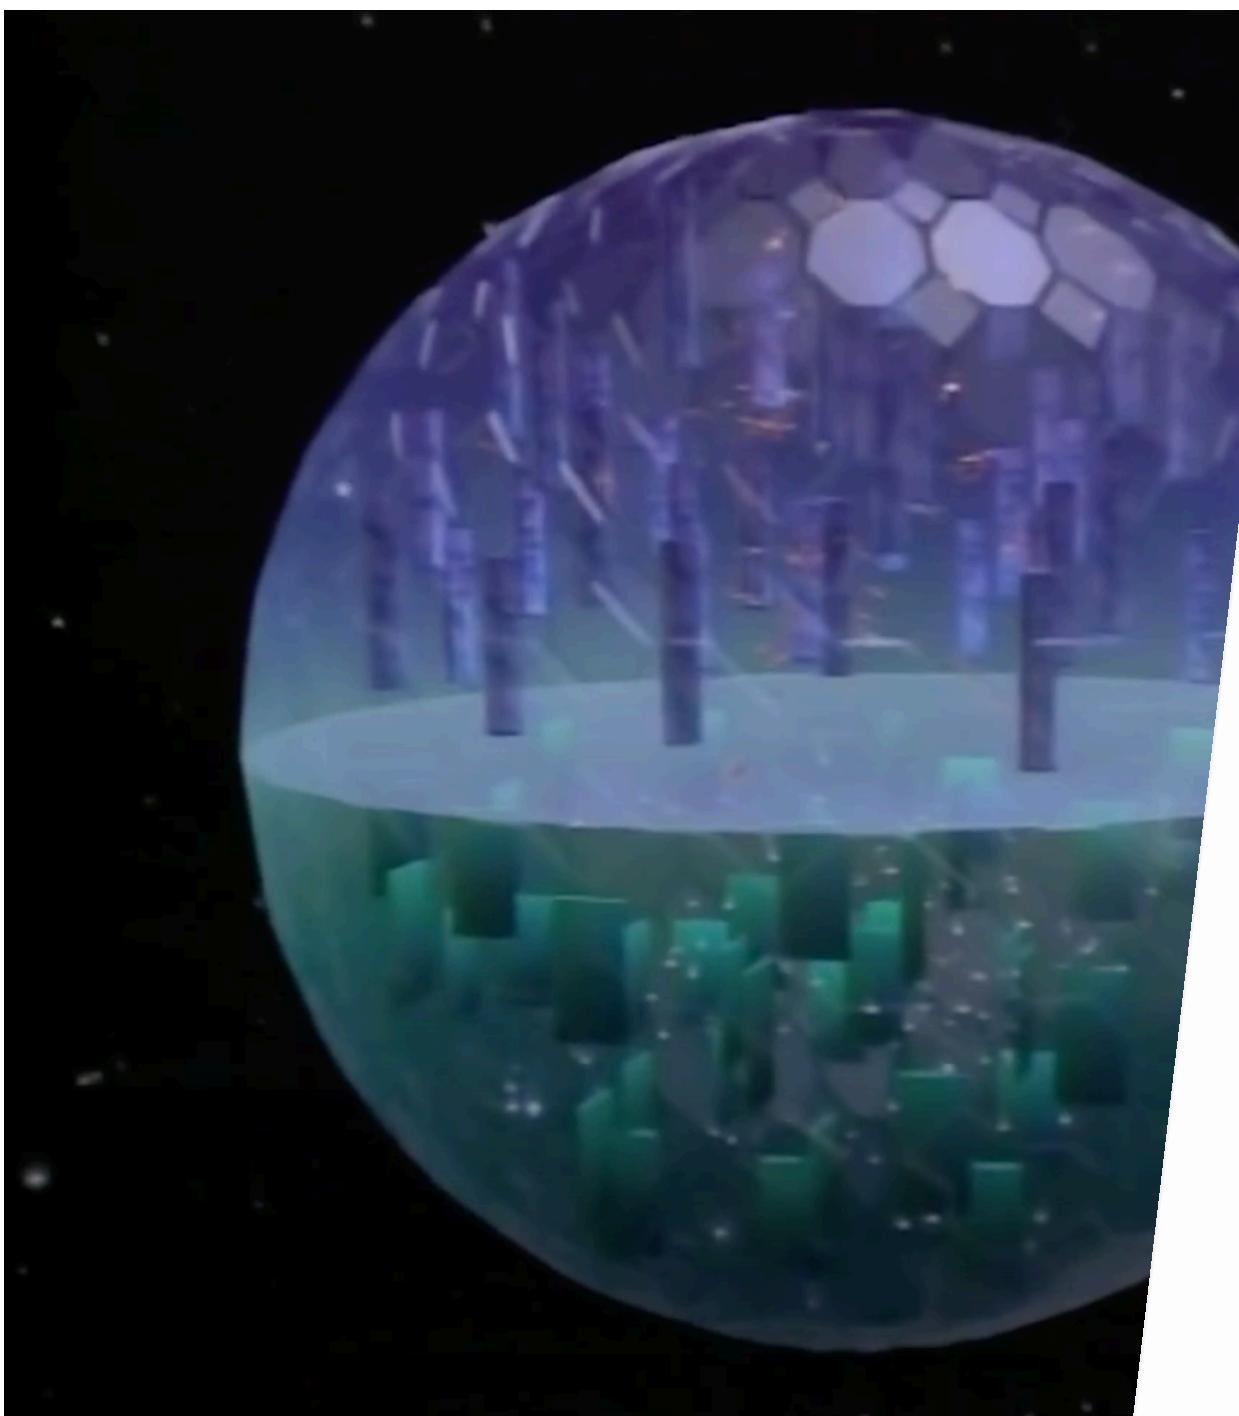


SIGGRAPH 1987

Breaking the Ice

Flocks, Herds, and Schools:
A Distributed Behavioral Model





Flocks, Herds, and Schools: A Distributed Behavioral Model

Craig W. Reynolds
Symbolics Graphics Division

1401 Westwood Boulevard
Los Angeles, California 90024

(Electronic mail: cwr@Symbolics.COM)

Abstract

The aggregate motion of a flock of birds, a herd of land animals, or a school of fish is a beautiful and familiar part of the natural world. But this type of complex motion is rarely seen in computer animation. This paper explores an approach based on simulation as an alternative to scripting the paths of each bird individually. The simulated flock is an elaboration of a particle system, with the simulated birds being the particles. The aggregate motion of the simulated flock is created by a distributed behavioral model much like that at work in a natural flock; the birds choose their own course. Each simulated bird is implemented as an independent actor that navigates according to its local perception of the dynamic environment, the laws of simulated physics that rule its motion, and a set of behaviors programmed into it by the "animator." The aggregate motion of the simulated flock is the result of the dense interaction of the relatively simple behaviors of the individual simulated birds.

Categories and Subject Descriptors: I.2.10 [Artificial Intelligence]: Vision and Scene Understanding; I.3.5 [Computer Graphics]: Computational Geometry and Object Modeling; I.3.7 [Computer Graphics]: Three-Dimensional Graphics and Realism—Animation; I.6.3 [Simulation and Modeling]: Applications.

General Terms: Algorithms, design.

Additional Key Words, and Phrases: flock, herd, school, bird, fish, aggregate motion, particle system, actor, flight, behavioral animation, constraints, path planning.

Introduction

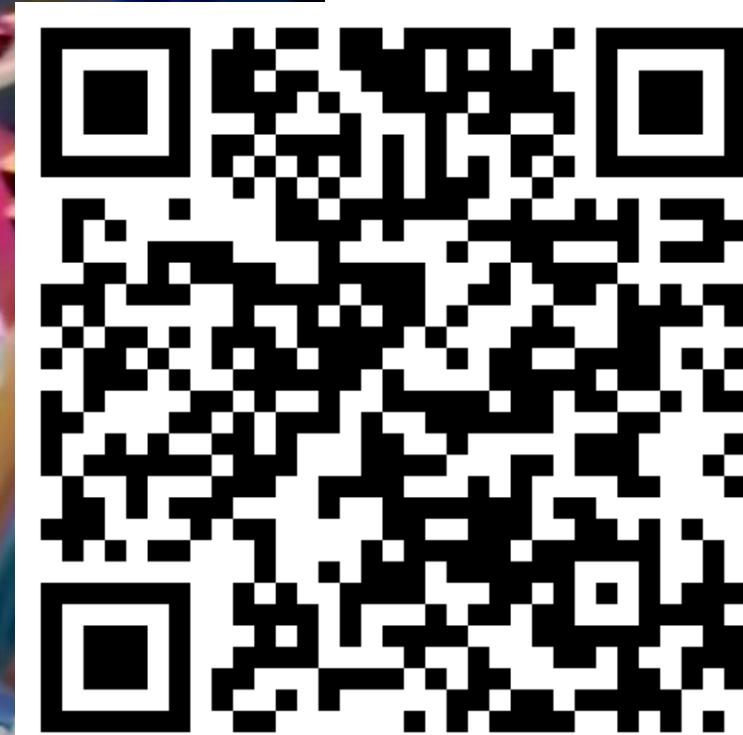
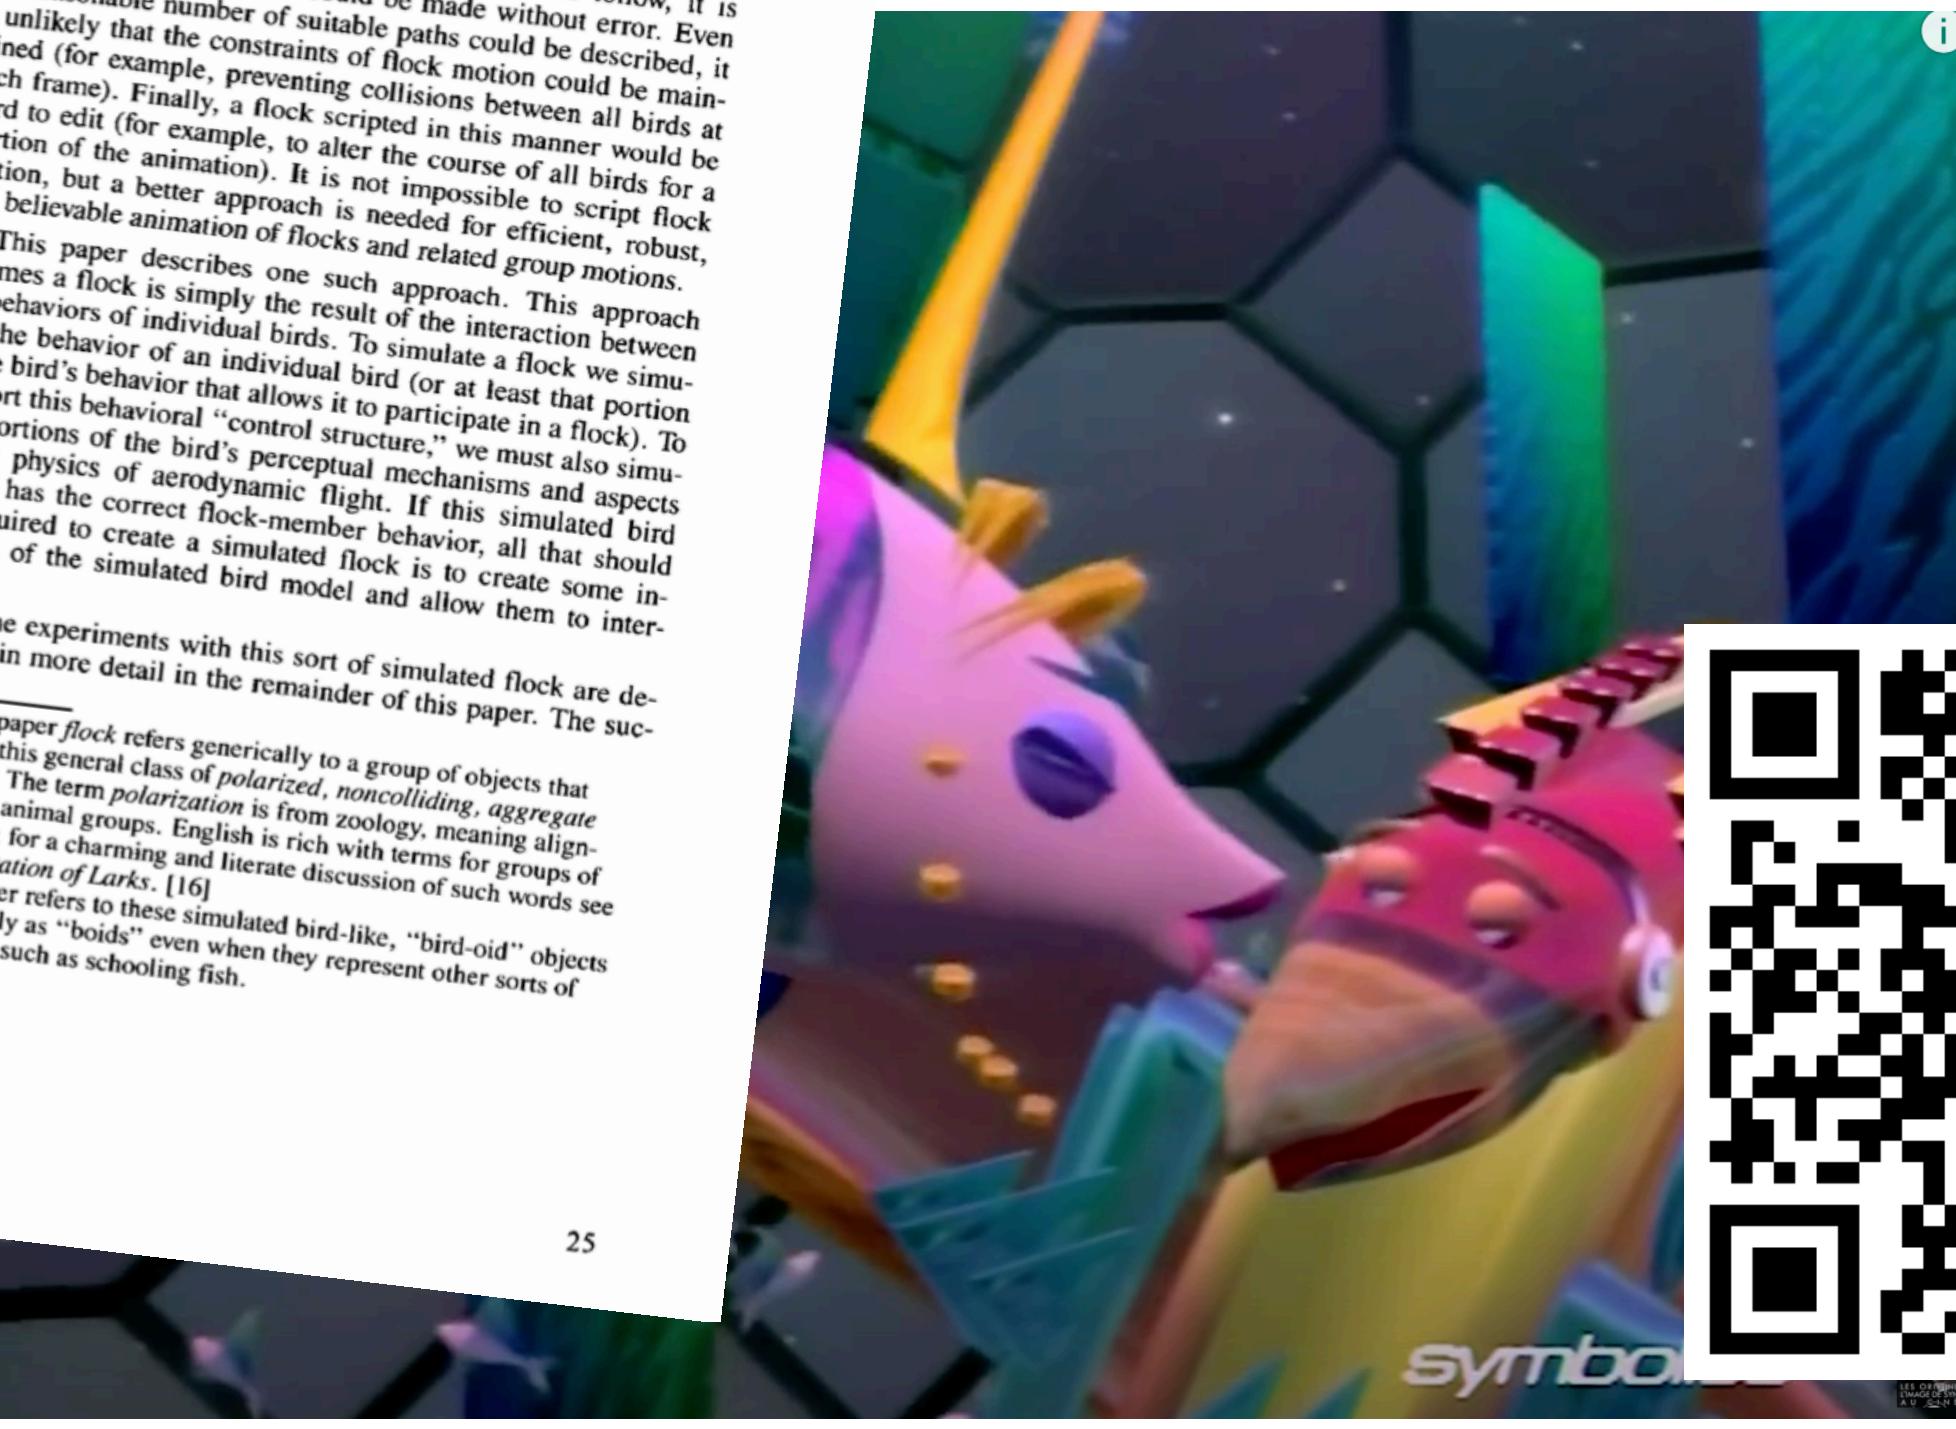
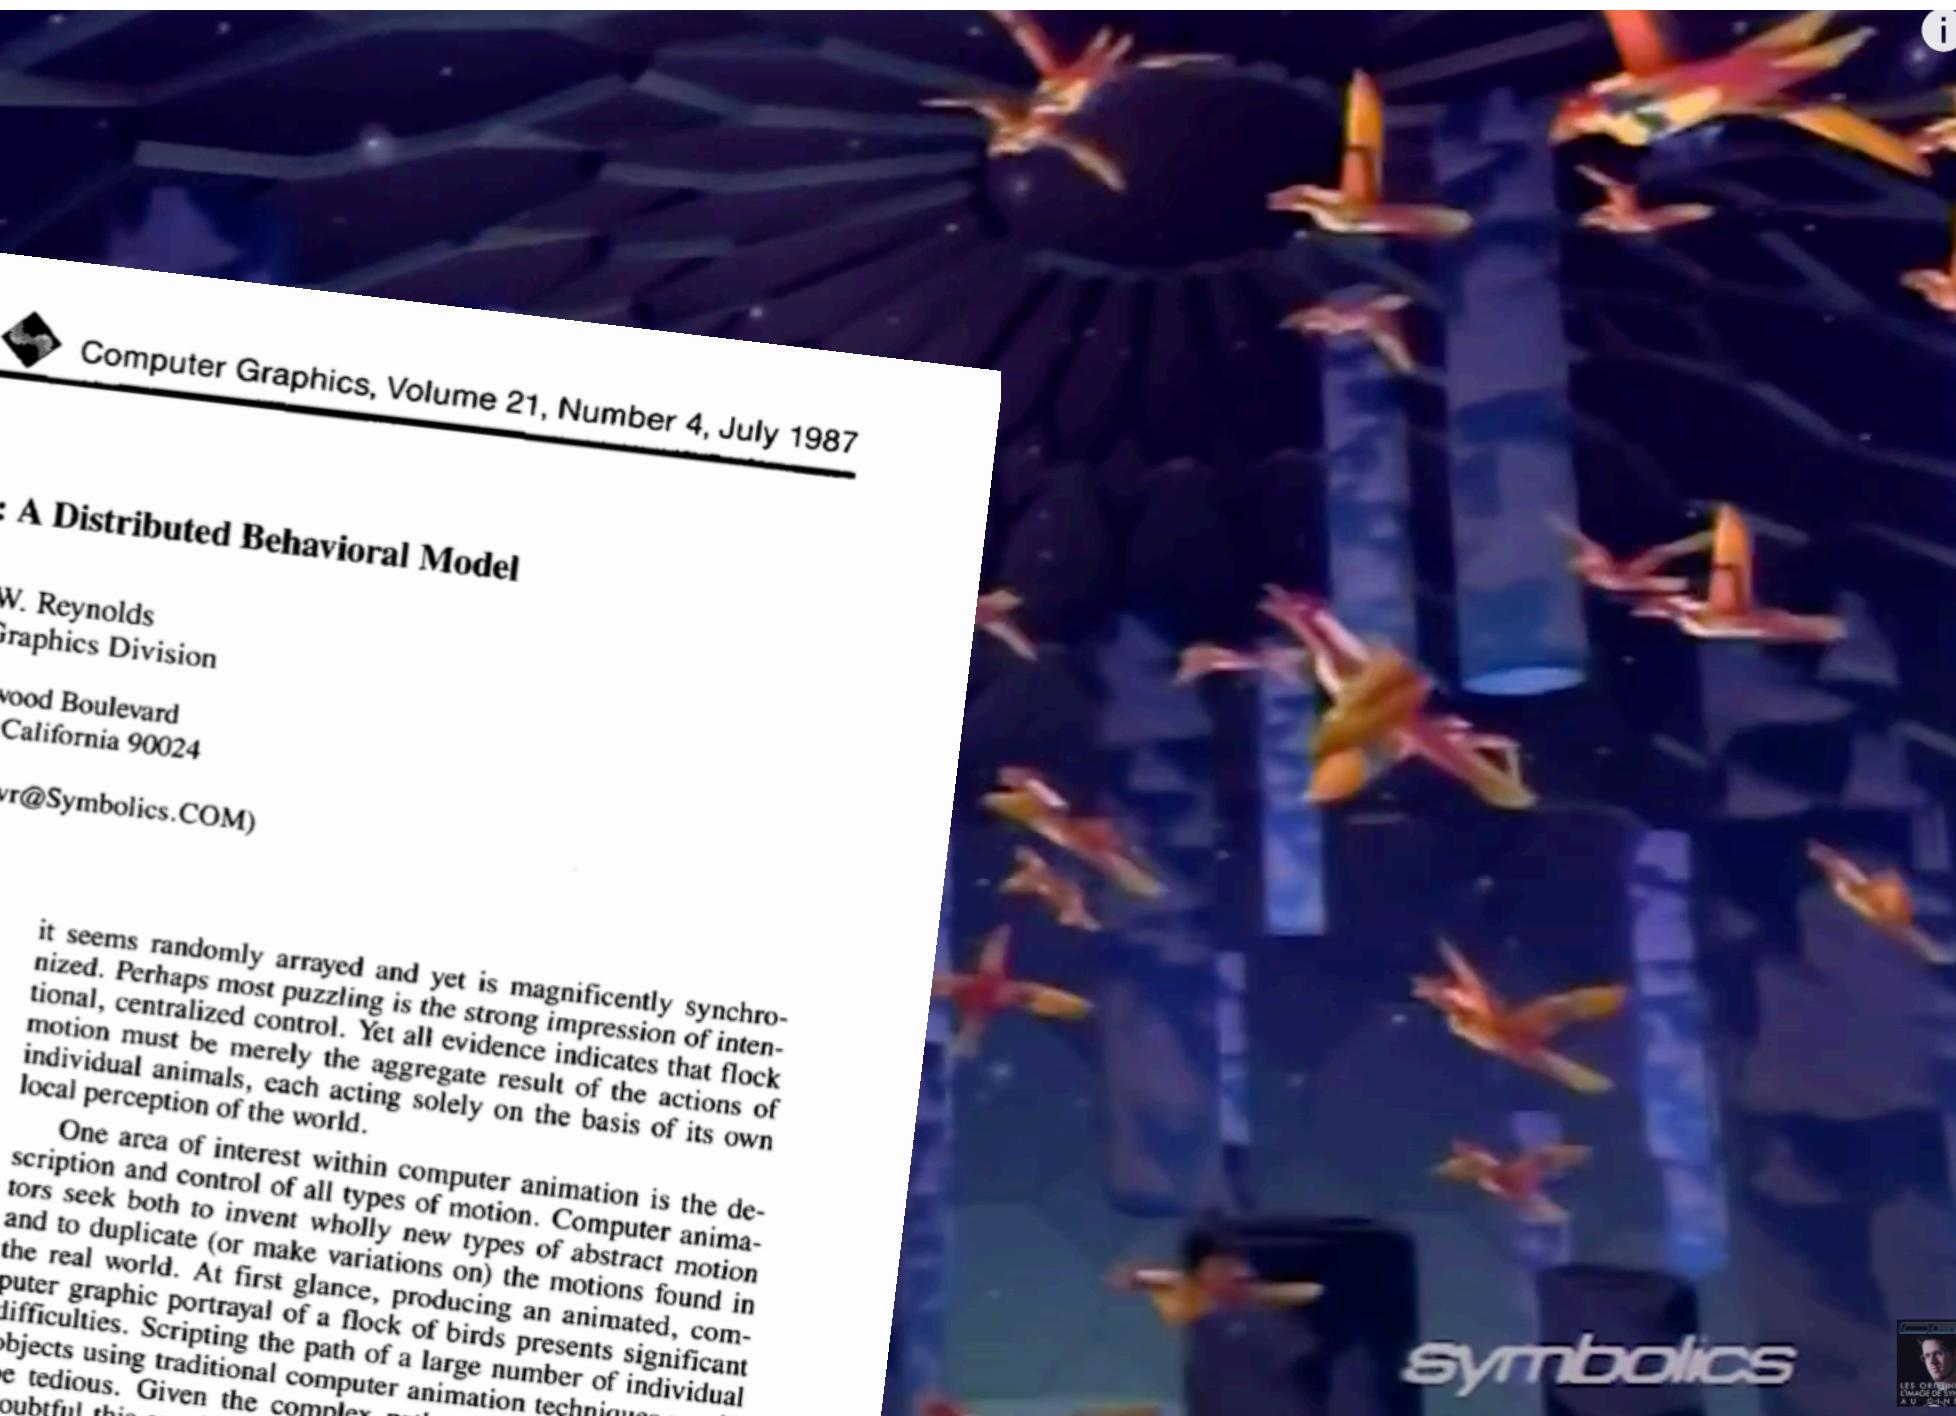
The motion of a flock of birds is one of nature's delights. Flocks and related synchronized group behaviors such as schools of fish or herds of land animals are both beautiful to watch and intriguing to contemplate. A flock* exhibits many contrasts. It is made up of discrete birds yet overall motion seems fluid; it is simple in concept yet is visually complex,

Permission to copy without fee all or part of this material is granted provided that the copies are not made or distributed for direct commercial advantage, the ACM copyright notice and the title of the publication and its date appear, and notice is given that copying is by permission of the Association for Computing Machinery. To copy otherwise, or to republish, requires a fee and/or specific permission.

© 1987

ACM-0-89791-227-6/87/007/0025

\$00.75



**SYMBOLICS
GRAPHICS DIVISION**

Boids and inverse design

2024-2025

Boids

Keywords: multi-agent, inverse design, optimization, evolutionary computation, genetic algorithm, boids, flocks, herds, schools, crowds, traffic

Introduction

Simulation models of multi-agent motion are used in many fields. These include animation, games, biology, robotic swarms, urban planning, training autonomous vehicles, and other applications. Since the 1980s, several models of group motion have been developed, including *boids* and others (see [Related Work](#)). These allow creating simulations of flocks and other group motions, which most observers will recognize as flocking. This paper will refer to bird

Evoflock: evolved inverse design of multi-agent motion

Craig Reynolds
unaffiliated researcher
cwr@red3d.com

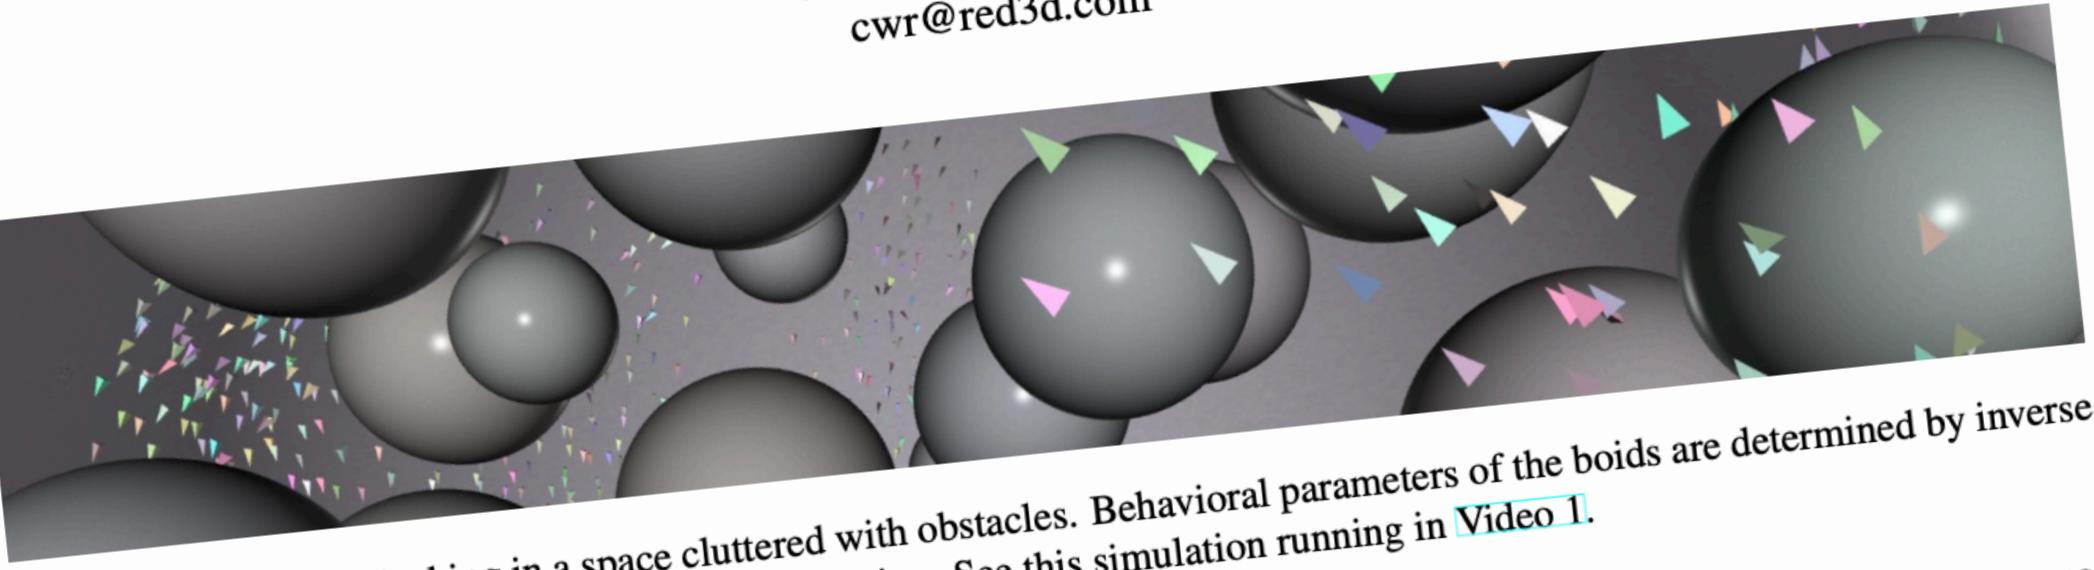


Figure 1: 1000 boids flocking in a space cluttered with obstacles. Behavioral parameters of the boids are determined by inverse design, using multi-objective evolutionary optimization. See this simulation running in [Video 1](#).

Abstract

This paper describes an automatic method for *adjusting* or *tuning* models of multi-agent motion. Simulating the motion of bird flocks, human crowds, vehicle traffic, and other multi-agent systems is a widely used technique. These simulations model the behavior of a single group member (bird, human, or vehicle). The group behaviors (flock, crowd, traffic) *emerge* from interactions between group members. These models typically have many numeric control parameters. Although each parameter may be intuitive in isolation, their interaction can be complex and nonlinear. It can be difficult to know how to adjust which parameters to change, leading to a long frustrating process of incremental changes. In this work, the desired group behavior is measured with an objective(/fitness/loss) function and optimized with a genetic algorithm. The objective function used here combines scoring that a flock's birds maintain proper spacing with neighbors, fly near a given speed, and avoid obstacles. Interestingly, the vivid alignment seen in bird flocks appears to emerge from maintaining proper spacing between flockmates.

flocks, with the assumption that other types of group motion (herds, schools, crowds, traffic, drone swarms) can be similarly modeled.

This project addresses the issue of *adjusting* or *tuning* multi-agent motion models, toward a given behavioral goal. For example, modifying an existing model of bird flocks to instead portray fish schools. Or starting from a model of flocking crows and change it to represent flocking sparrows. Similarly, starting from a generic abstract flock model and fitting it to field observations of a particular species of real birds in nature. Or to take a plausibly realistic model of a natural bird flock, and change it, say for storytelling purposes, to convey a flock of birds that are happy, or angry.

A boids-like simulation model typically has a collection of numeric parameters (“knobs”) that control its action. EvoFlock is a software framework that automatically finds a set of near-optimal parameters for a simulation-based multi-agent system. The flock model used for these experiments is a predefined hand-written “black box” model whose input is a *parameter set* (consisting of 15 numbers, see Table 2). The user provides a *objective* (also known as *fitness* or *loss*) function. It takes a candidate parameter set, runs a simulation, and returns a *score* reflecting how well the behavioral goals were met. The optimization process runs (for about two hours on a laptop) and produces a high quality parameter set, as measured by the objective function.

For a user of this optimization framework, it is convenient that the flock model is treated as a black box. Almost no restrictions are imposed on a preexisting model by the optimization framework. There are no restrictions on model design, source code availability, or the programming language used. (As could be an issue if optimization required language tools such as automatic differentiation (Baydin et al., 2018).)

sign

Boid models are hard to “tune”

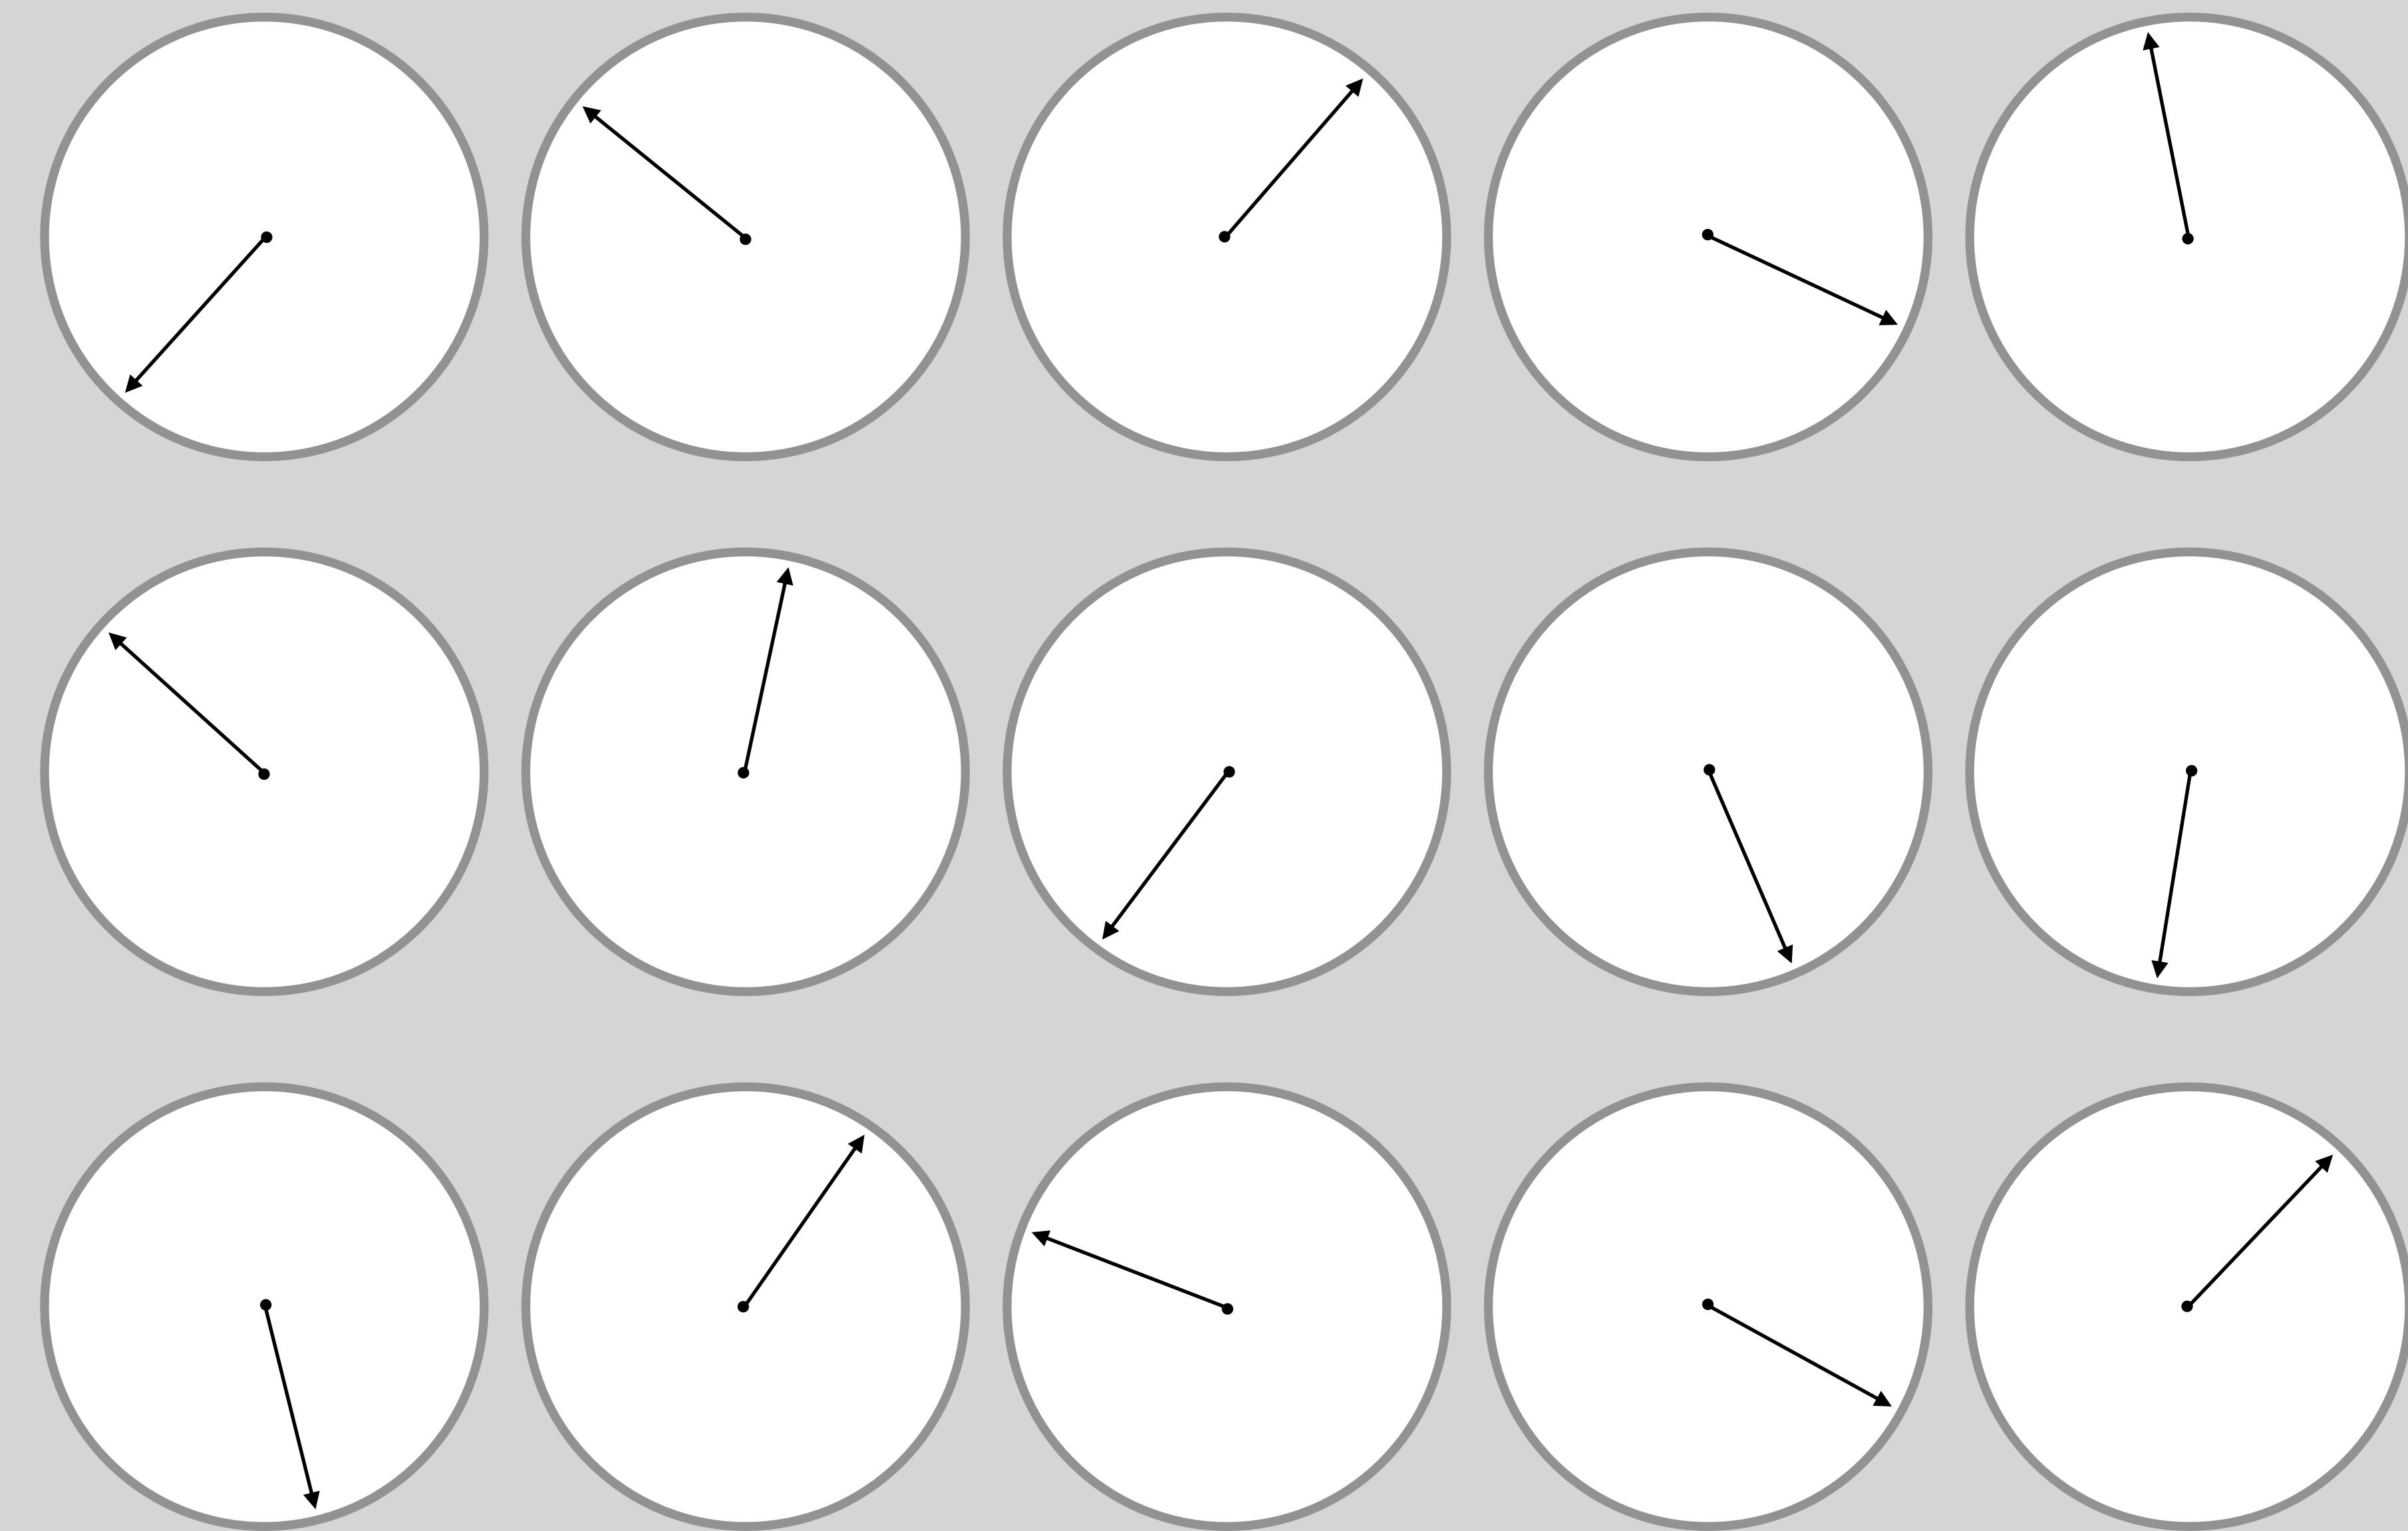
Boid models are hard to “tune”

They have lots of parameters,
which have nonlinear effect,
and all interact with each other.

Boid models are hard to “tune”

adjusting one knob requires
adjusting others to compensate

flock control panel



**How can we use optimization to tune
boid flocks to get a desired result?**

How can we use optimization to tune boid flocks to get a desired result?

Biological example: finch versus crow.
Animation example: happy versus afraid.

How can we use optimization to tune boid flocks to get a desired result?

“Inverse design.”

Create a metric/loss/fitness function.

Optimize toward design goal.

Preliminary work in progress

- Joint work with Gilbert Bernstein, Matthew Shang, and Jennifer Luo at University of Washington.
- Made reference boids model, first in Python then in C++.
- Two approaches being pursued:
 - Gradient descent with differentiable programming.
 - Gradient-free evolutionary optimization.
 - **Genetic algorithm: parameters for a “black box” boid model**
 - Genetic programming: steering programs from scratch

Preliminary work in progress

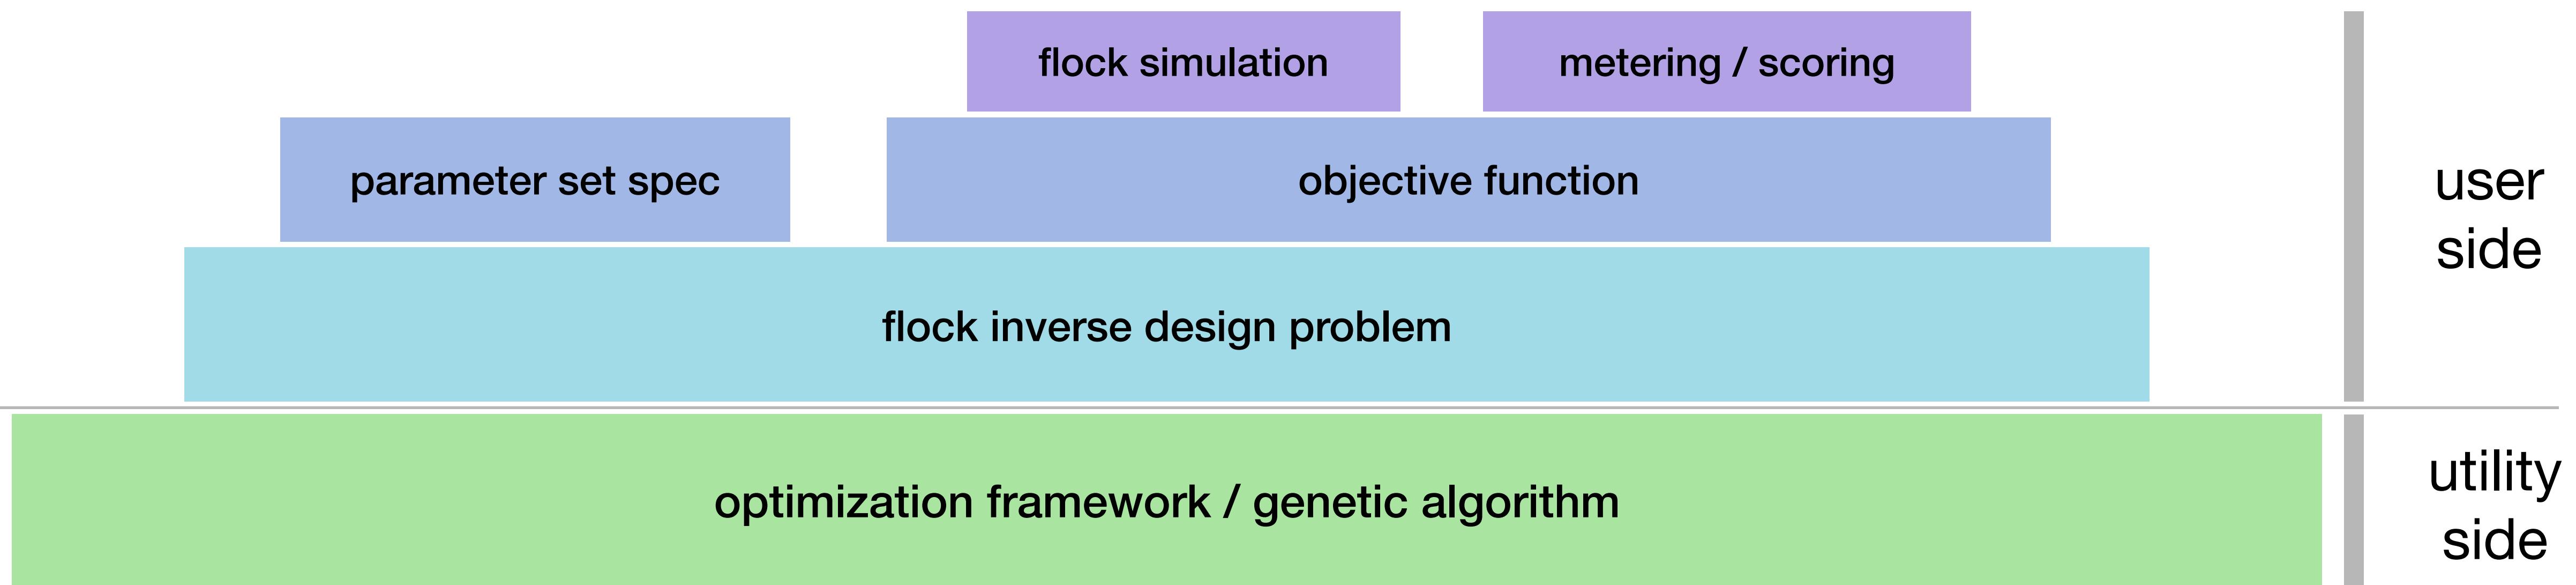
- Joint work with Gilbert Bernstein, Matthew Shang, and Jennifer Luo at University of Washington.
- Made reference boids model, first in Python then in C++.
- Two approaches being pursued:
 - Gradient descent with differentiable programming.
 - Gradient-free evolutionary optimization.
 - **Genetic algorithm: parameters for a “black box” boid model** 
 - Genetic programming: steering programs from scratch

Difficult to find gradient of chaotic system

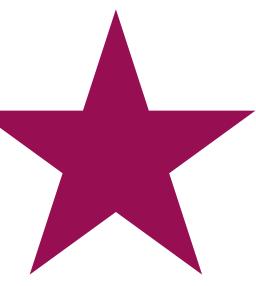
Multi-agent systems are chaotic.
Finding a principled gradient is hard.
That analysis is ongoing.

In parallel, using evolutionary optimization

EvoFlock components



Multi-objective optimization



Brooklyn

Firth of Forth



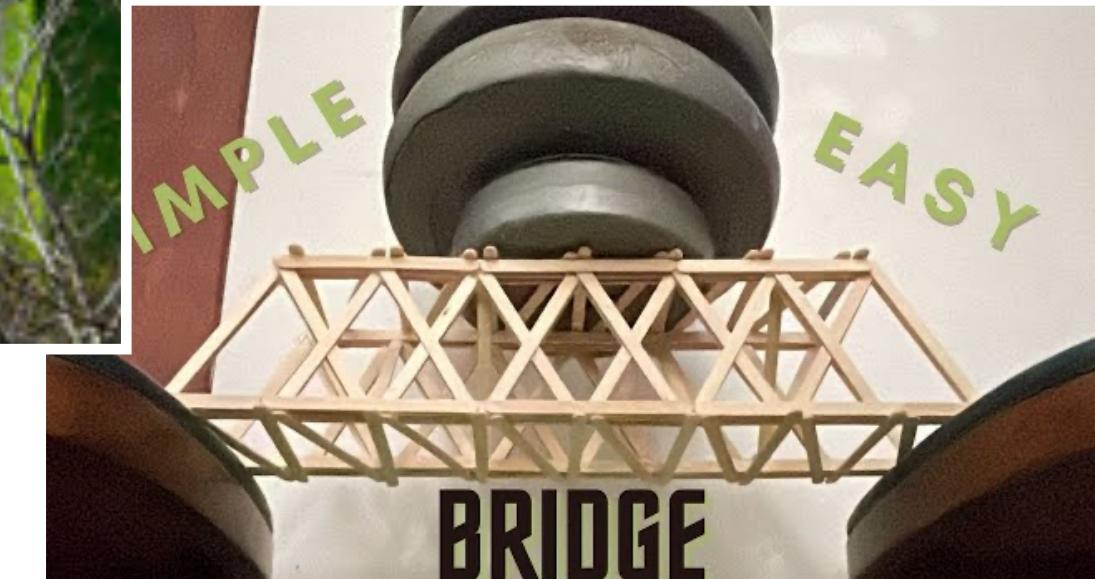
expensive

Strong



Natchez Trace Parkway

popsicle sticks

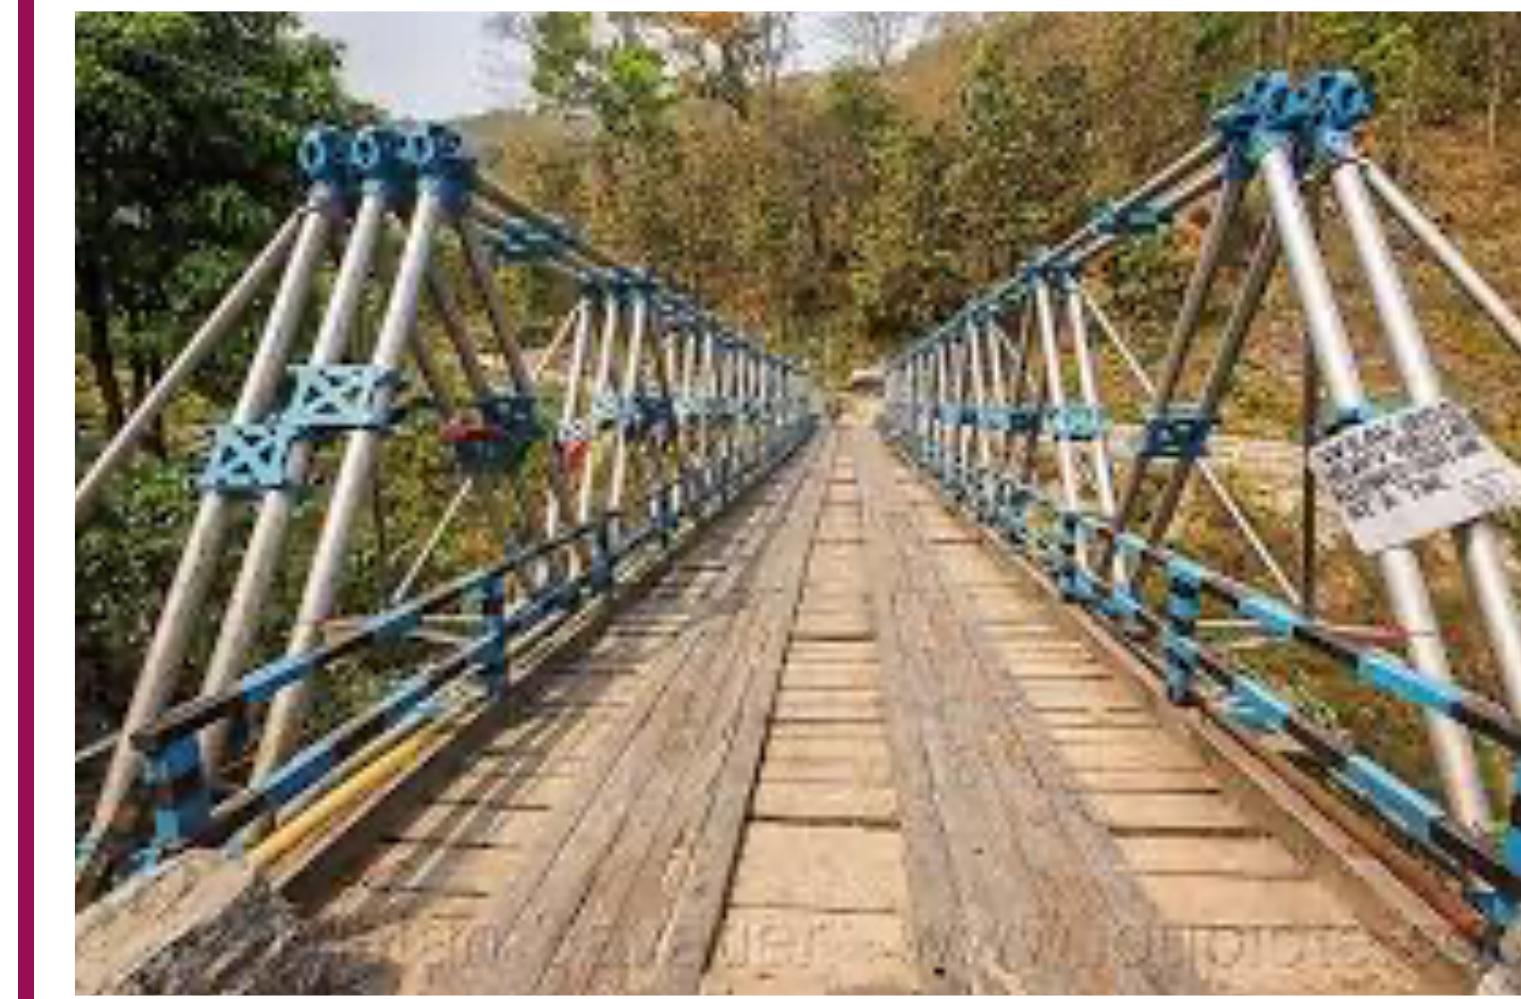


cheap

Tacoma Narrows



Firth of Tay



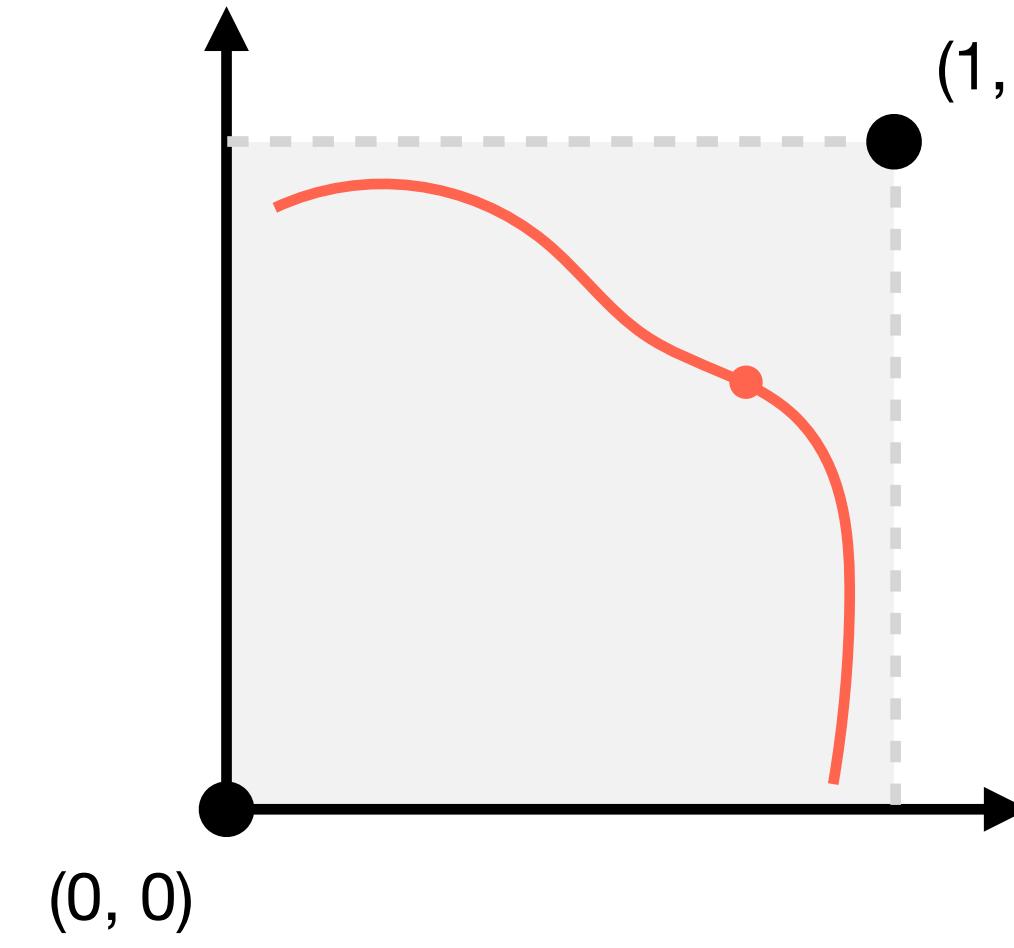
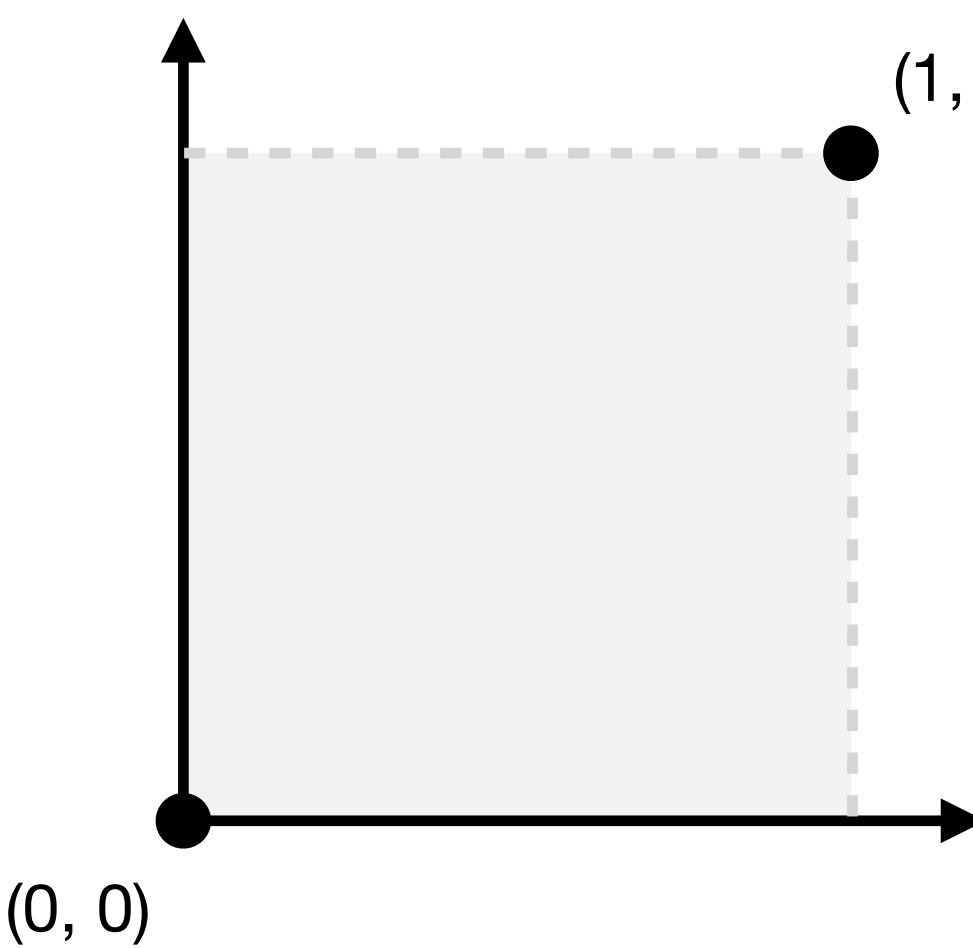
metal tubes, wood deck



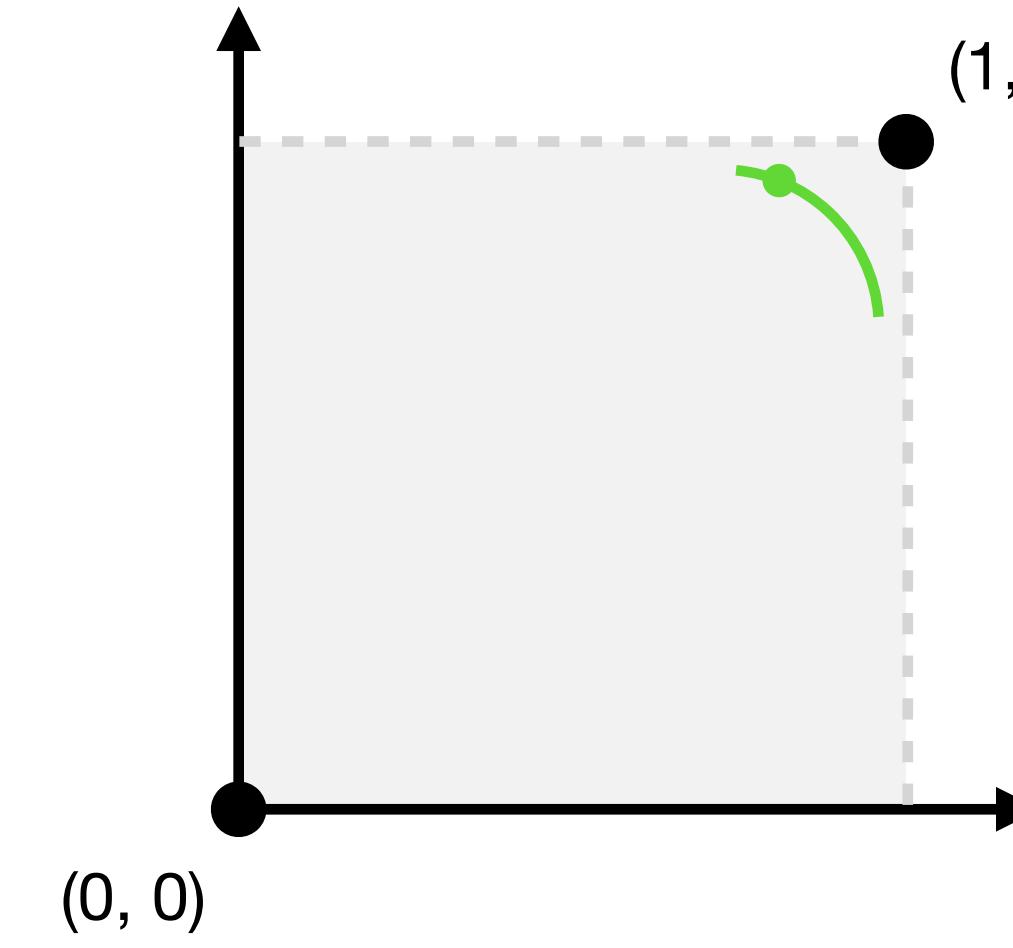
rickety wood bridge

weak

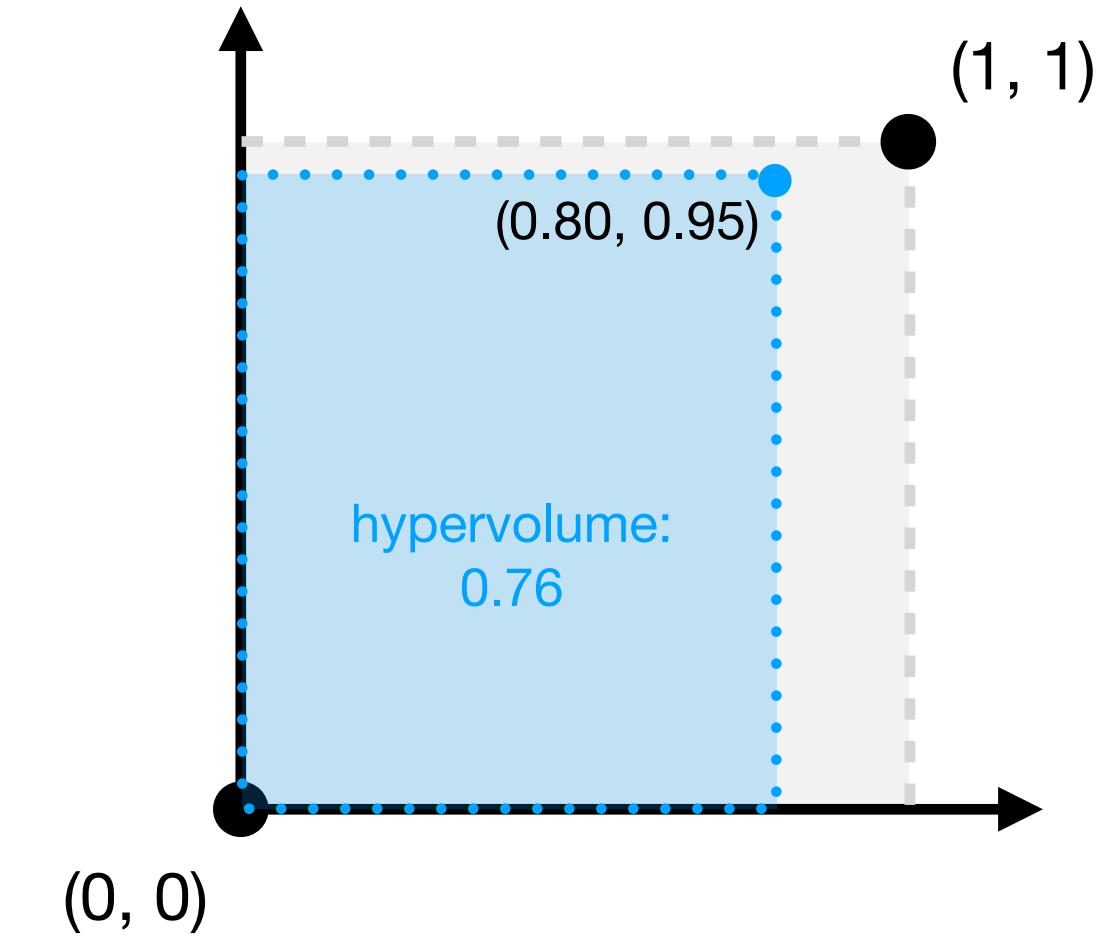
Multi-objective optimization – normalized fitness space.



Pareto front for two
mutually conflicting
objectives.

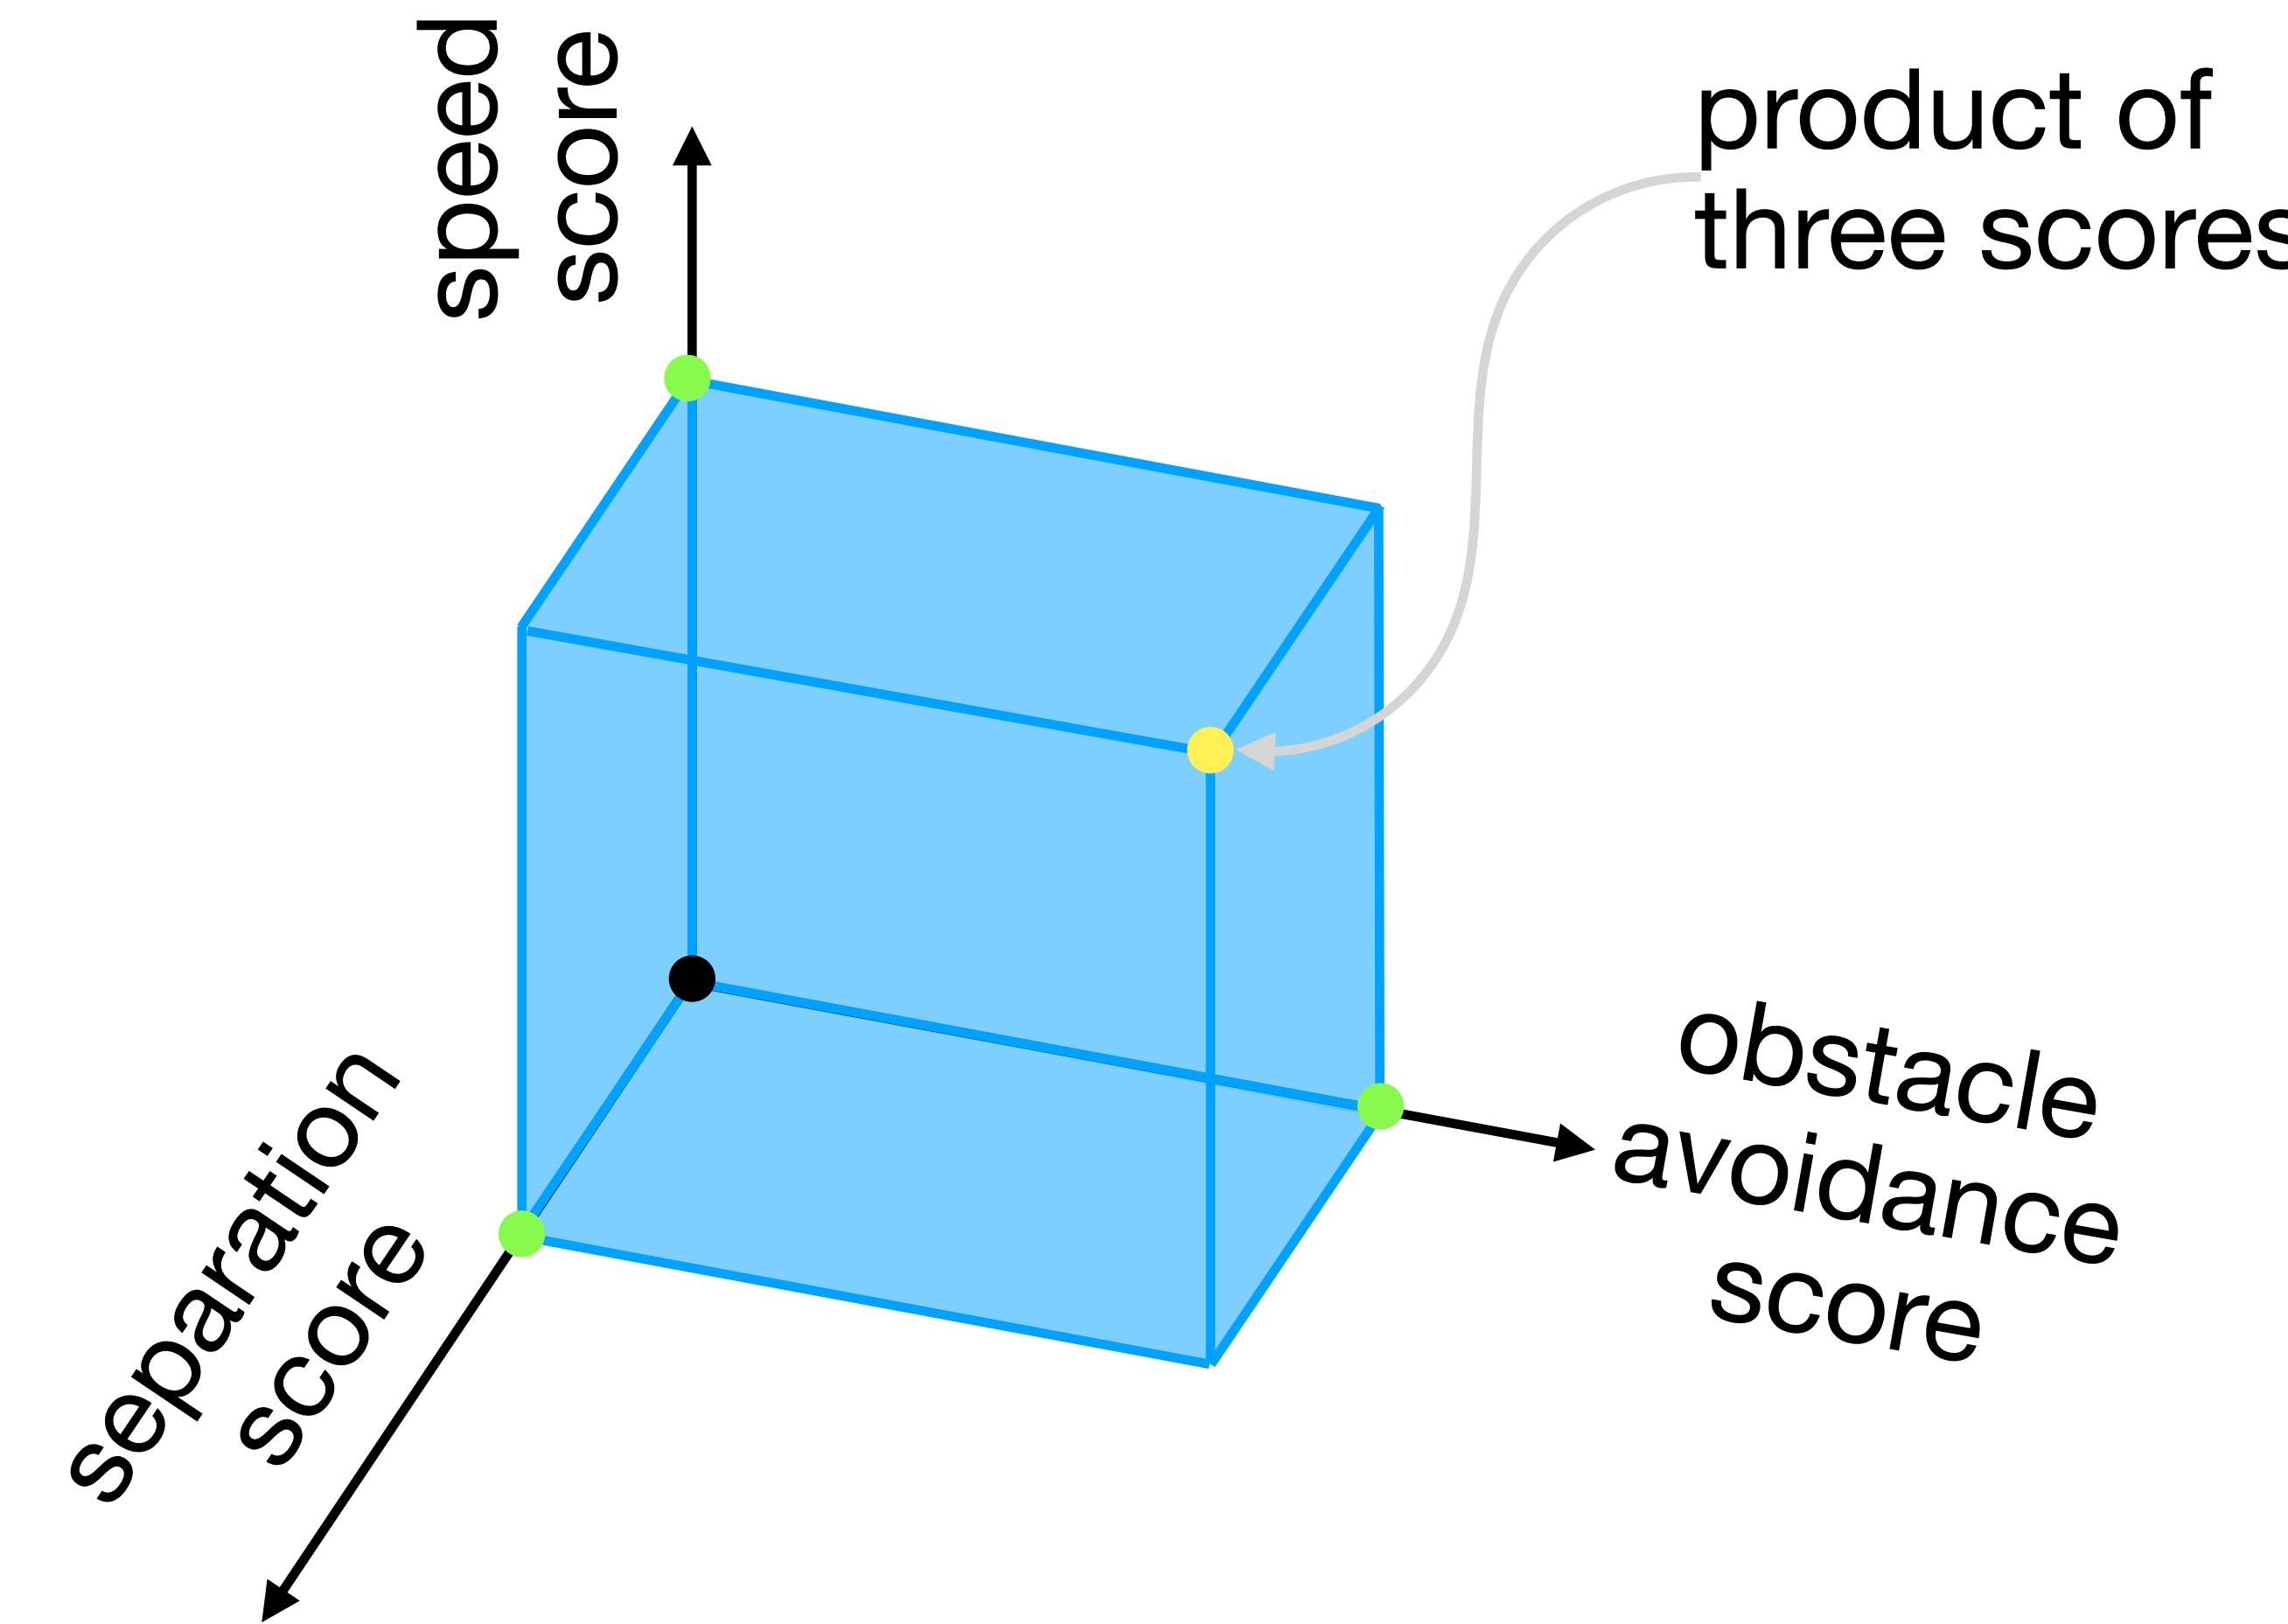


Pareto front for two
“mostly compatible”
objectives.

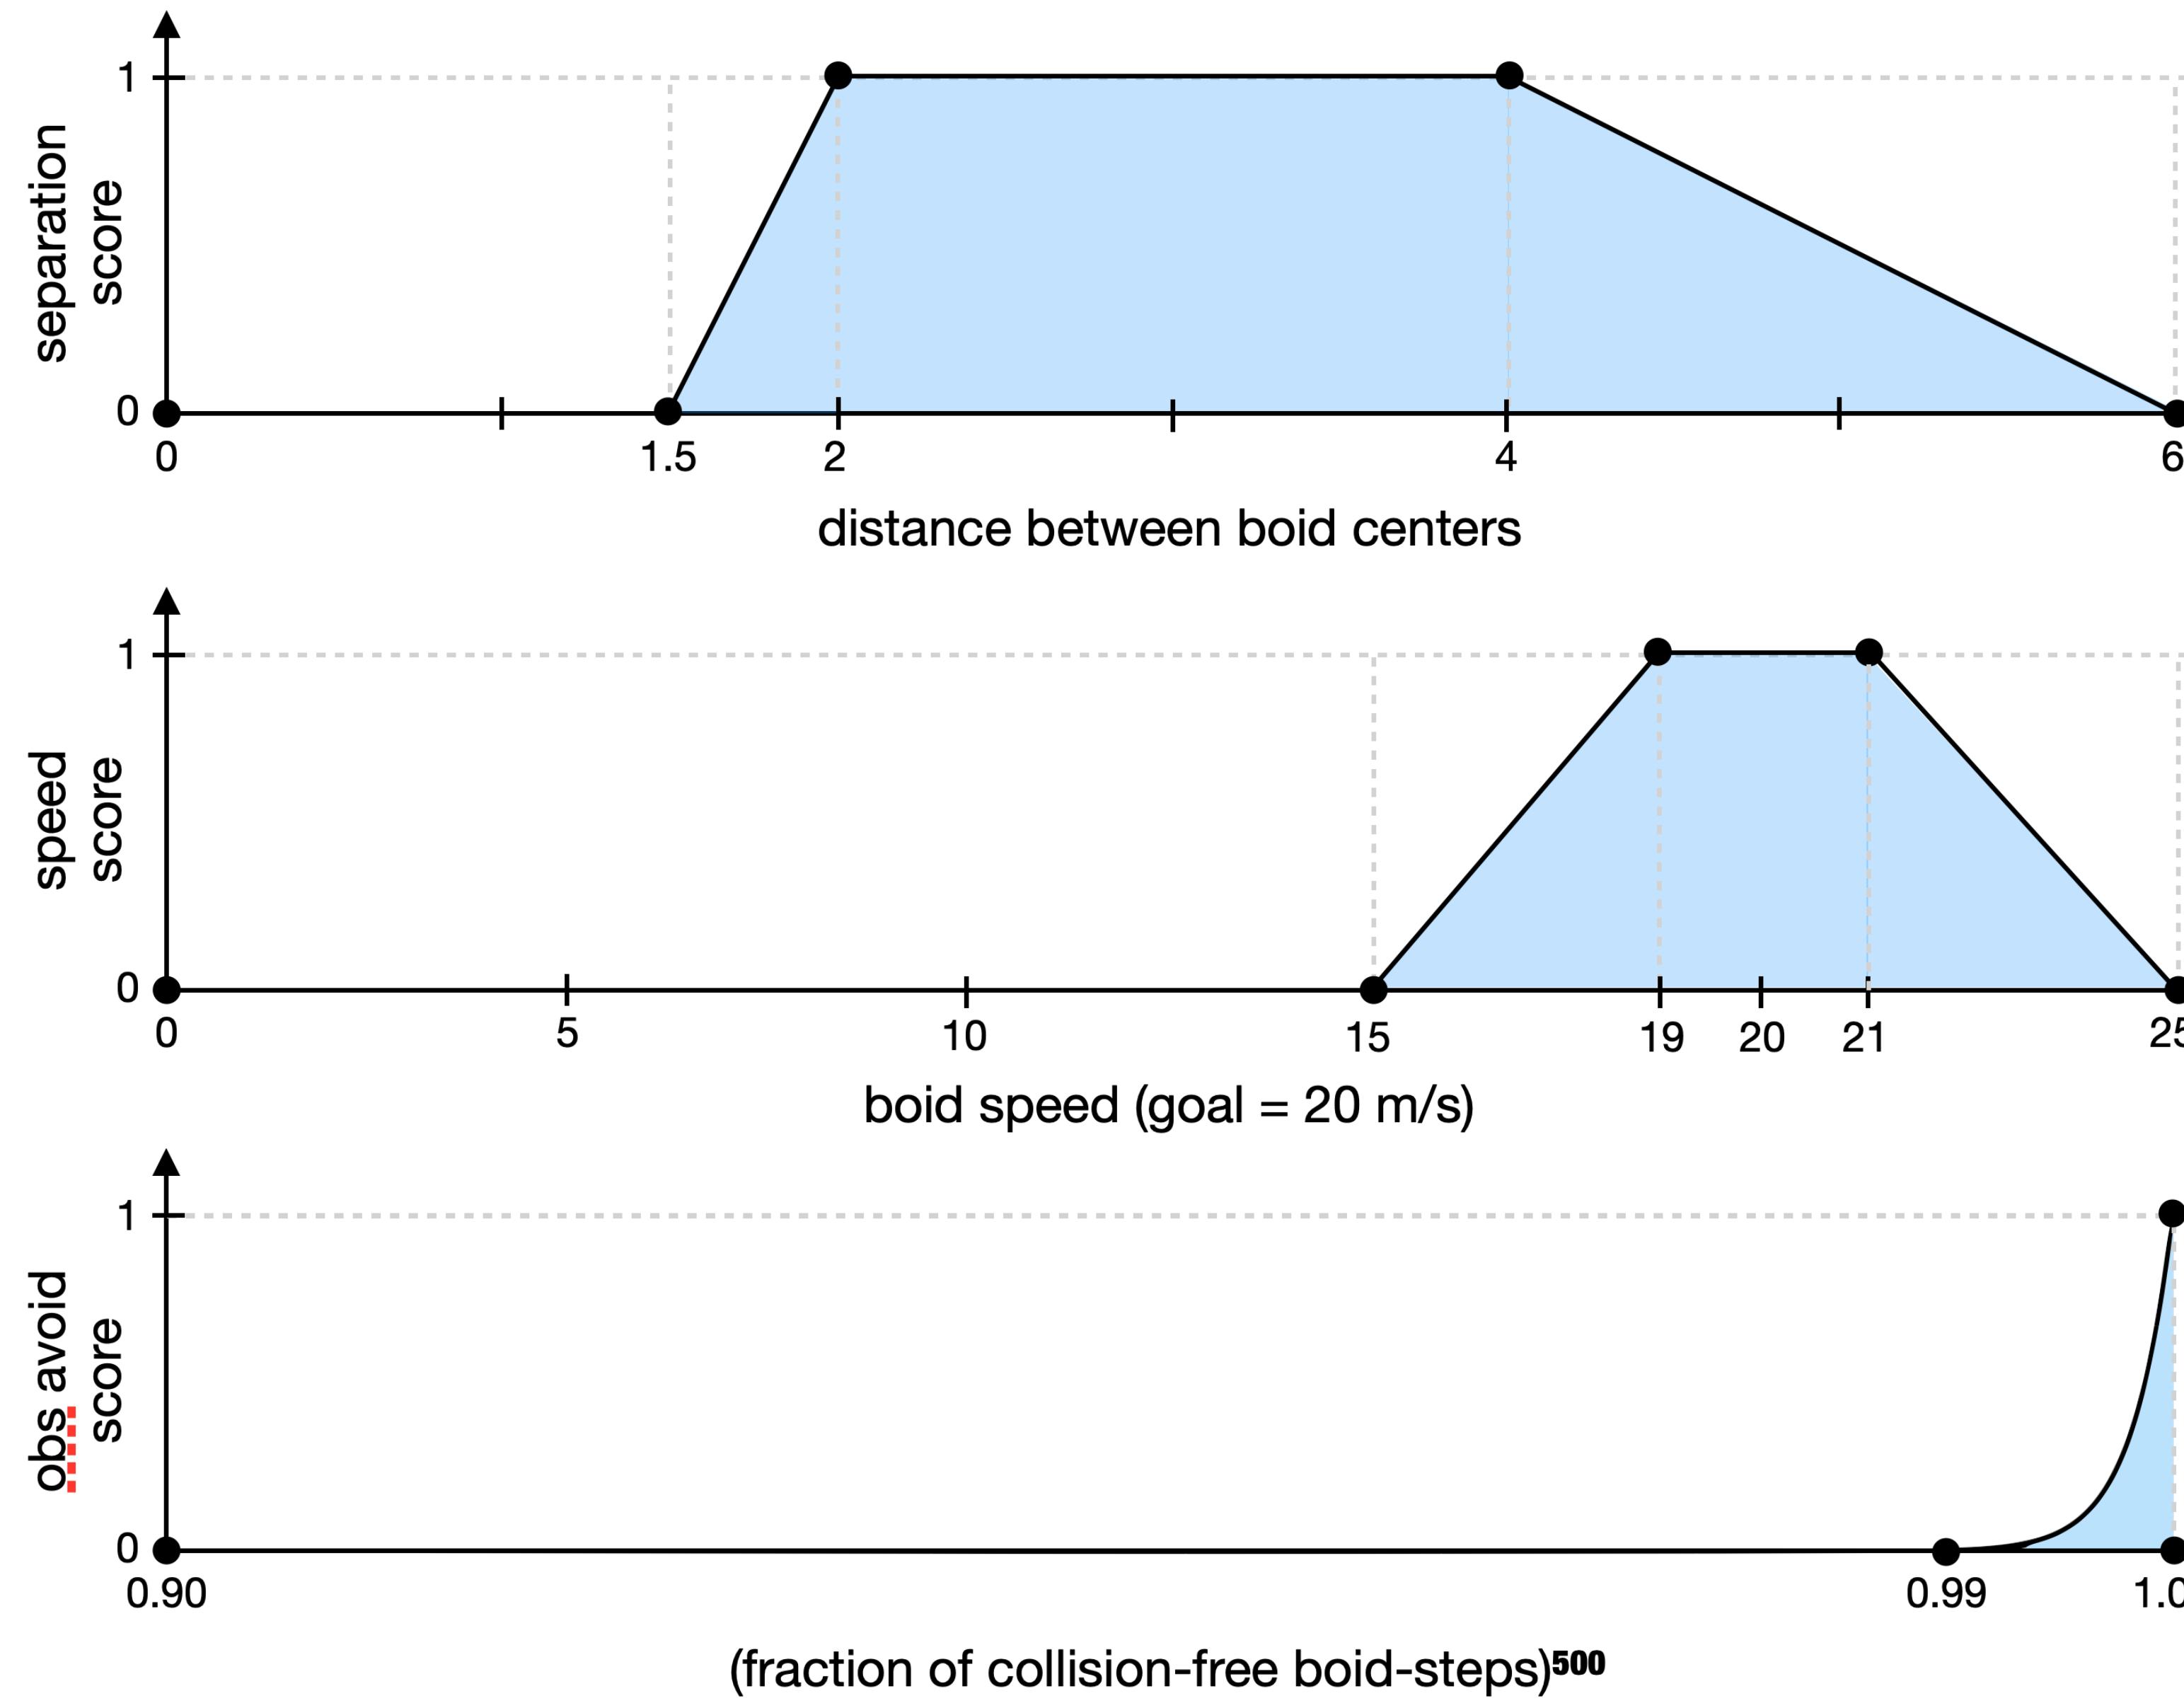


Scalarization (area)
for two objectives.

Hypervolume of three objectives



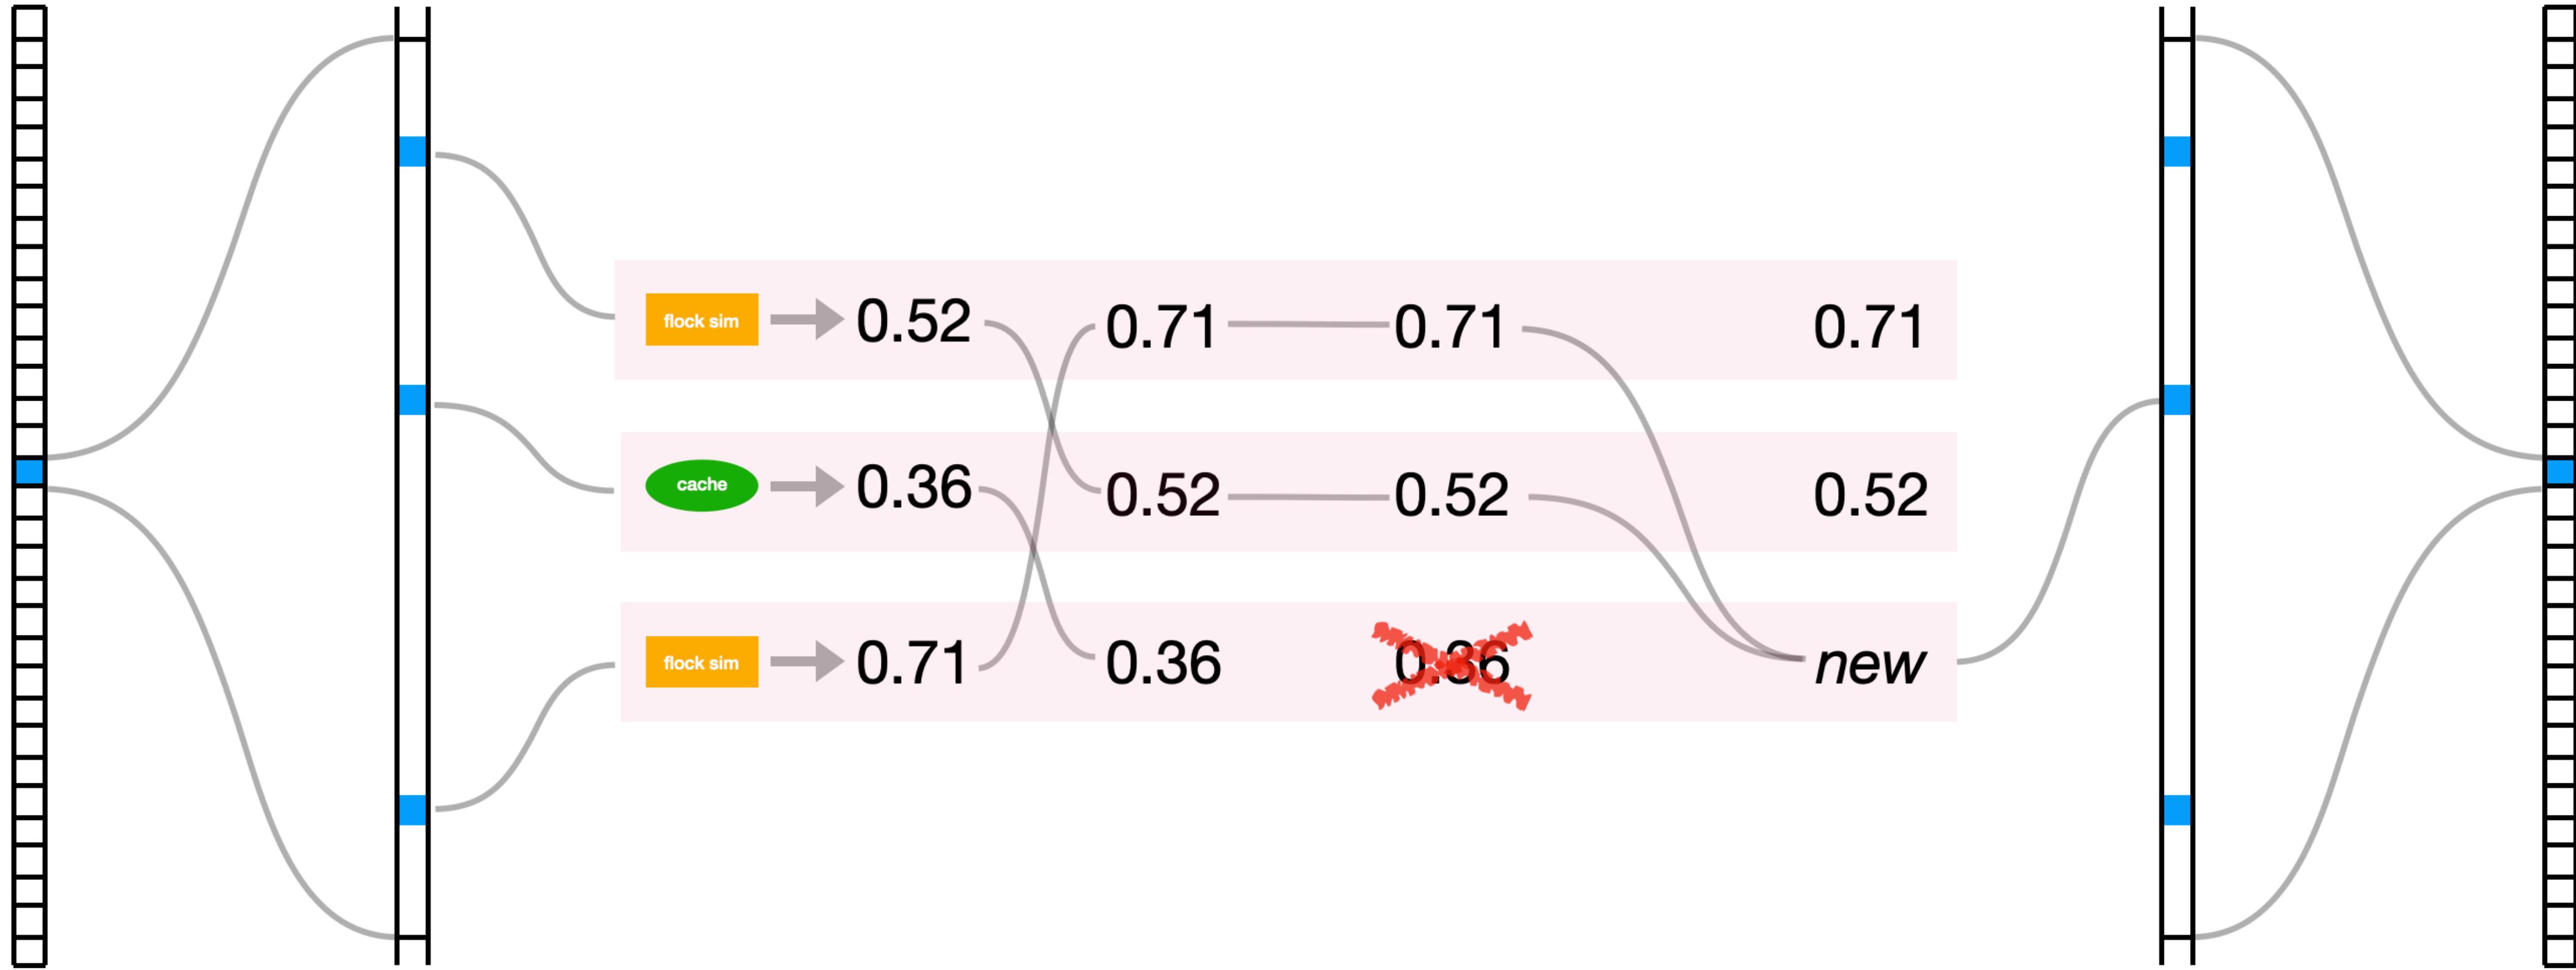
Score weighting functions: separation, speed, avoid



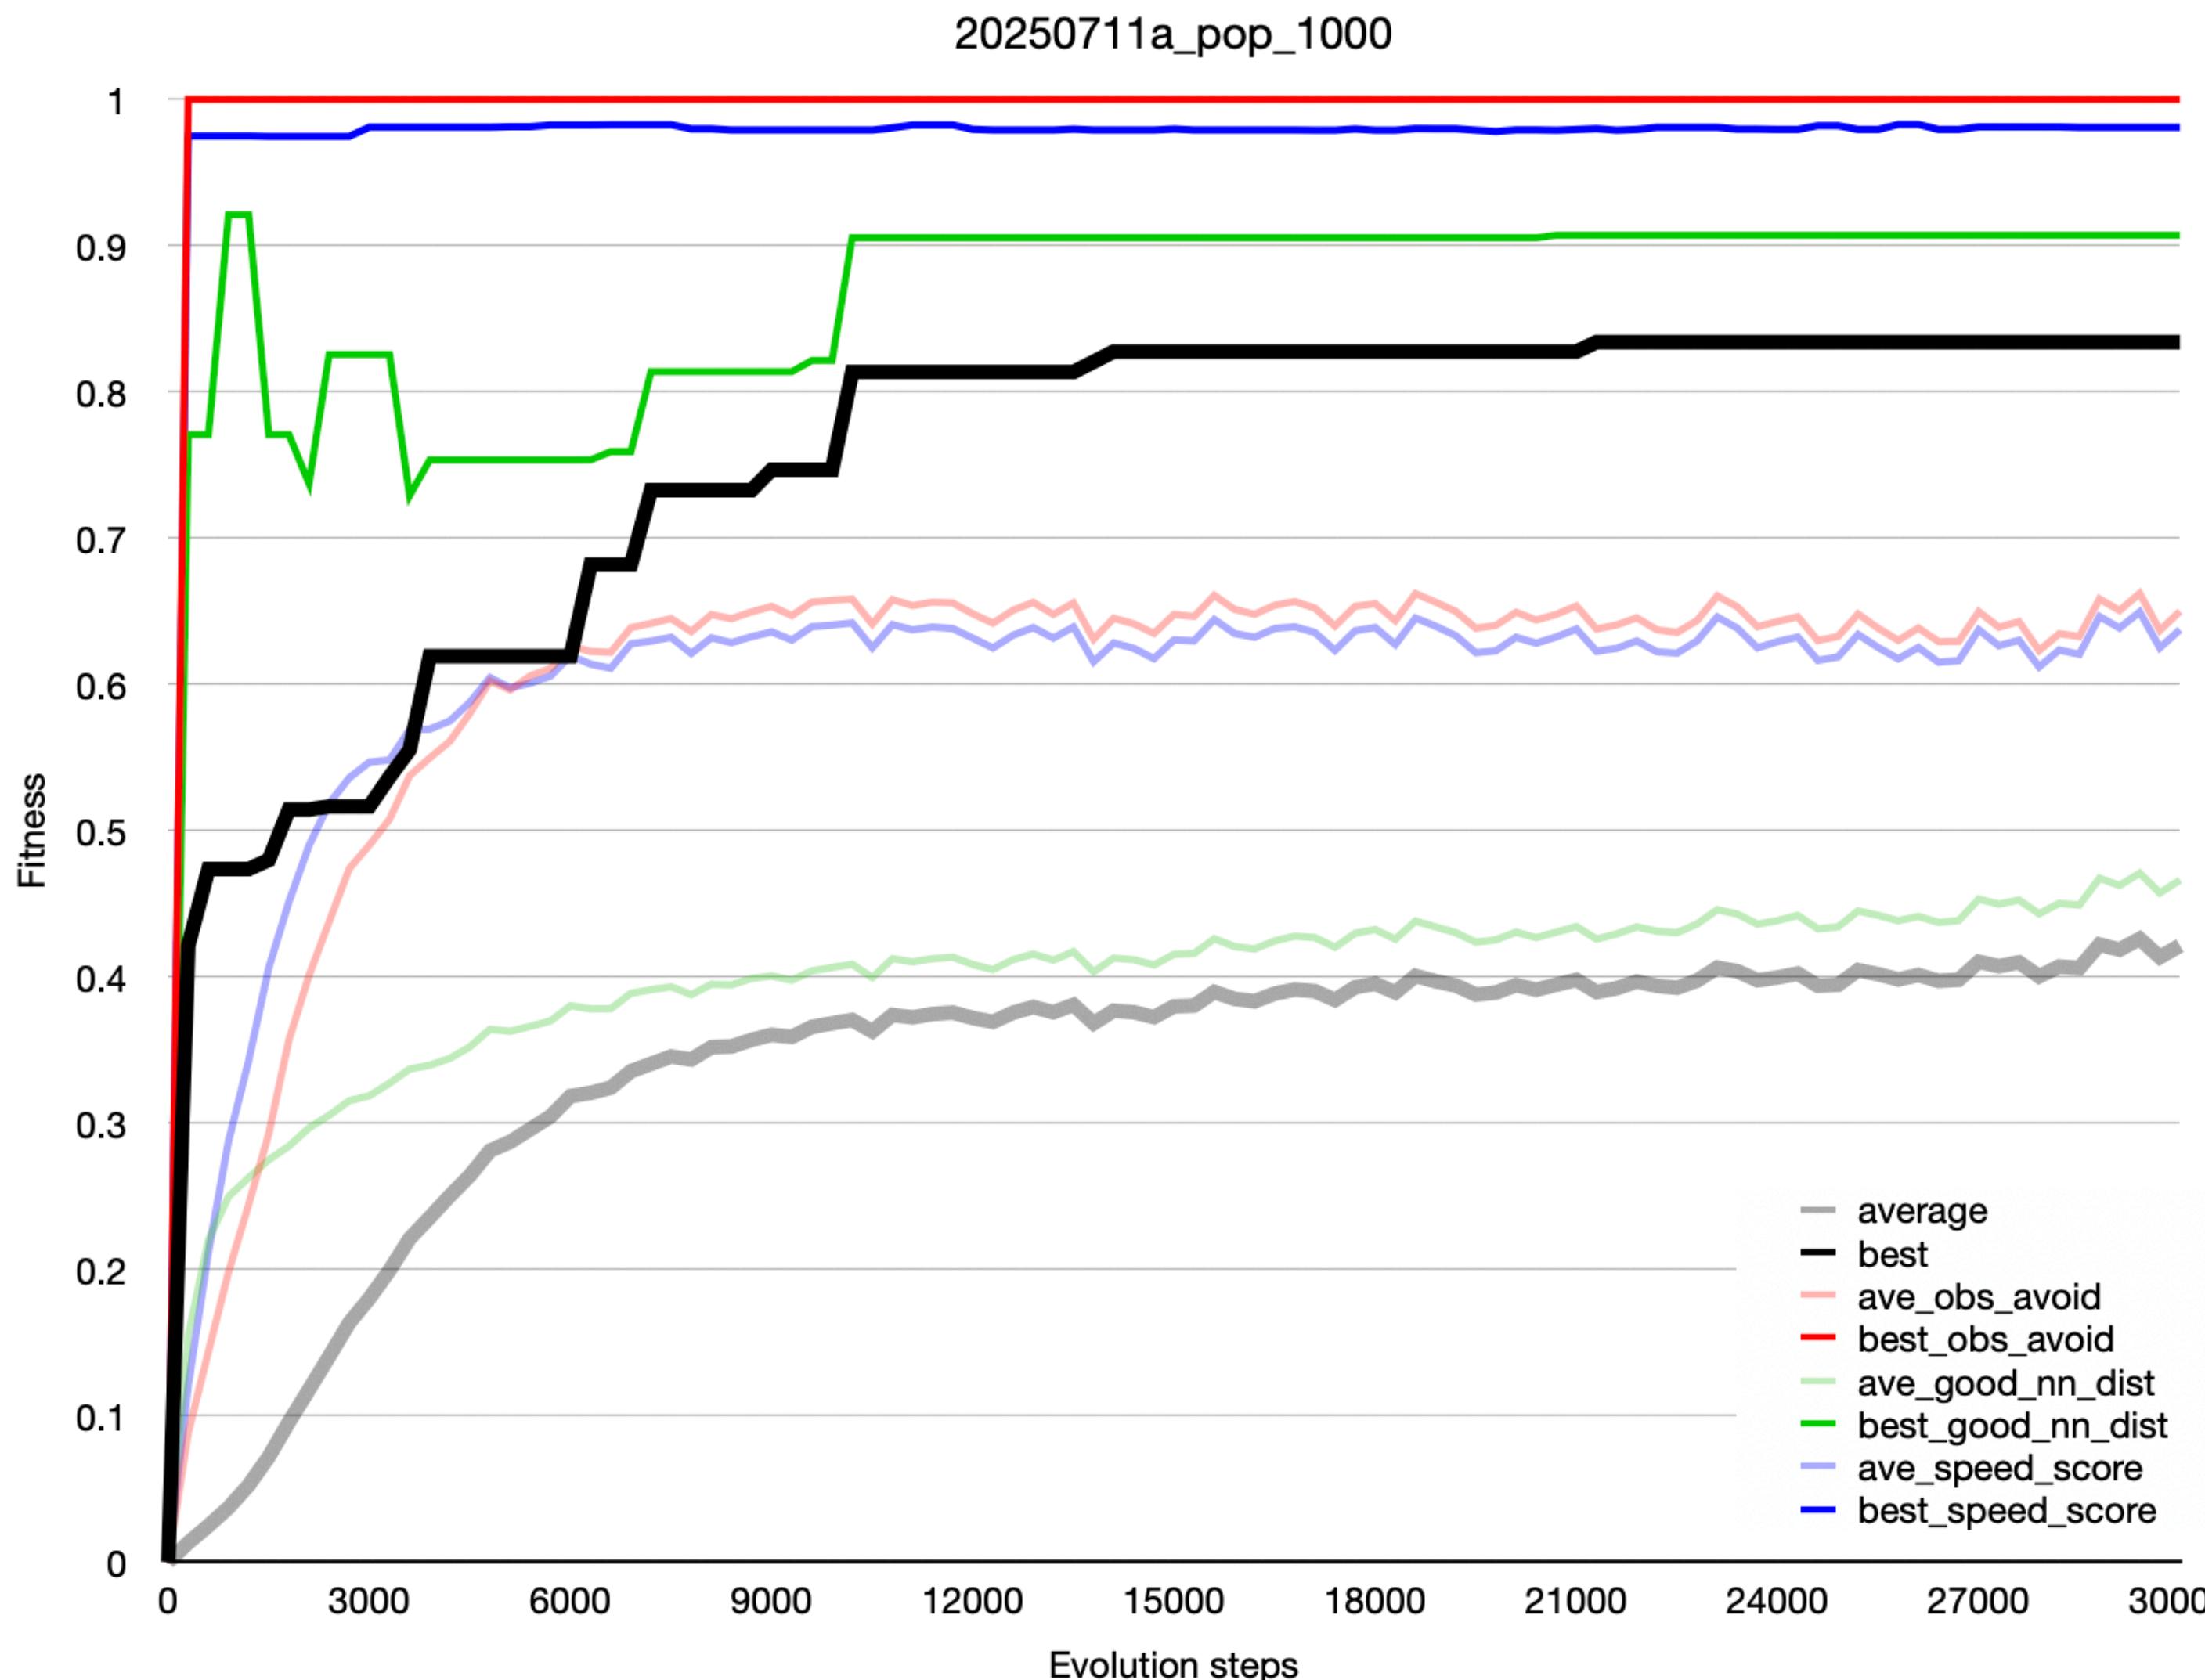
Evolutionary optimization

Single step of steady-state genetic algorithm (SSGA)

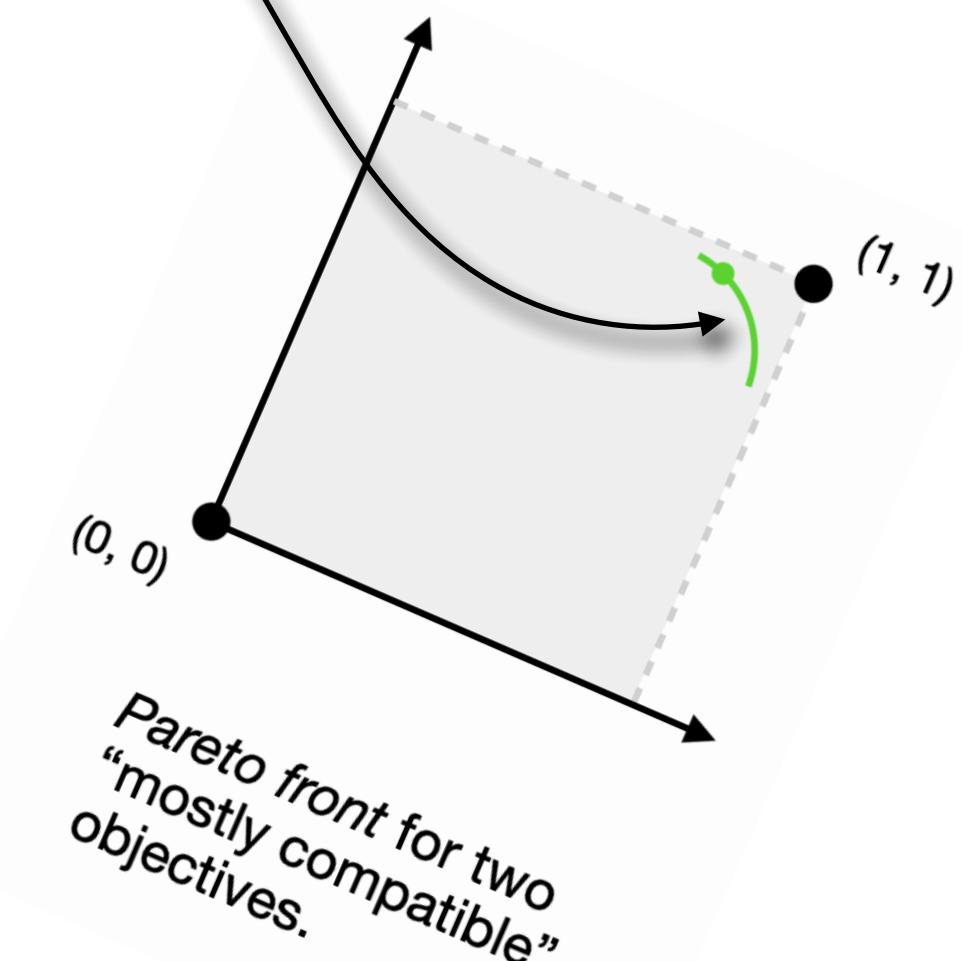
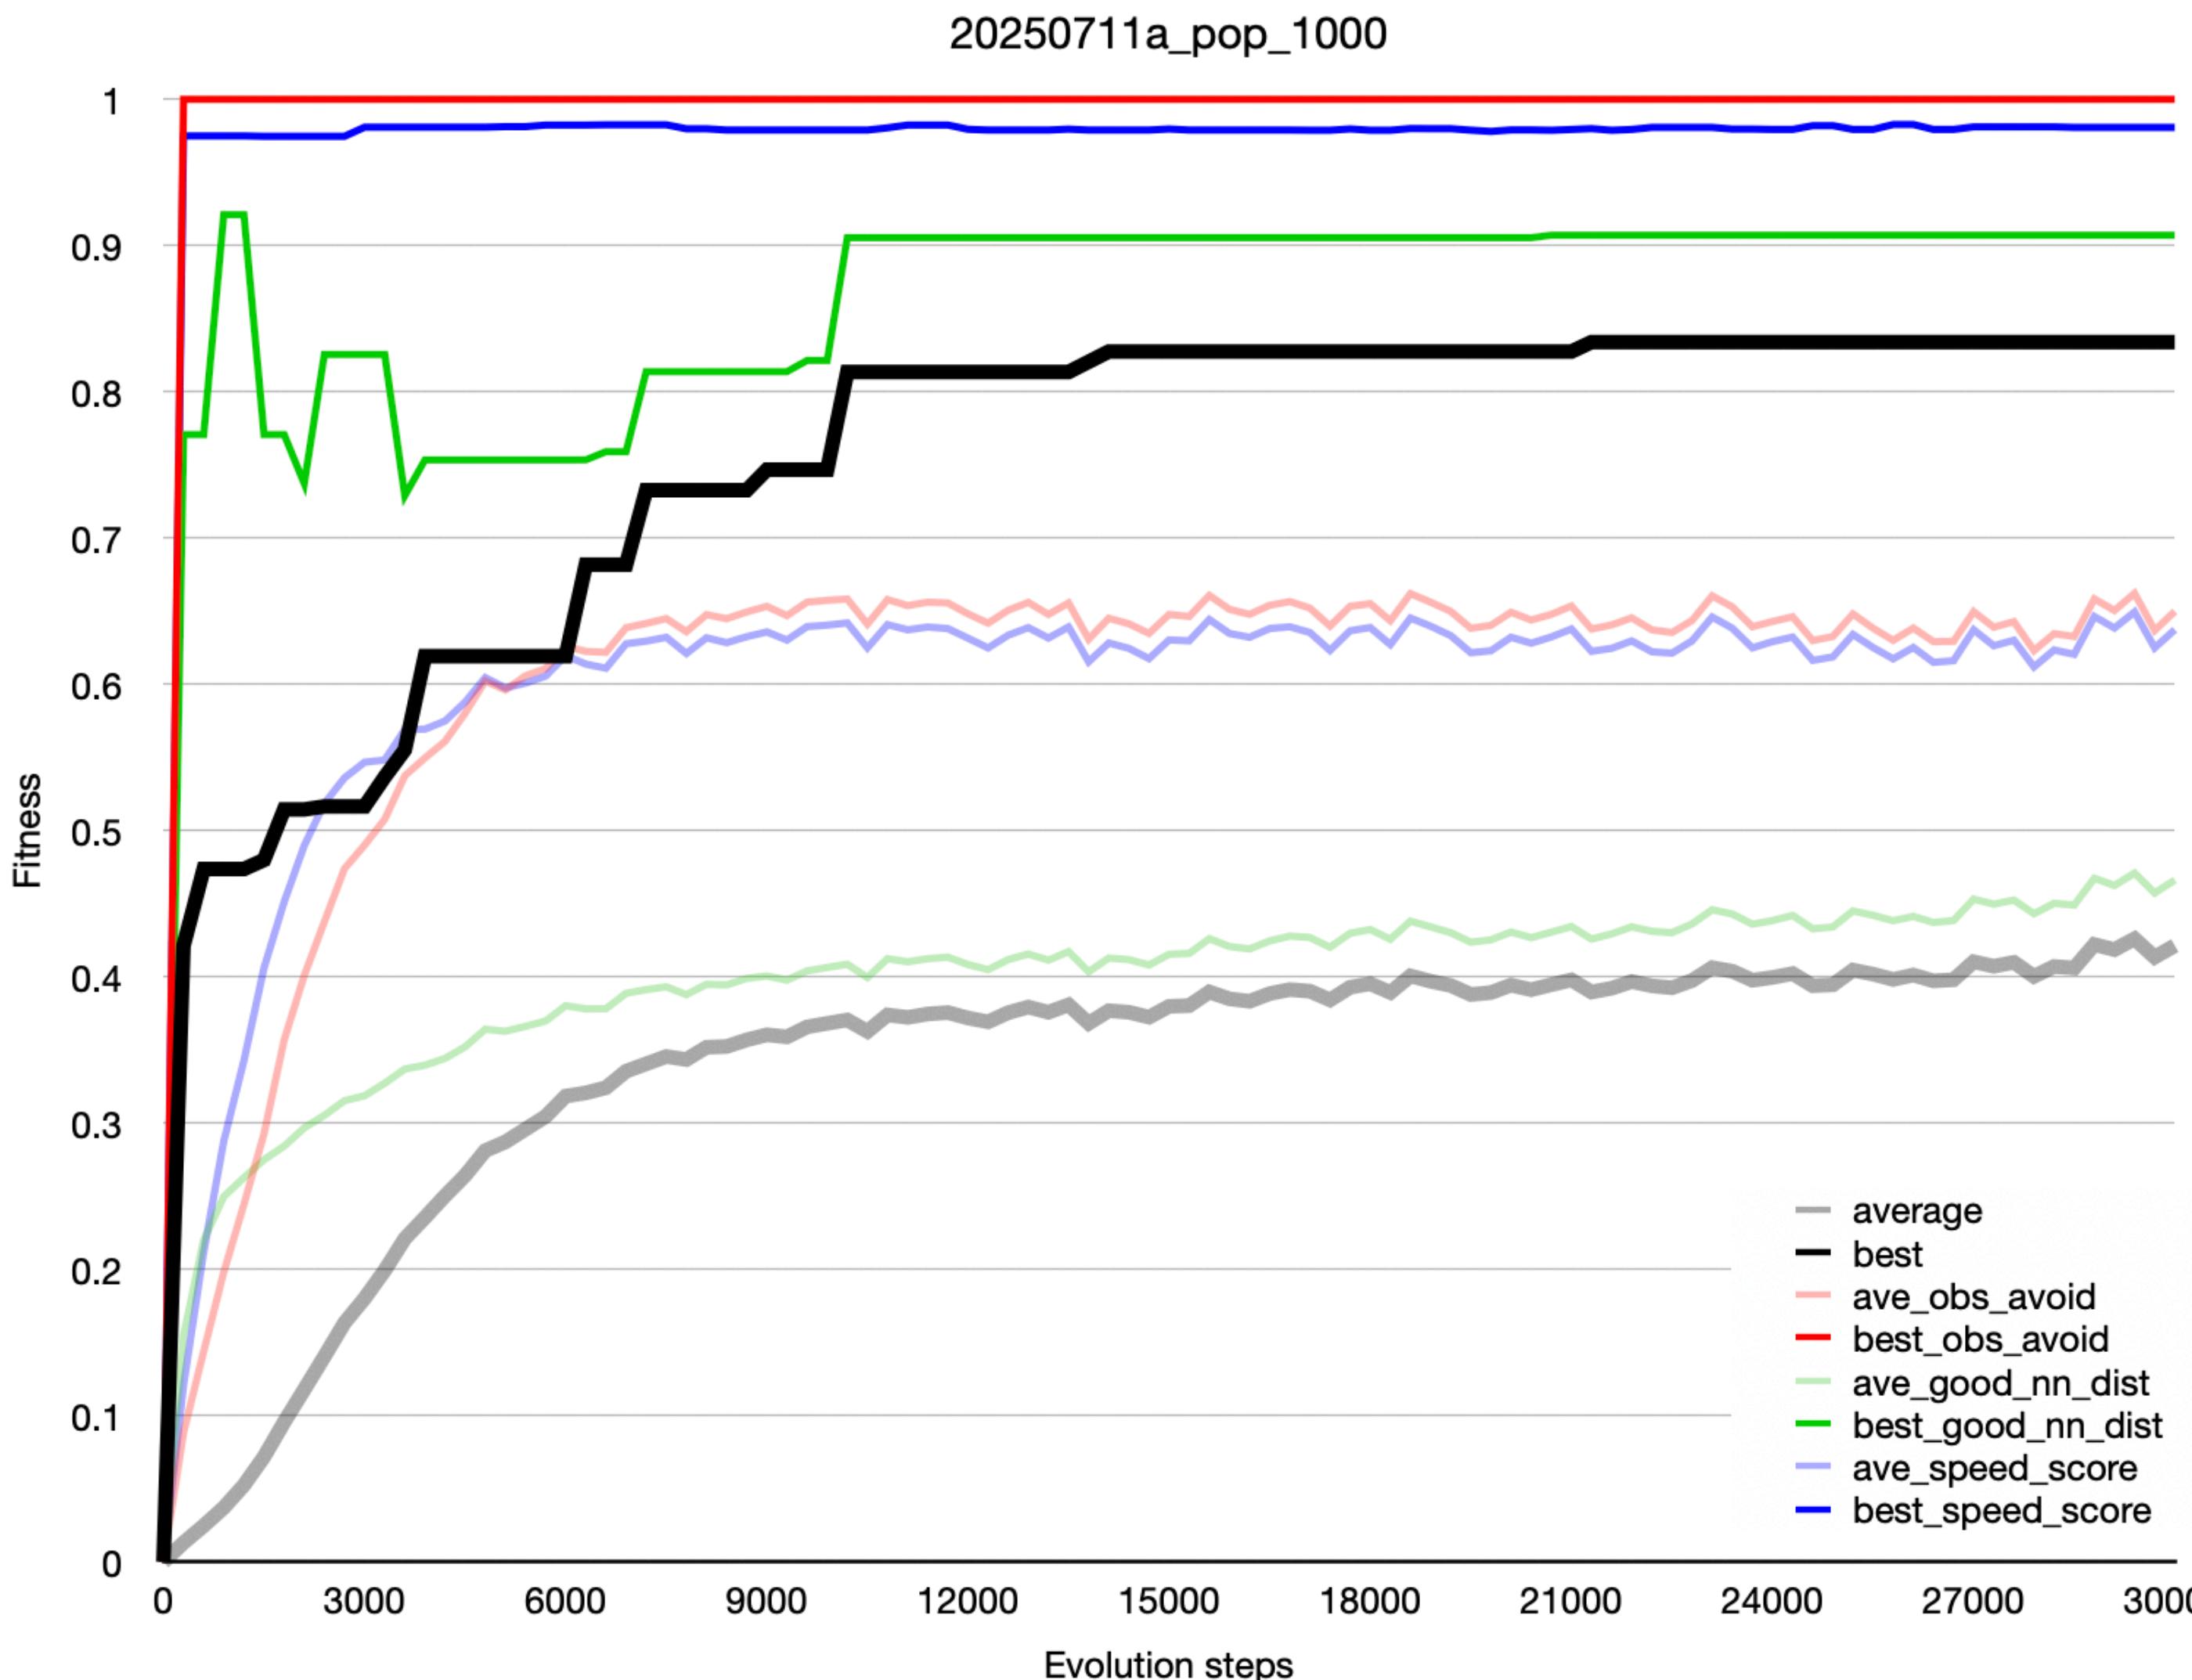
3 way tournament, negative selection, preserves diversity



EvoFlock fitness (and multiple objectives)



EvoFlock fitness (and multiple objectives)



Video 1

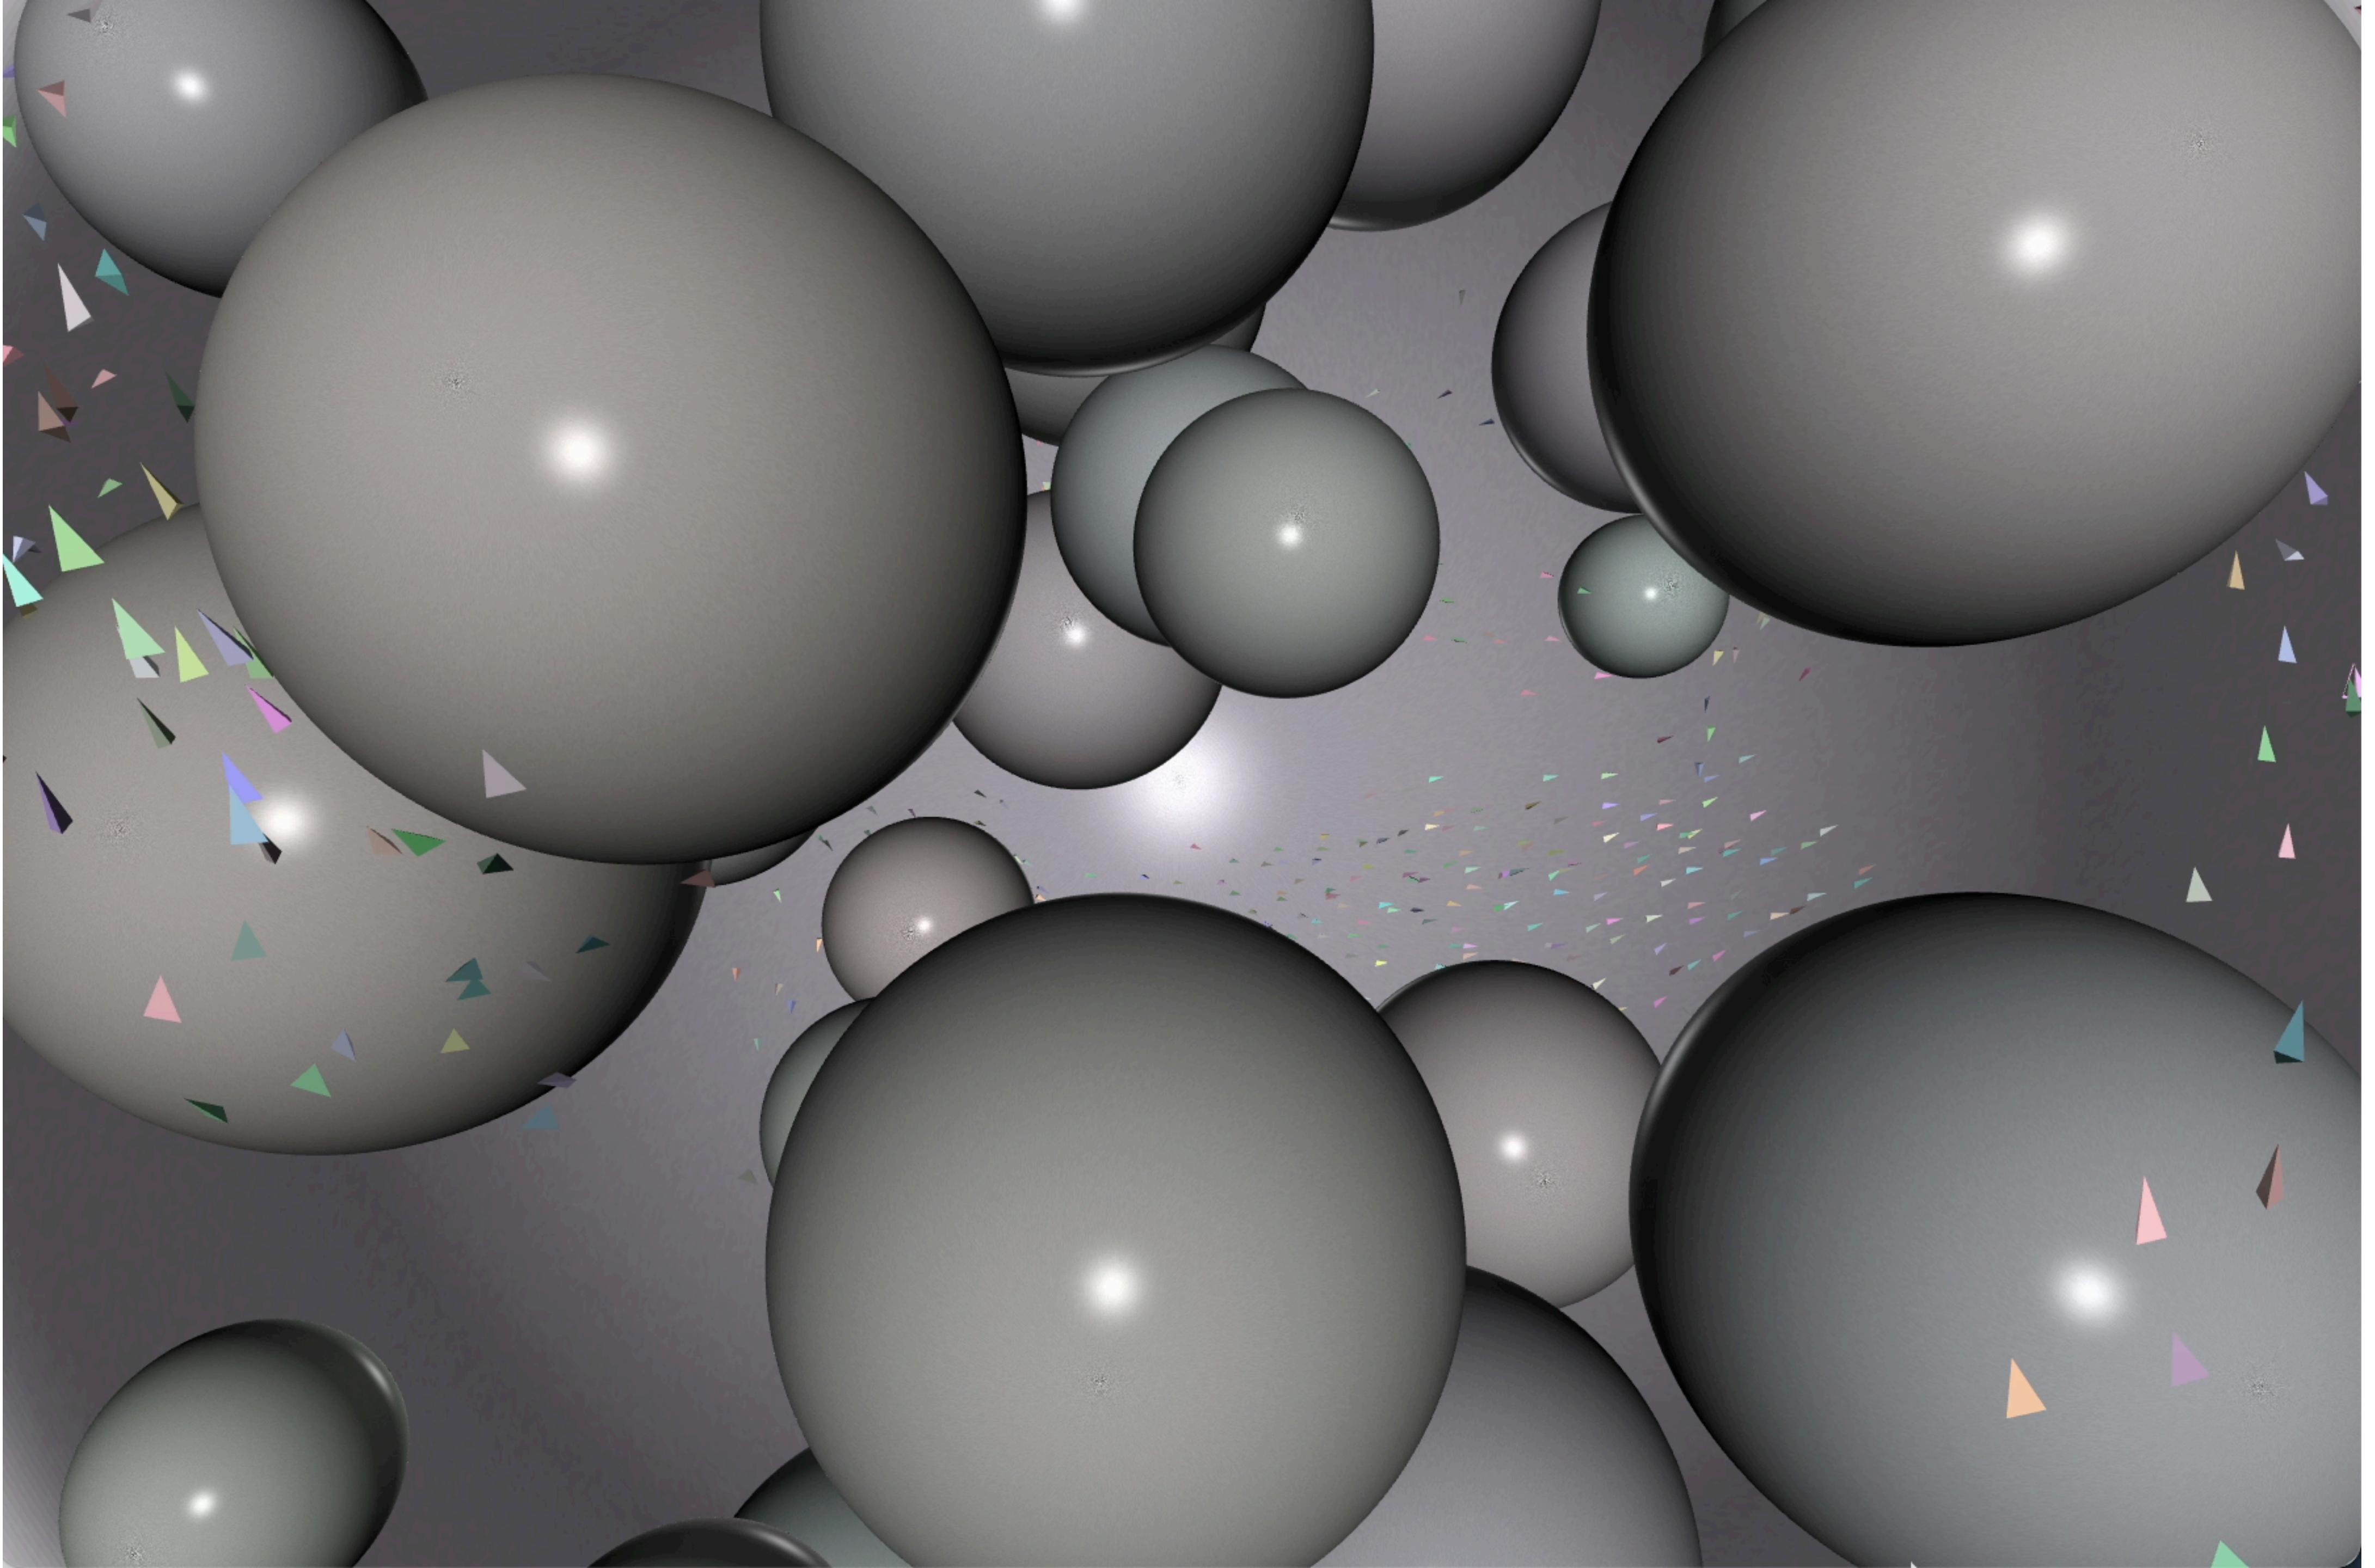
- Three objectives:
 - speed
 - separation
 - obstacle avoidance
- edit of several environments

Video 1 on web

Adding an objective: curvature

Video 2

- Four objectives:
 - speed
 - separation
 - obstacle avoidance
 - **curvature**
- edit of several environments



Video 2 on web

Thank you!

contact:

cwr@red3d.com

<https://www.red3d.com/cwr/>

these slides:

https://www.red3d.com/cwr/presentations/20250806_UCSC_inverse_flocks.pdf

